
Ray - Scalability from a Laptop to a Cluster

Dean Wampler - April 8, 2020

dean@anyscale.com

@deanwampler

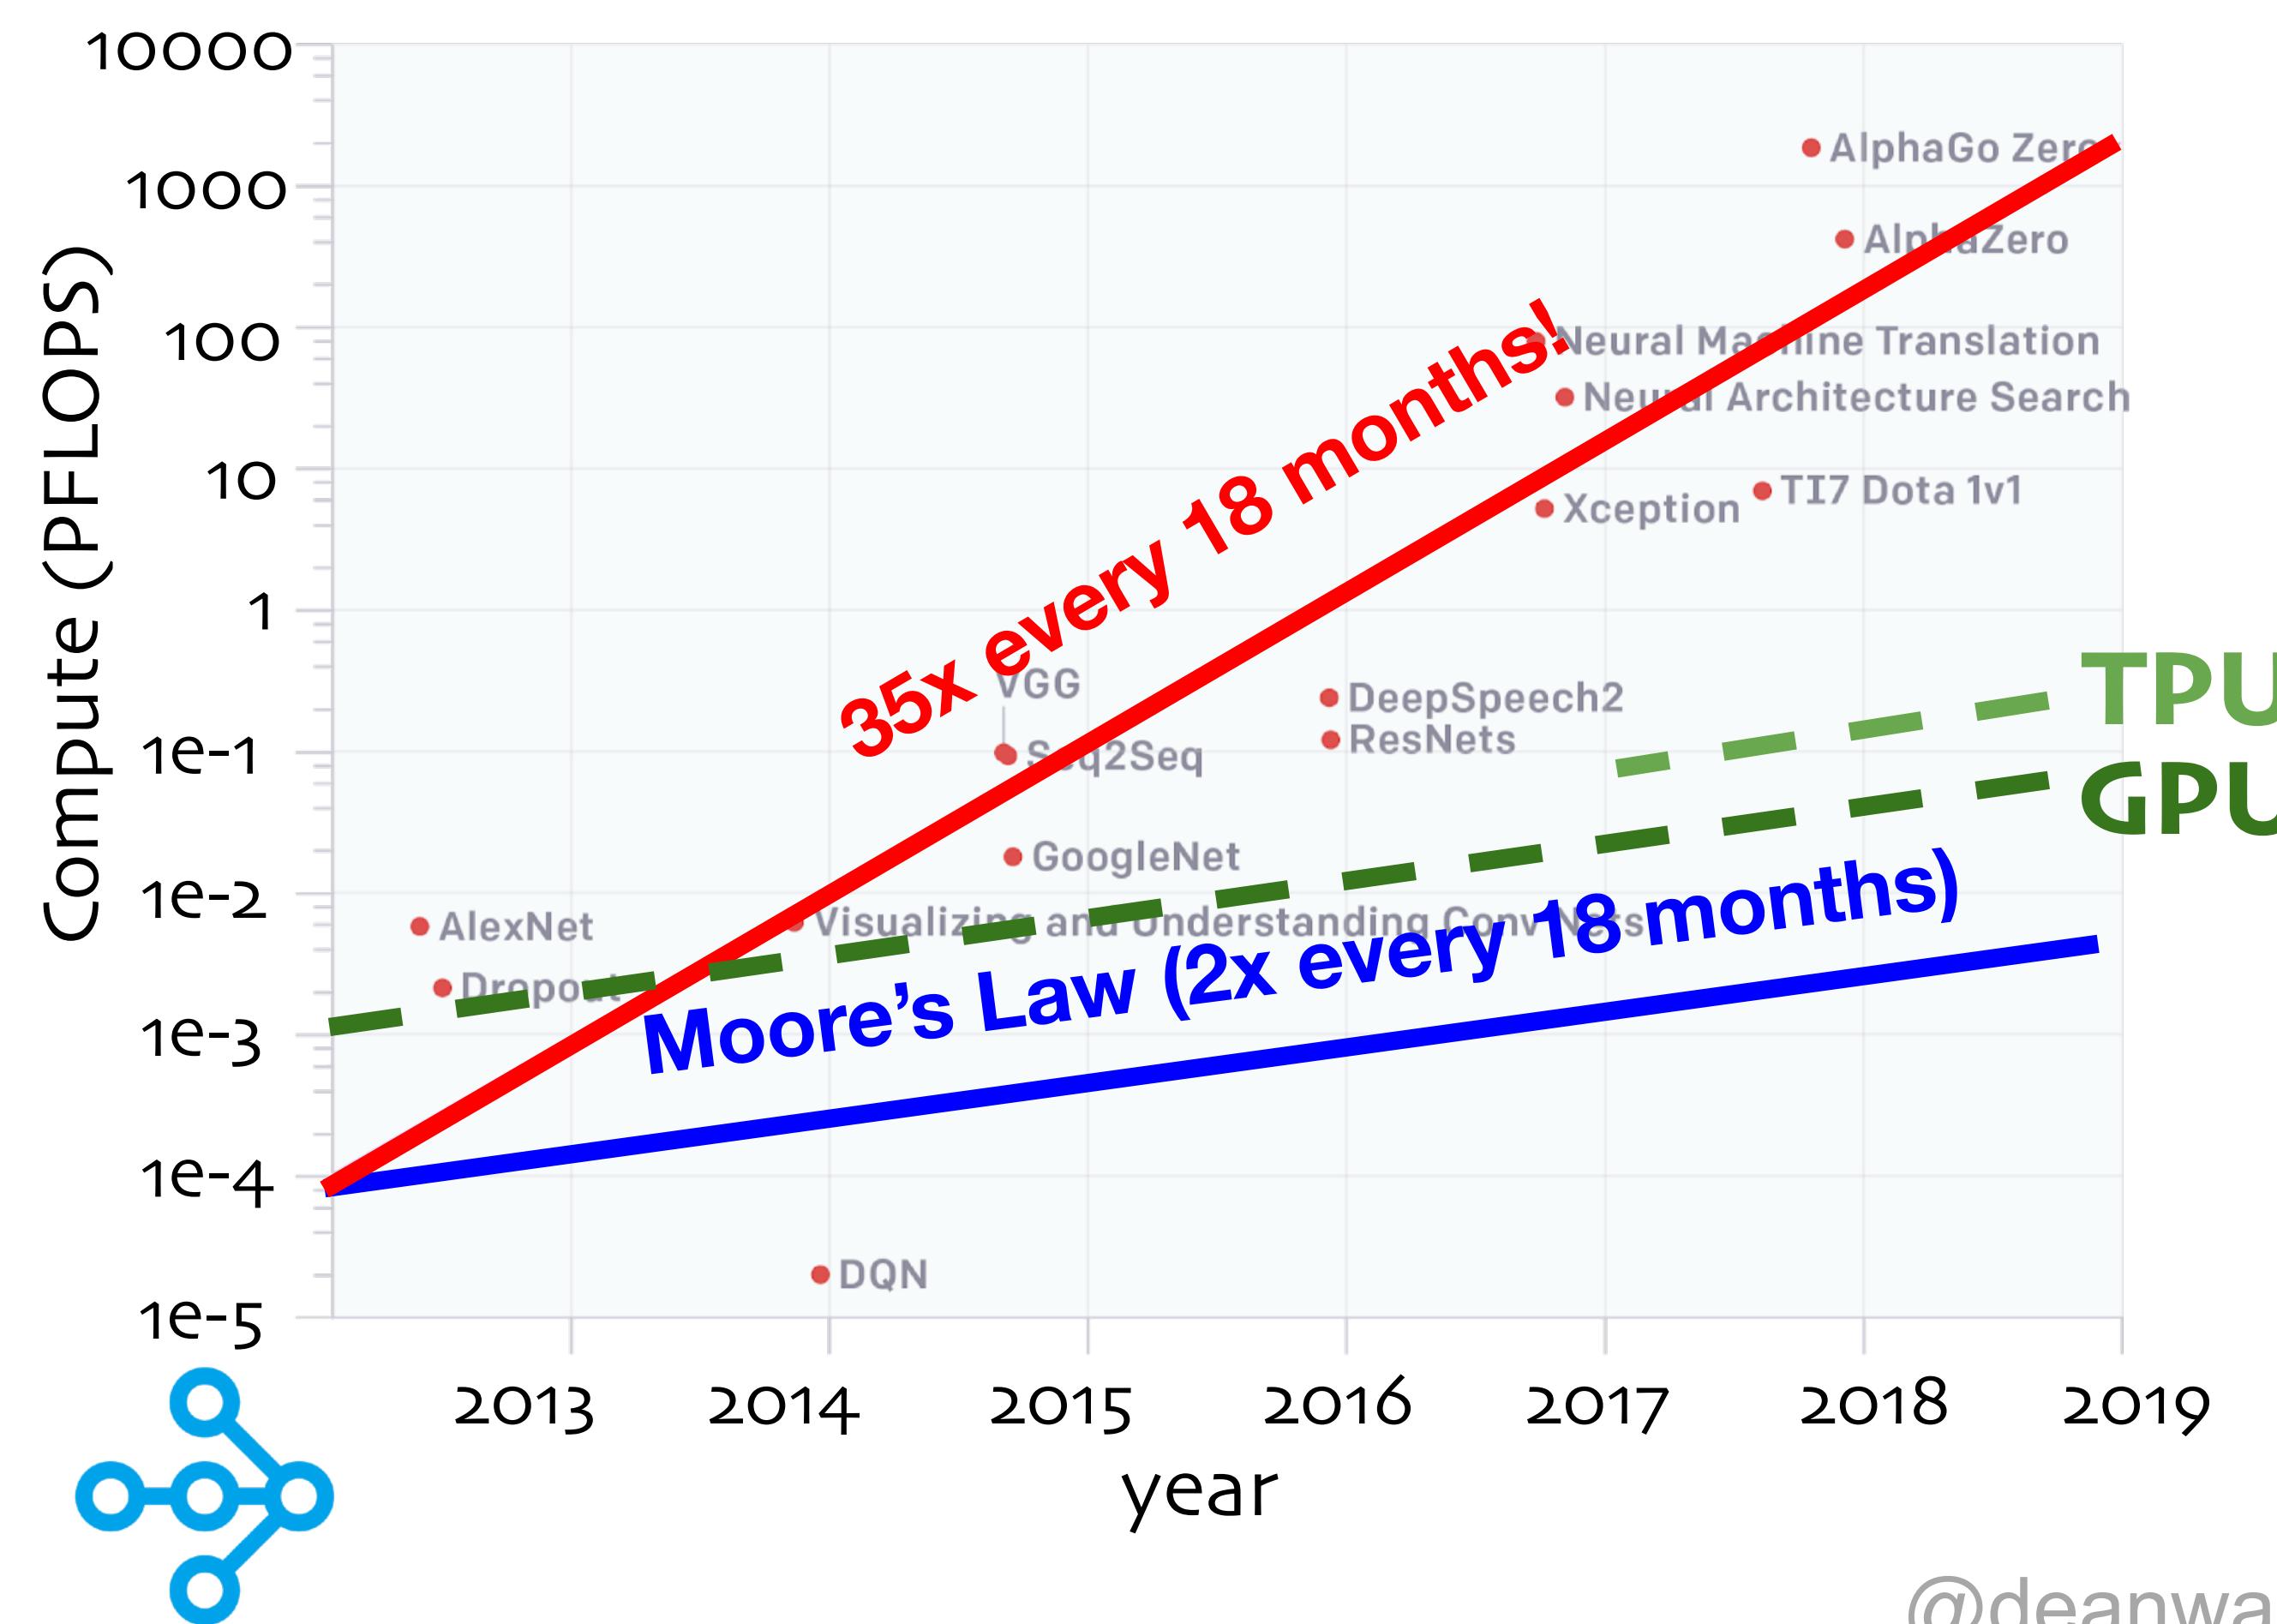
https://ray.io

https://anyscale.com

**Checkout our online events
this Summer:**

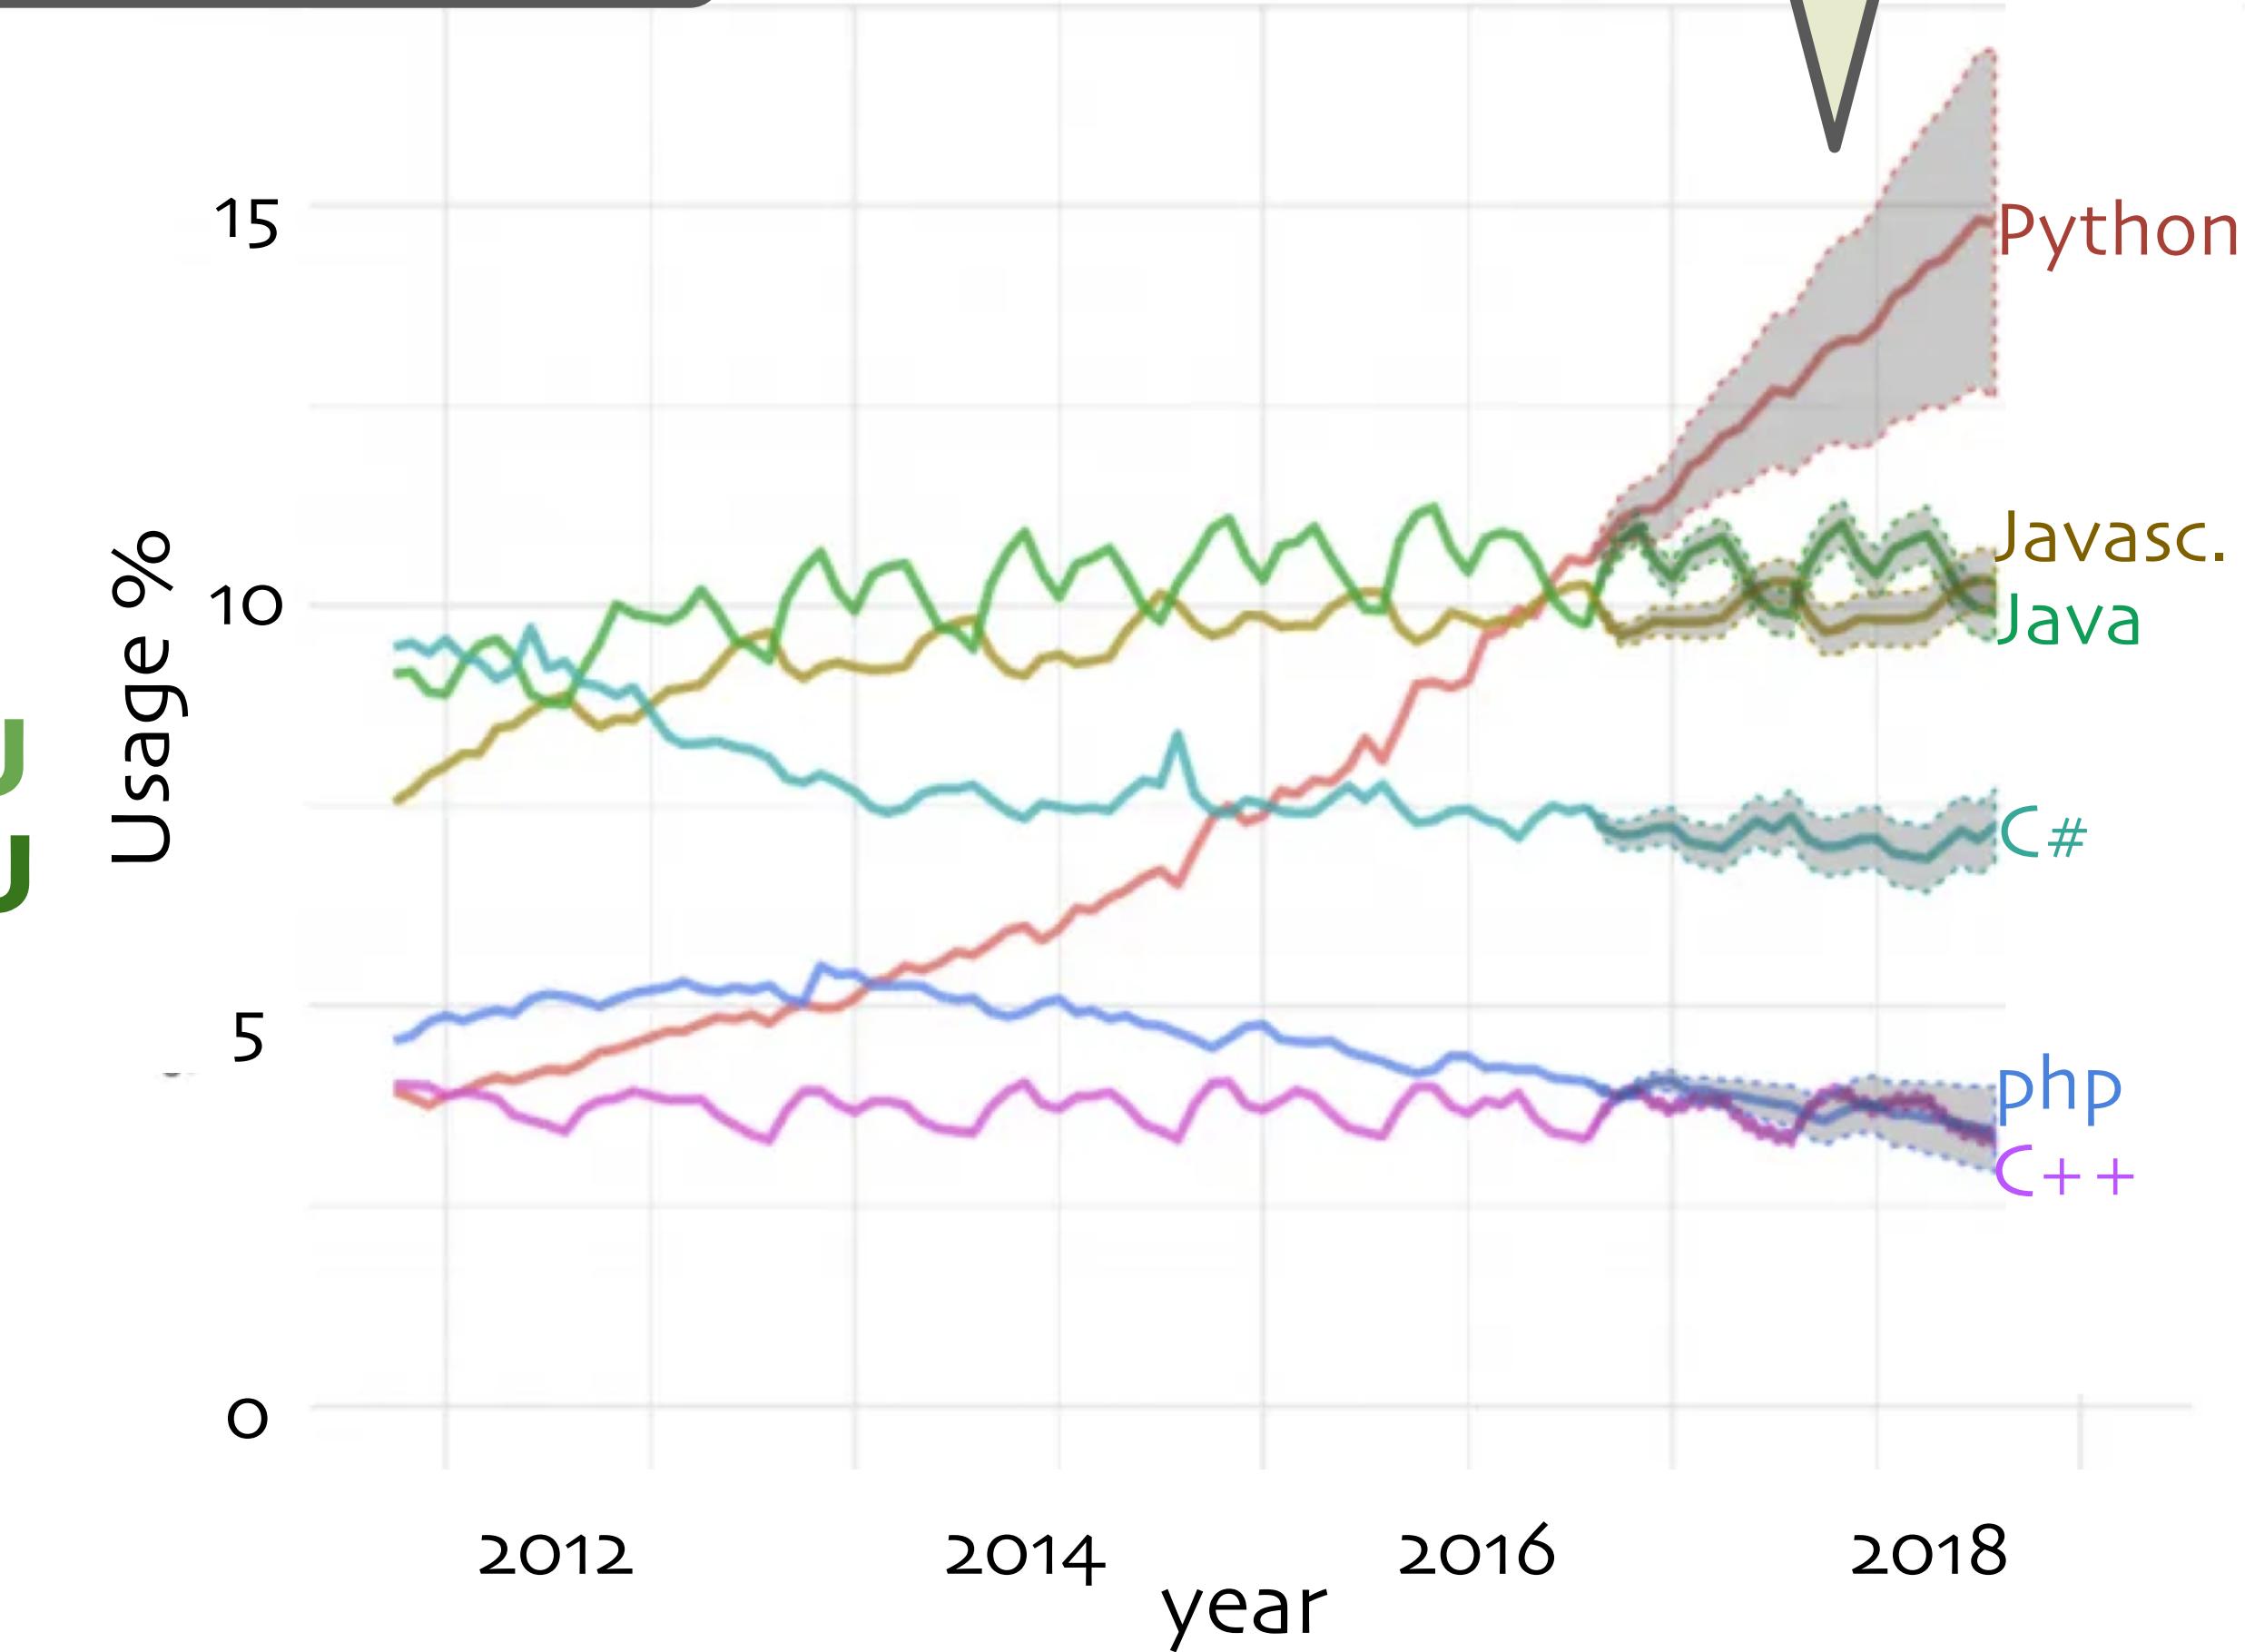
https://anyscale.com/events

Two Major Trends

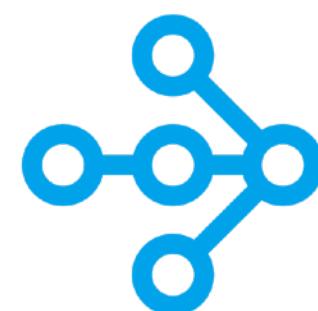
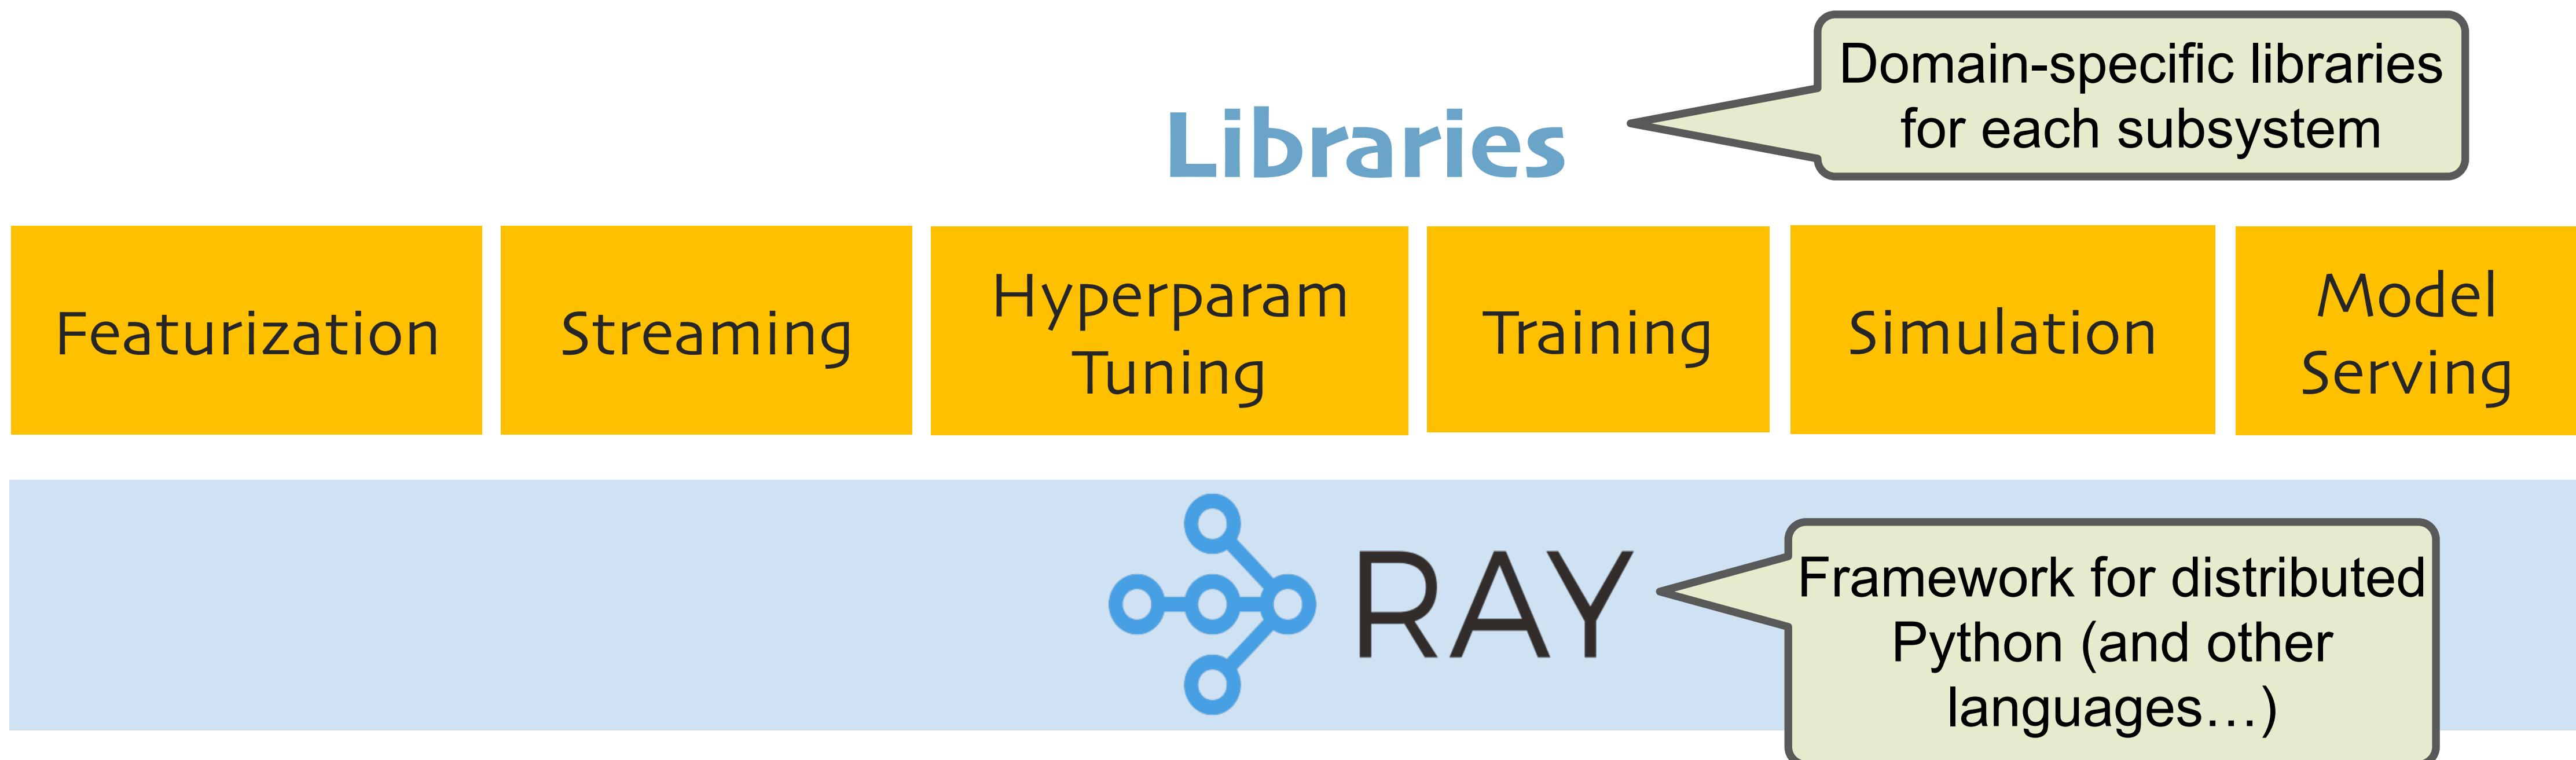


There is a pressing need
for robust, easy to use
solutions for distributed
Python

Driven in large part by
ML/AI and other data
science workloads



The Ray Vision: Sharing a Common Framework



API

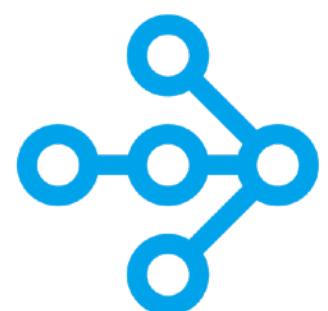
Functions -> Tasks

```
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

Designed to
intuitive and
concise

```
def add_arrays(a, b):  
    return np.add(a, b)
```

The Python you know
already...



API

Functions -> Tasks

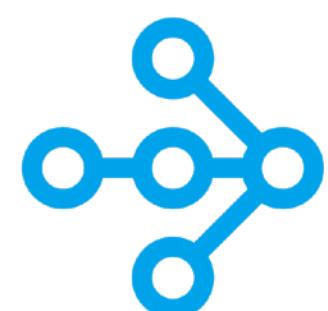
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

For completeness, add these First:

```
import ray  
import numpy as np  
ray.init()
```

Now these functions are
remote “tasks”



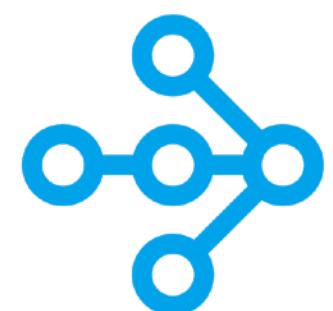
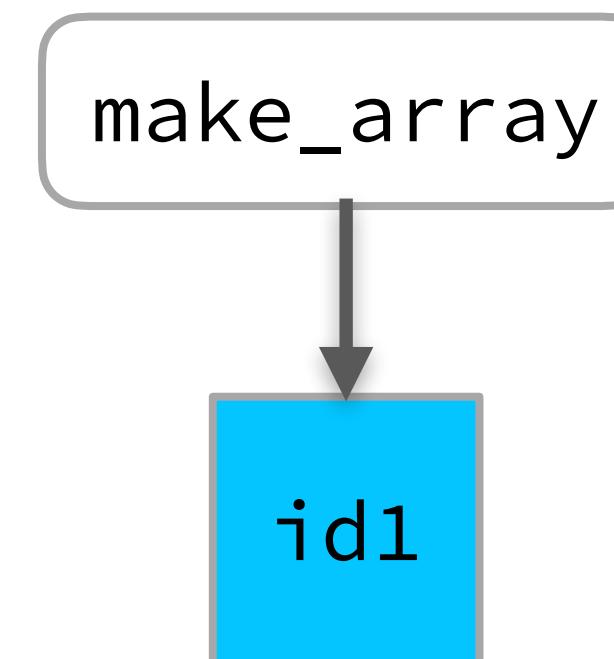
API

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)
```



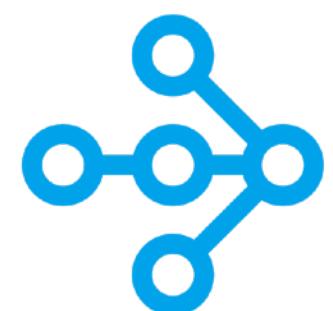
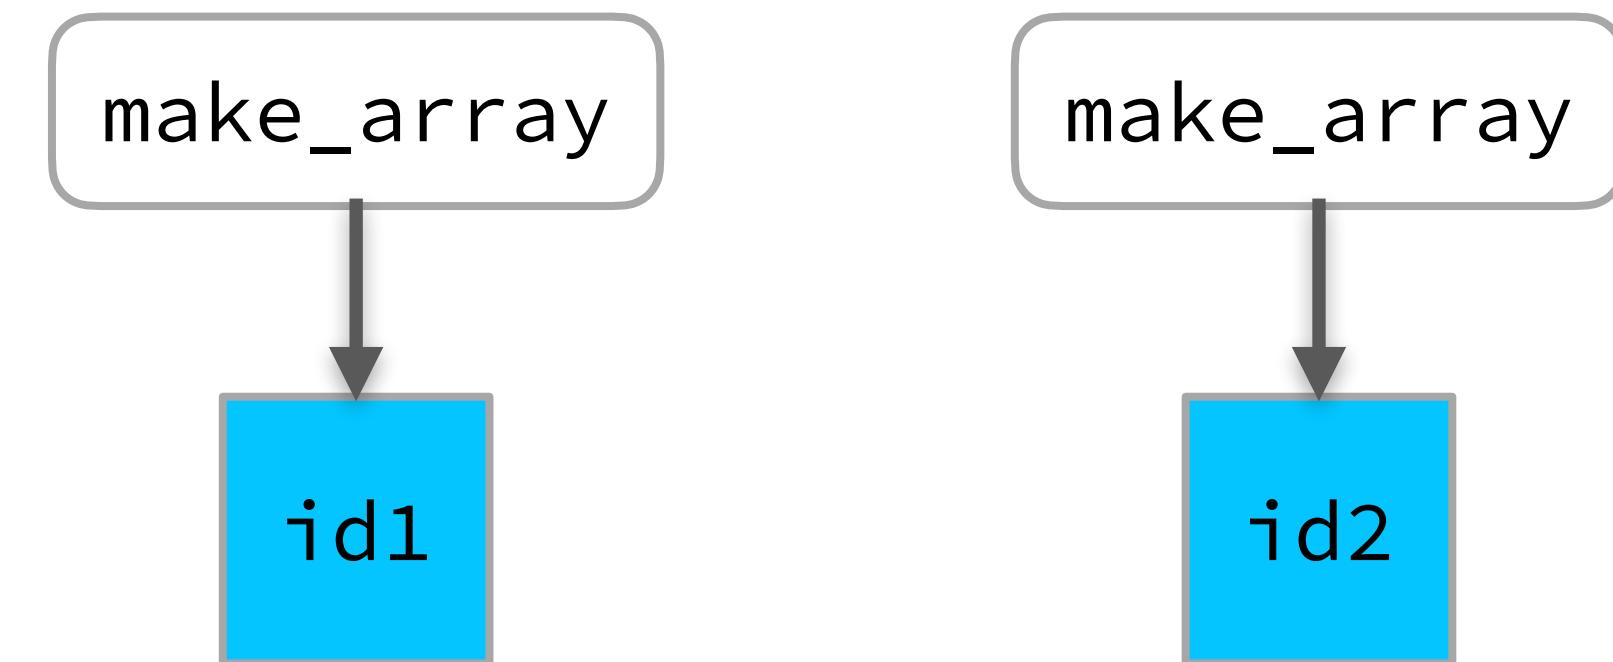
API

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)
```



API

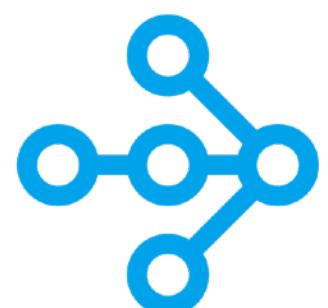
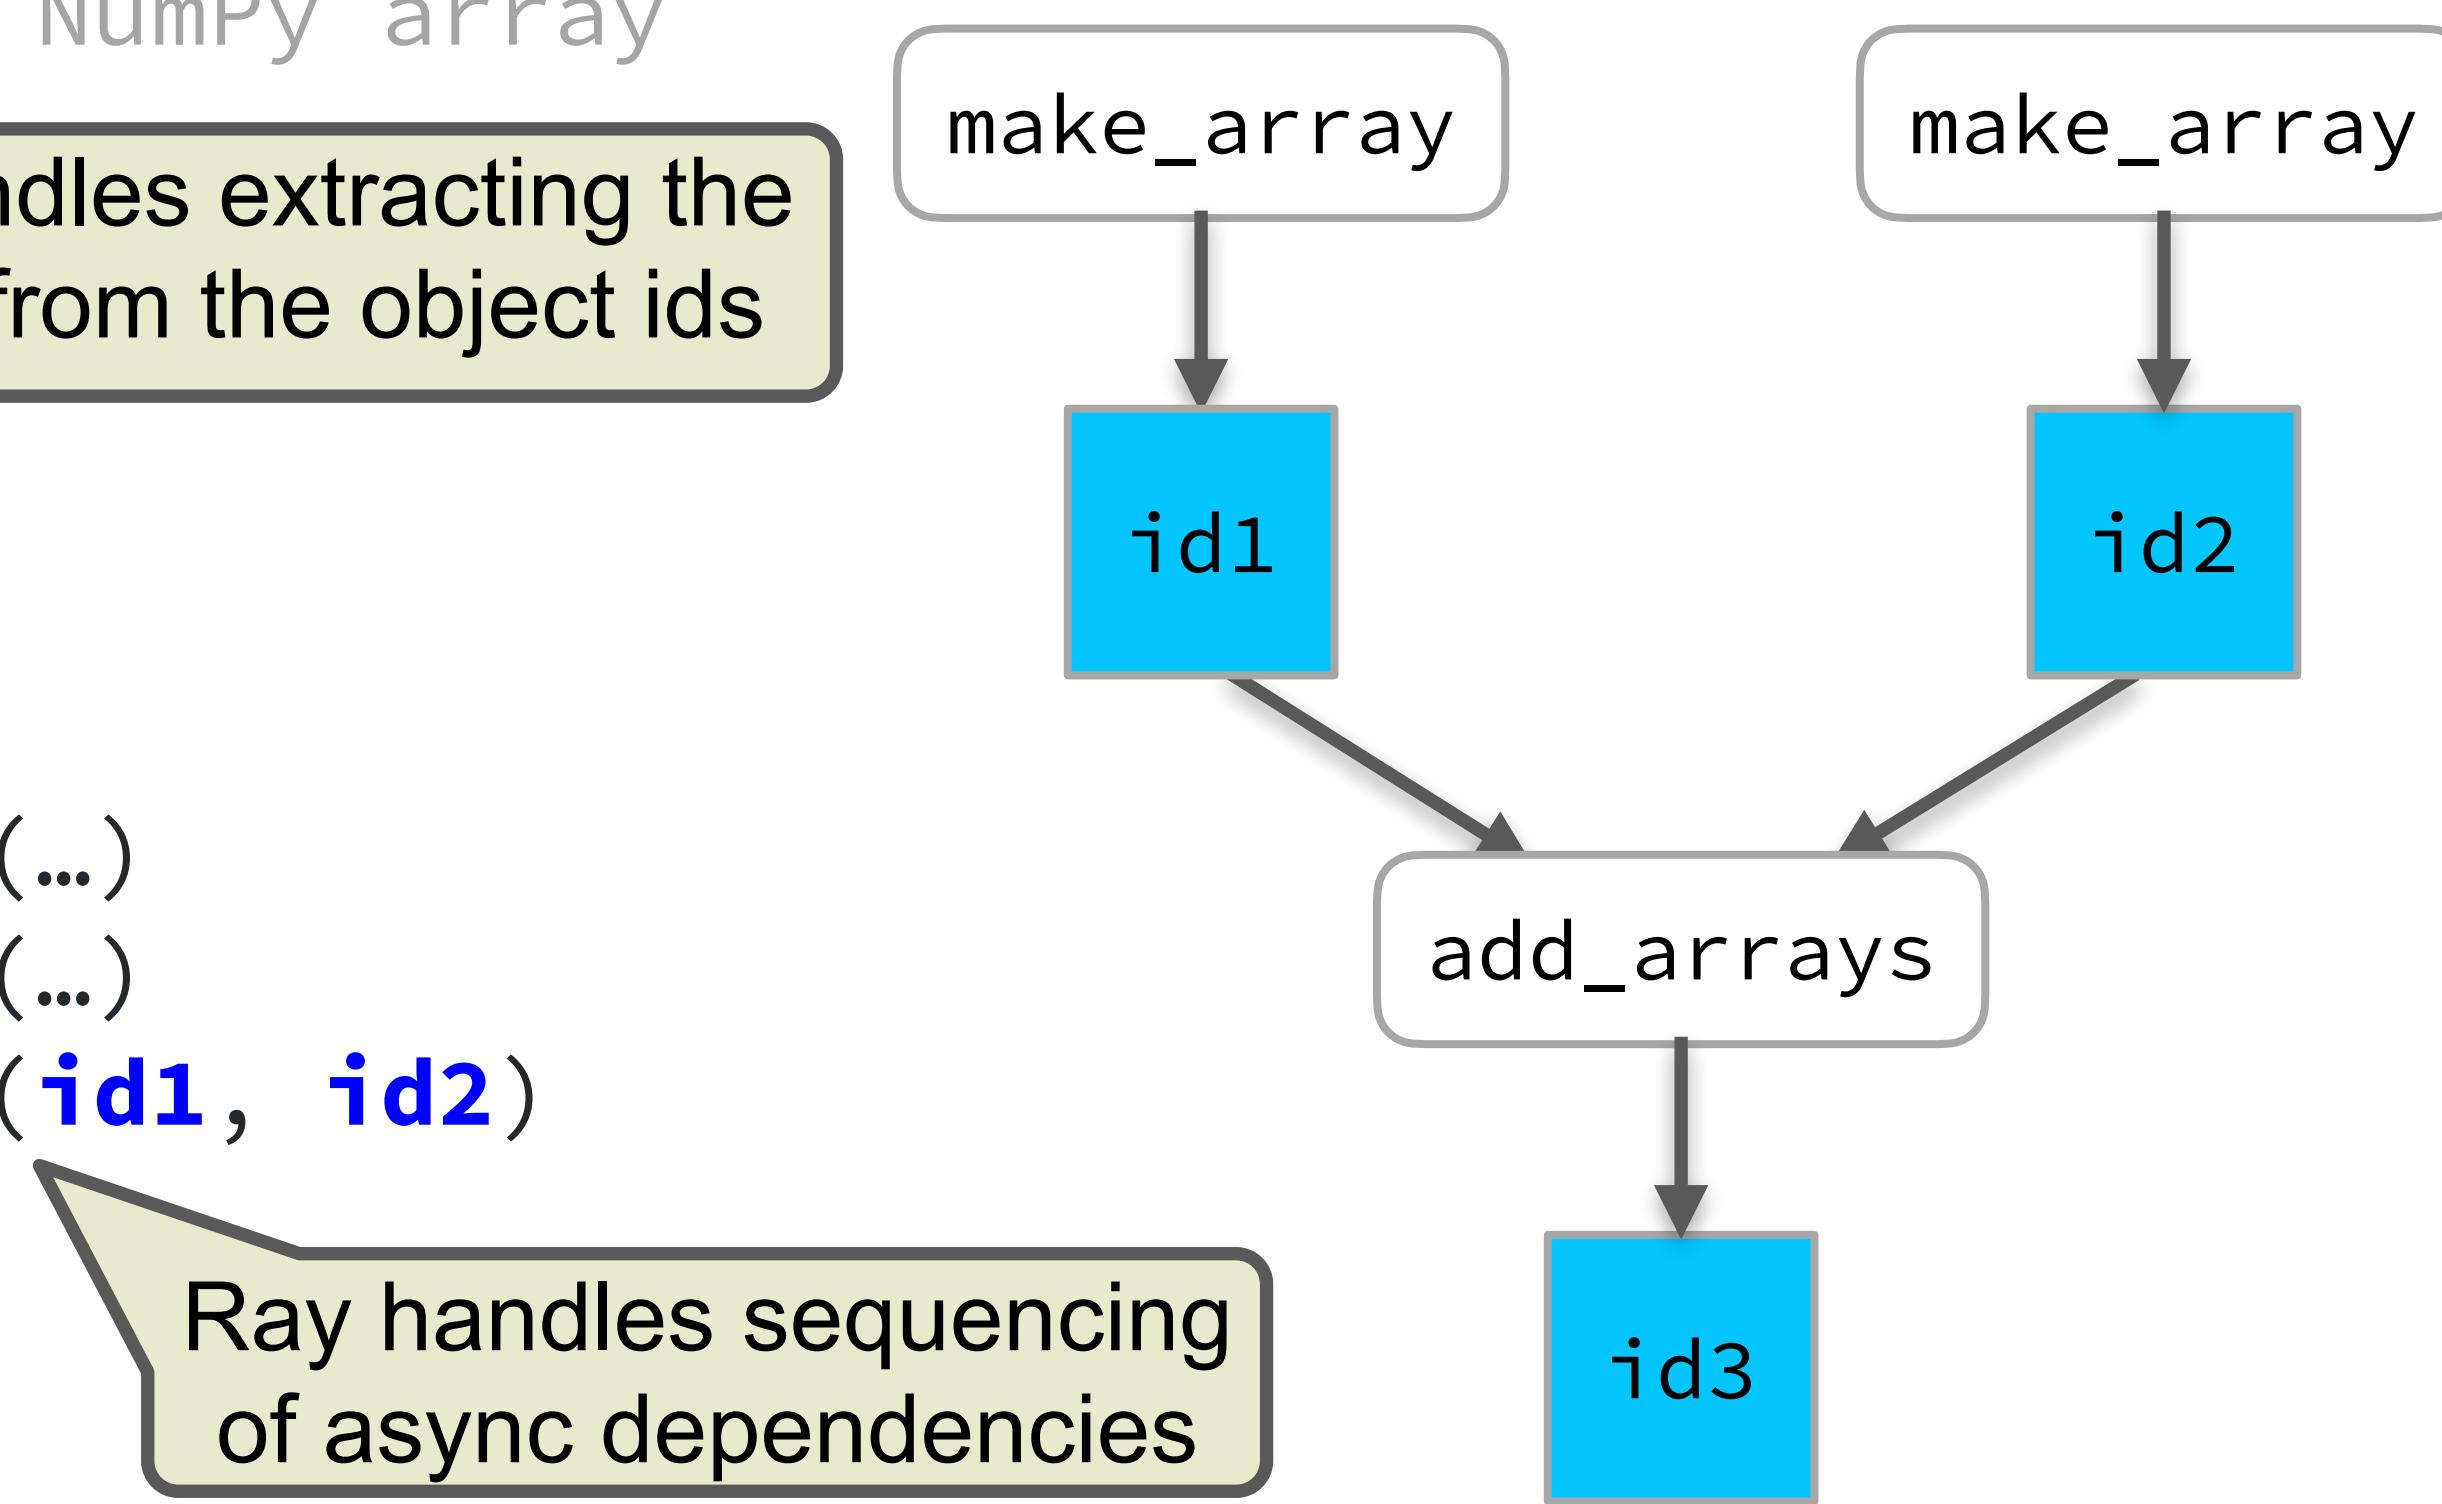
Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

Ray handles extracting the arrays from the object ids

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)
```



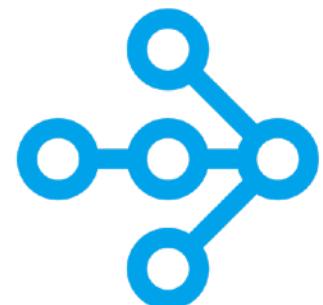
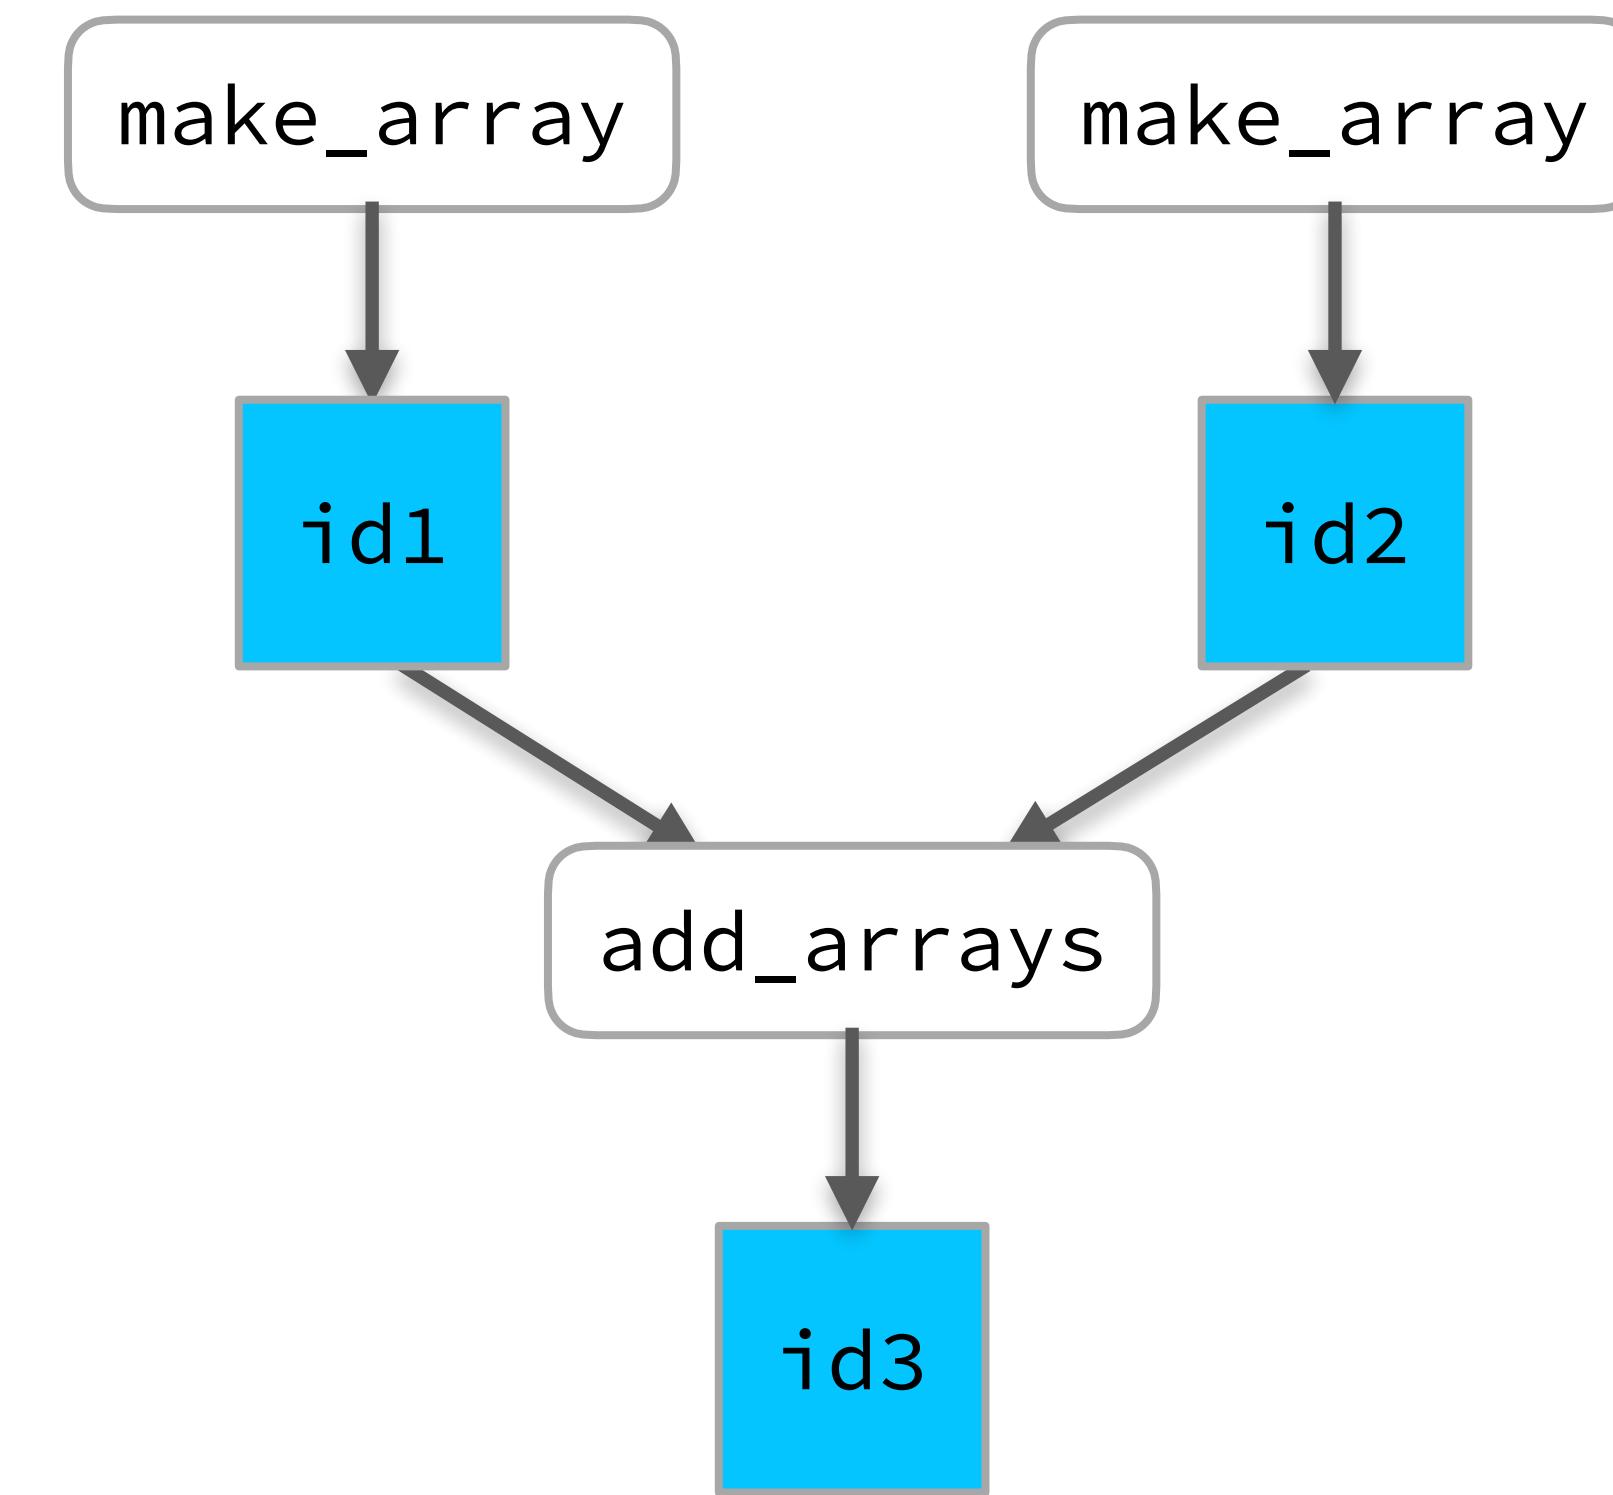
API

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```



API

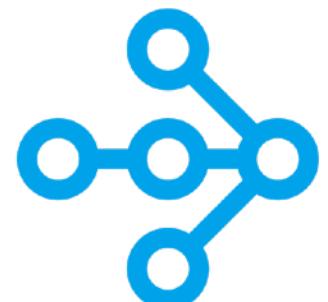
Functions -> Tasks

What about
distributed
state?

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a
```

```
@ray.remote
def add_arrays(a, b):
    return np.add(a, b)
```

```
id1 = make_array.remote(...)
id2 = make_array.remote(...)
id3 = add_arrays.remote(id1, id2)
ray.get(id3)
```



API

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

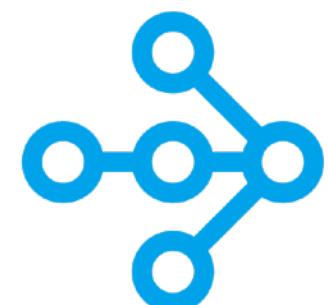
```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```

The Python
classes you
love...

Classes -> Actors

```
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
    return self.value
```

What about
distributed
state?



API

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

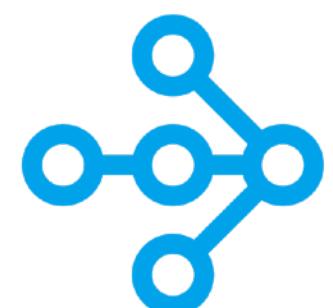
```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
id1 = make_array.remote(...)  
id2 = make_array.remote(...)  
id3 = add_arrays.remote(id1, id2)  
ray.get(id3)
```

Classes -> Actors

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
    return self.value
```

... now a remote
“actor”



API

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

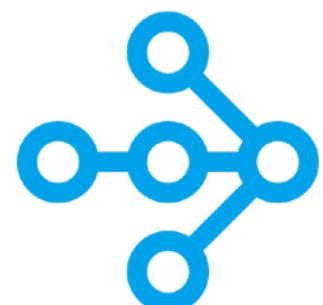
@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

id1 = make_array.remote(...)
id2 = make_array.remote(...)
id3 = add_arrays.remote(id1, id2)
ray.get(id3)
```

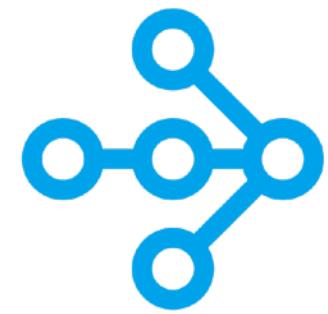
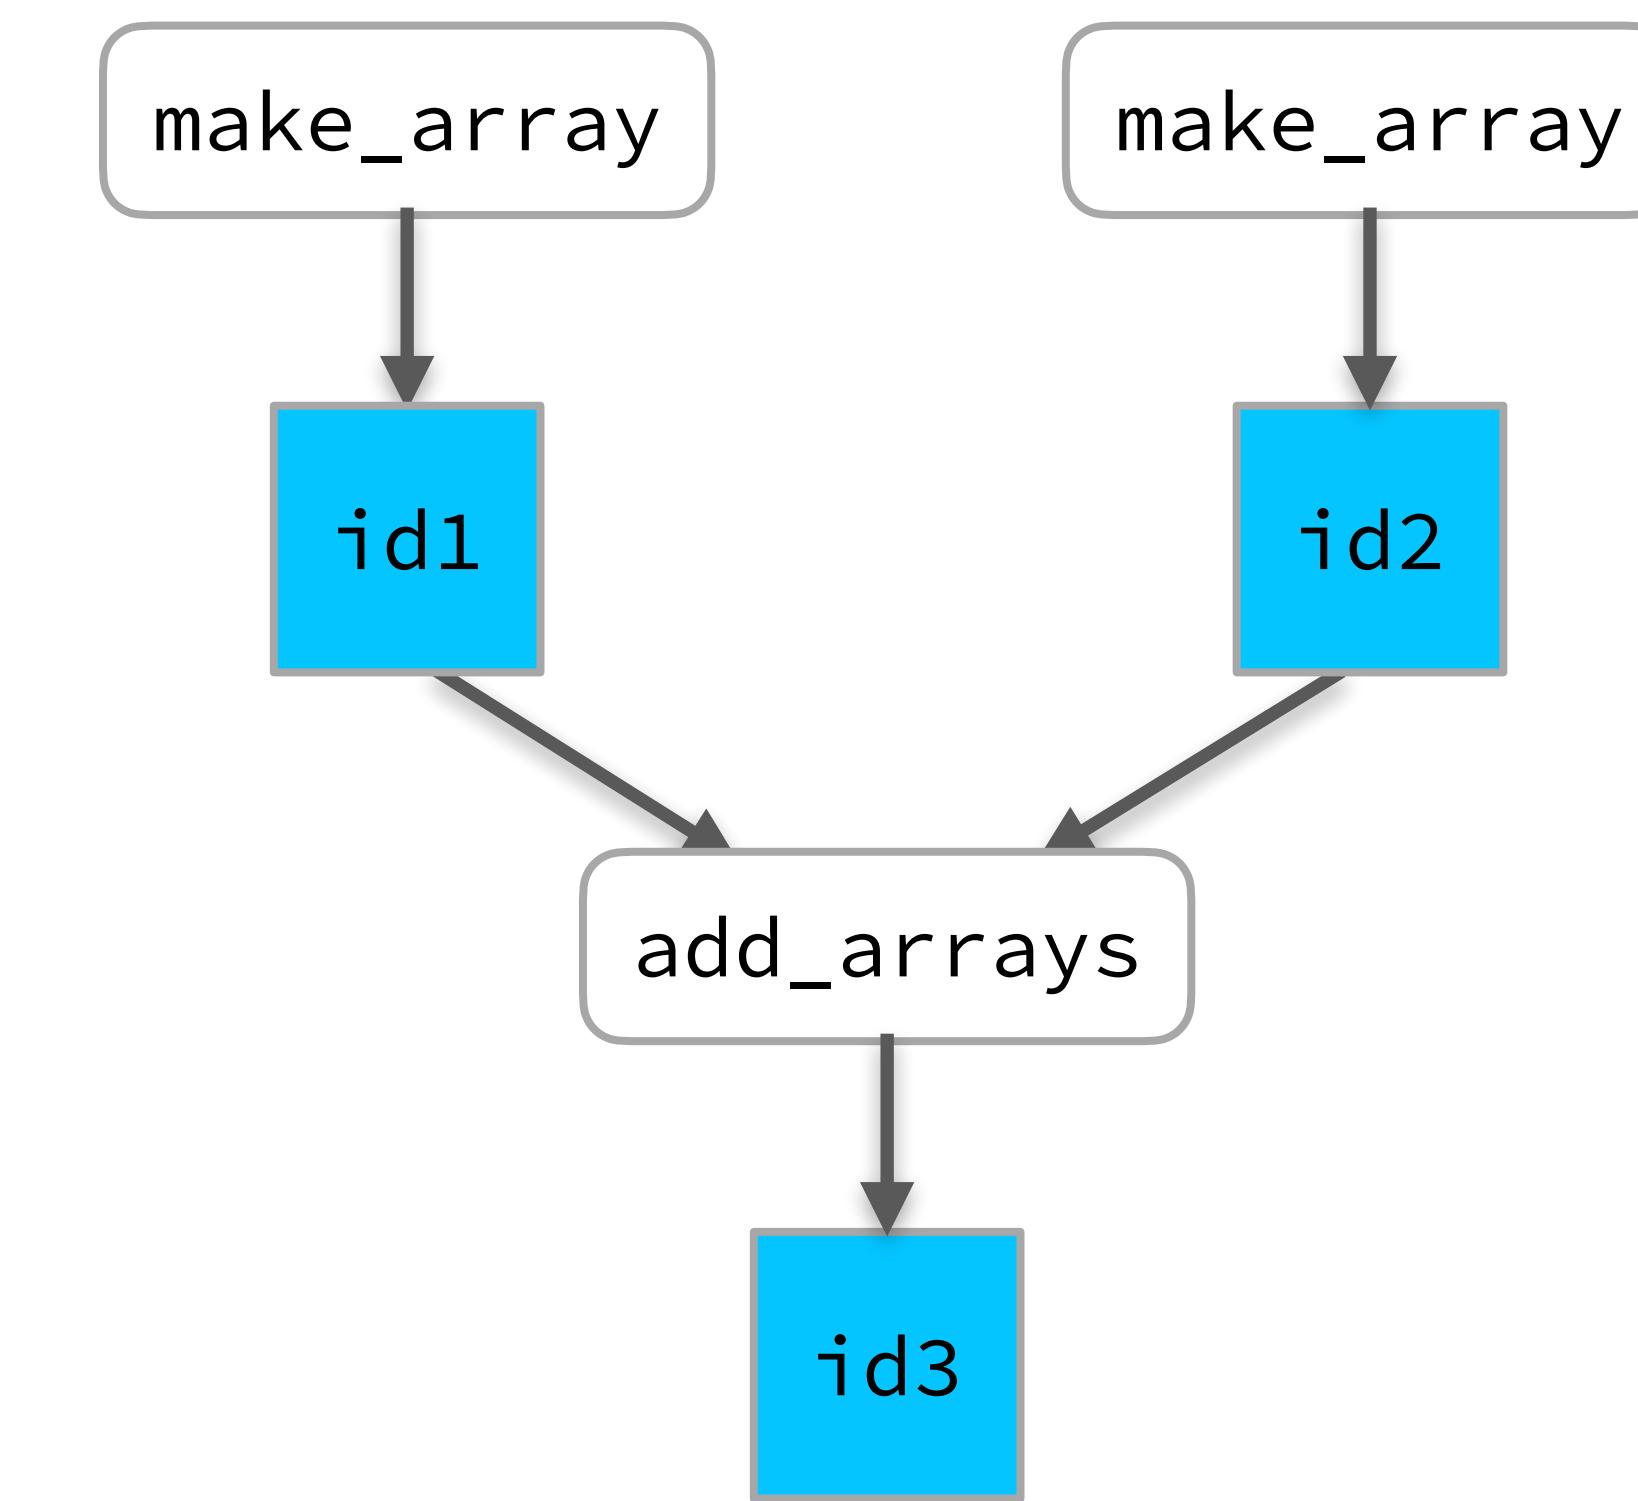
Classes -> Actors

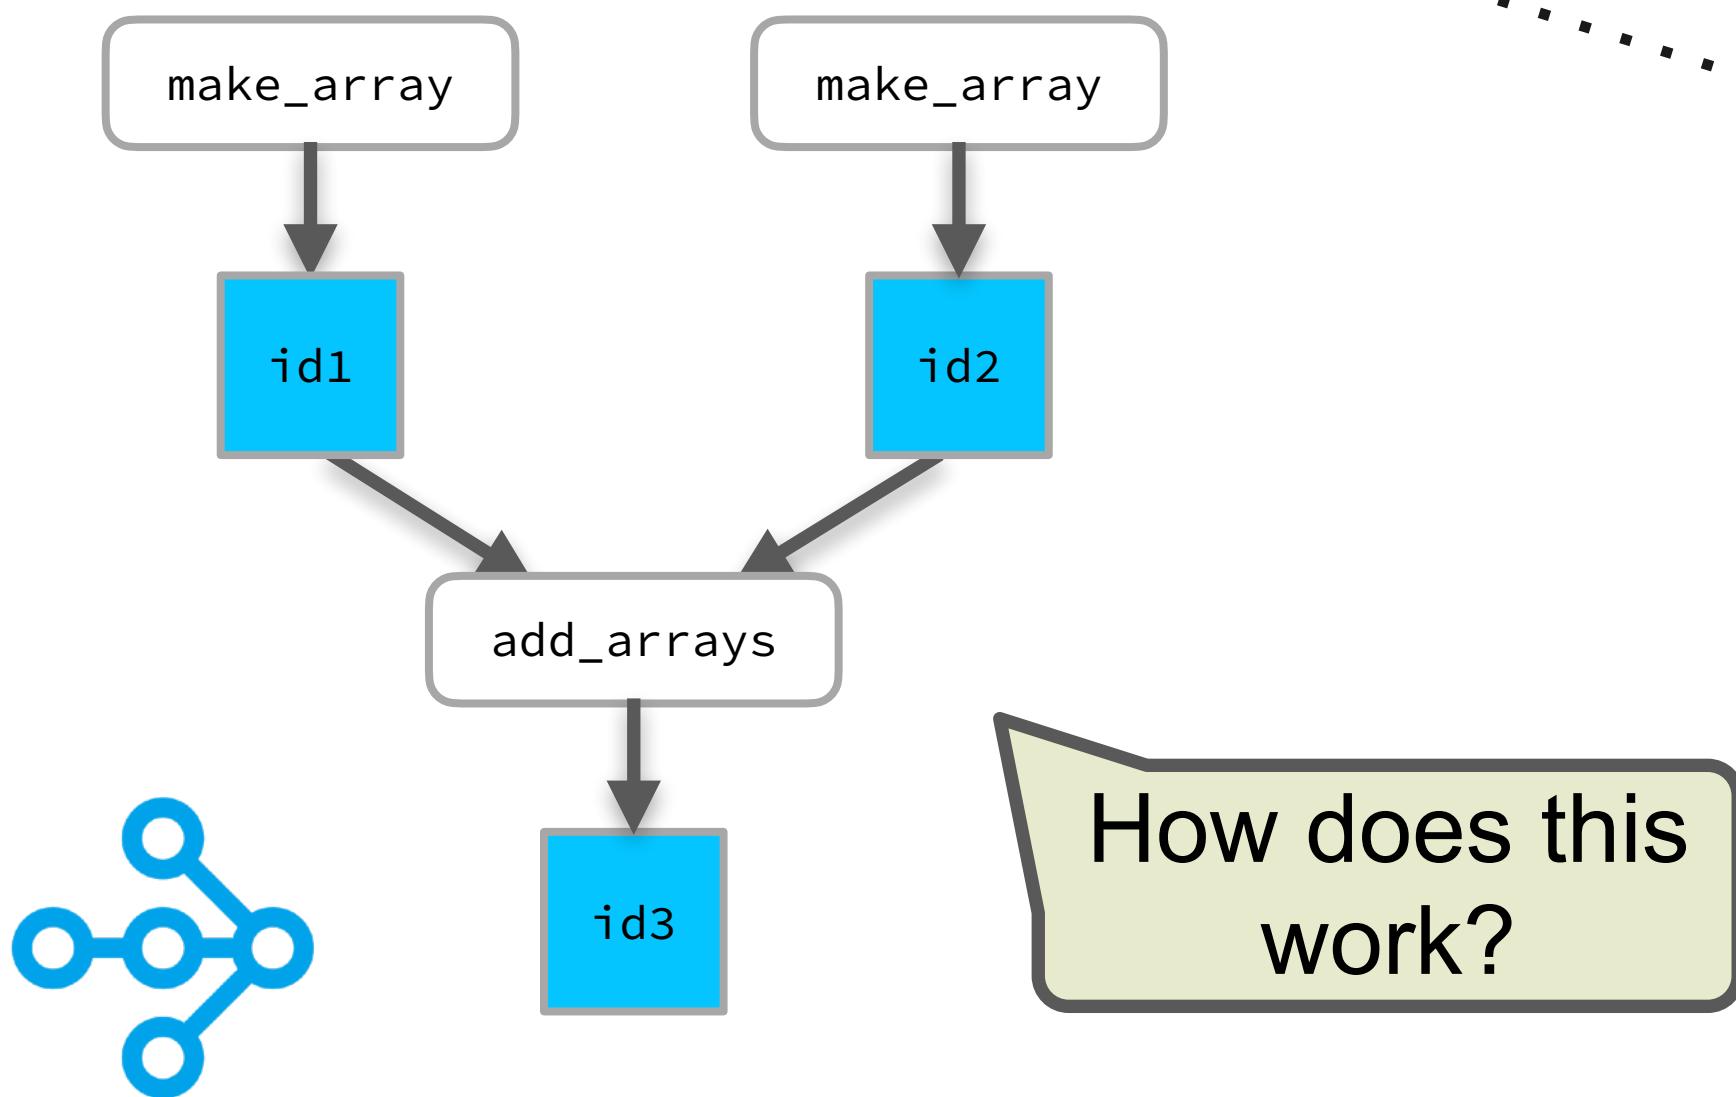
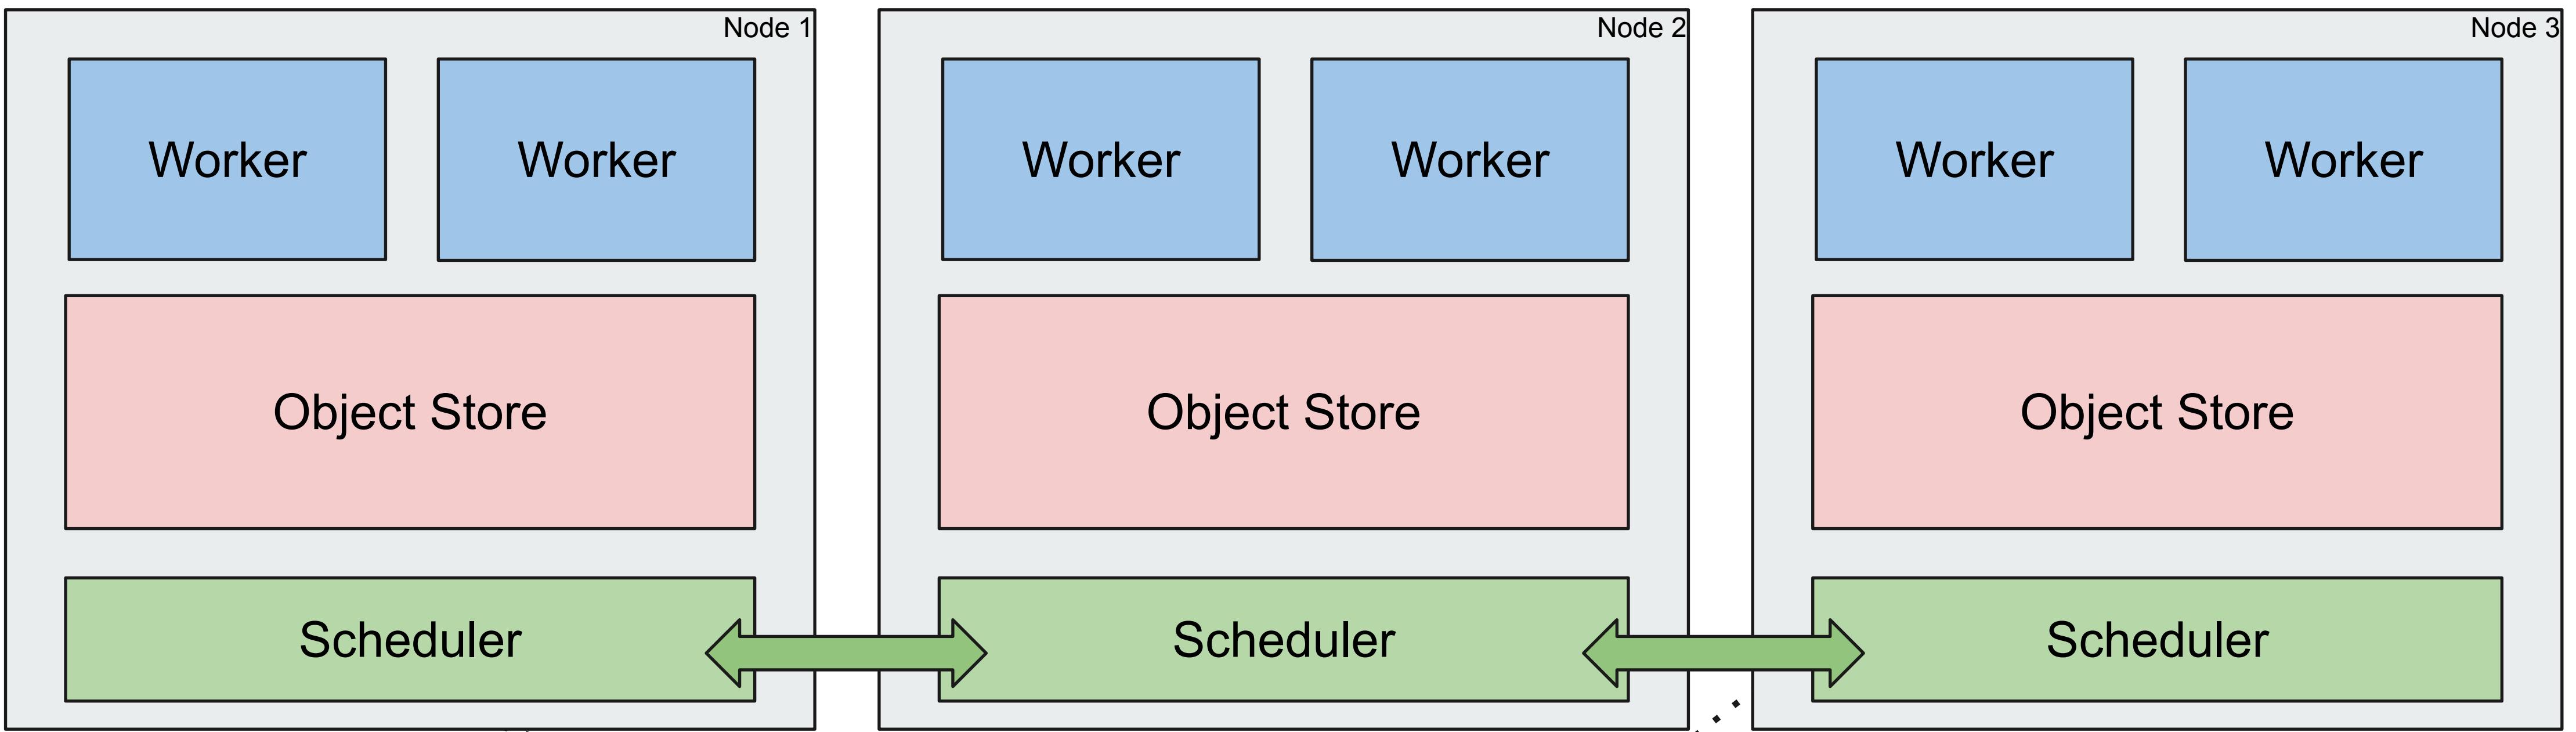
```
@ray.remote
class Counter(object):
    def __init__(self):
        self.value = 0
    def increment(self):
        self.value += 1
        return self.value

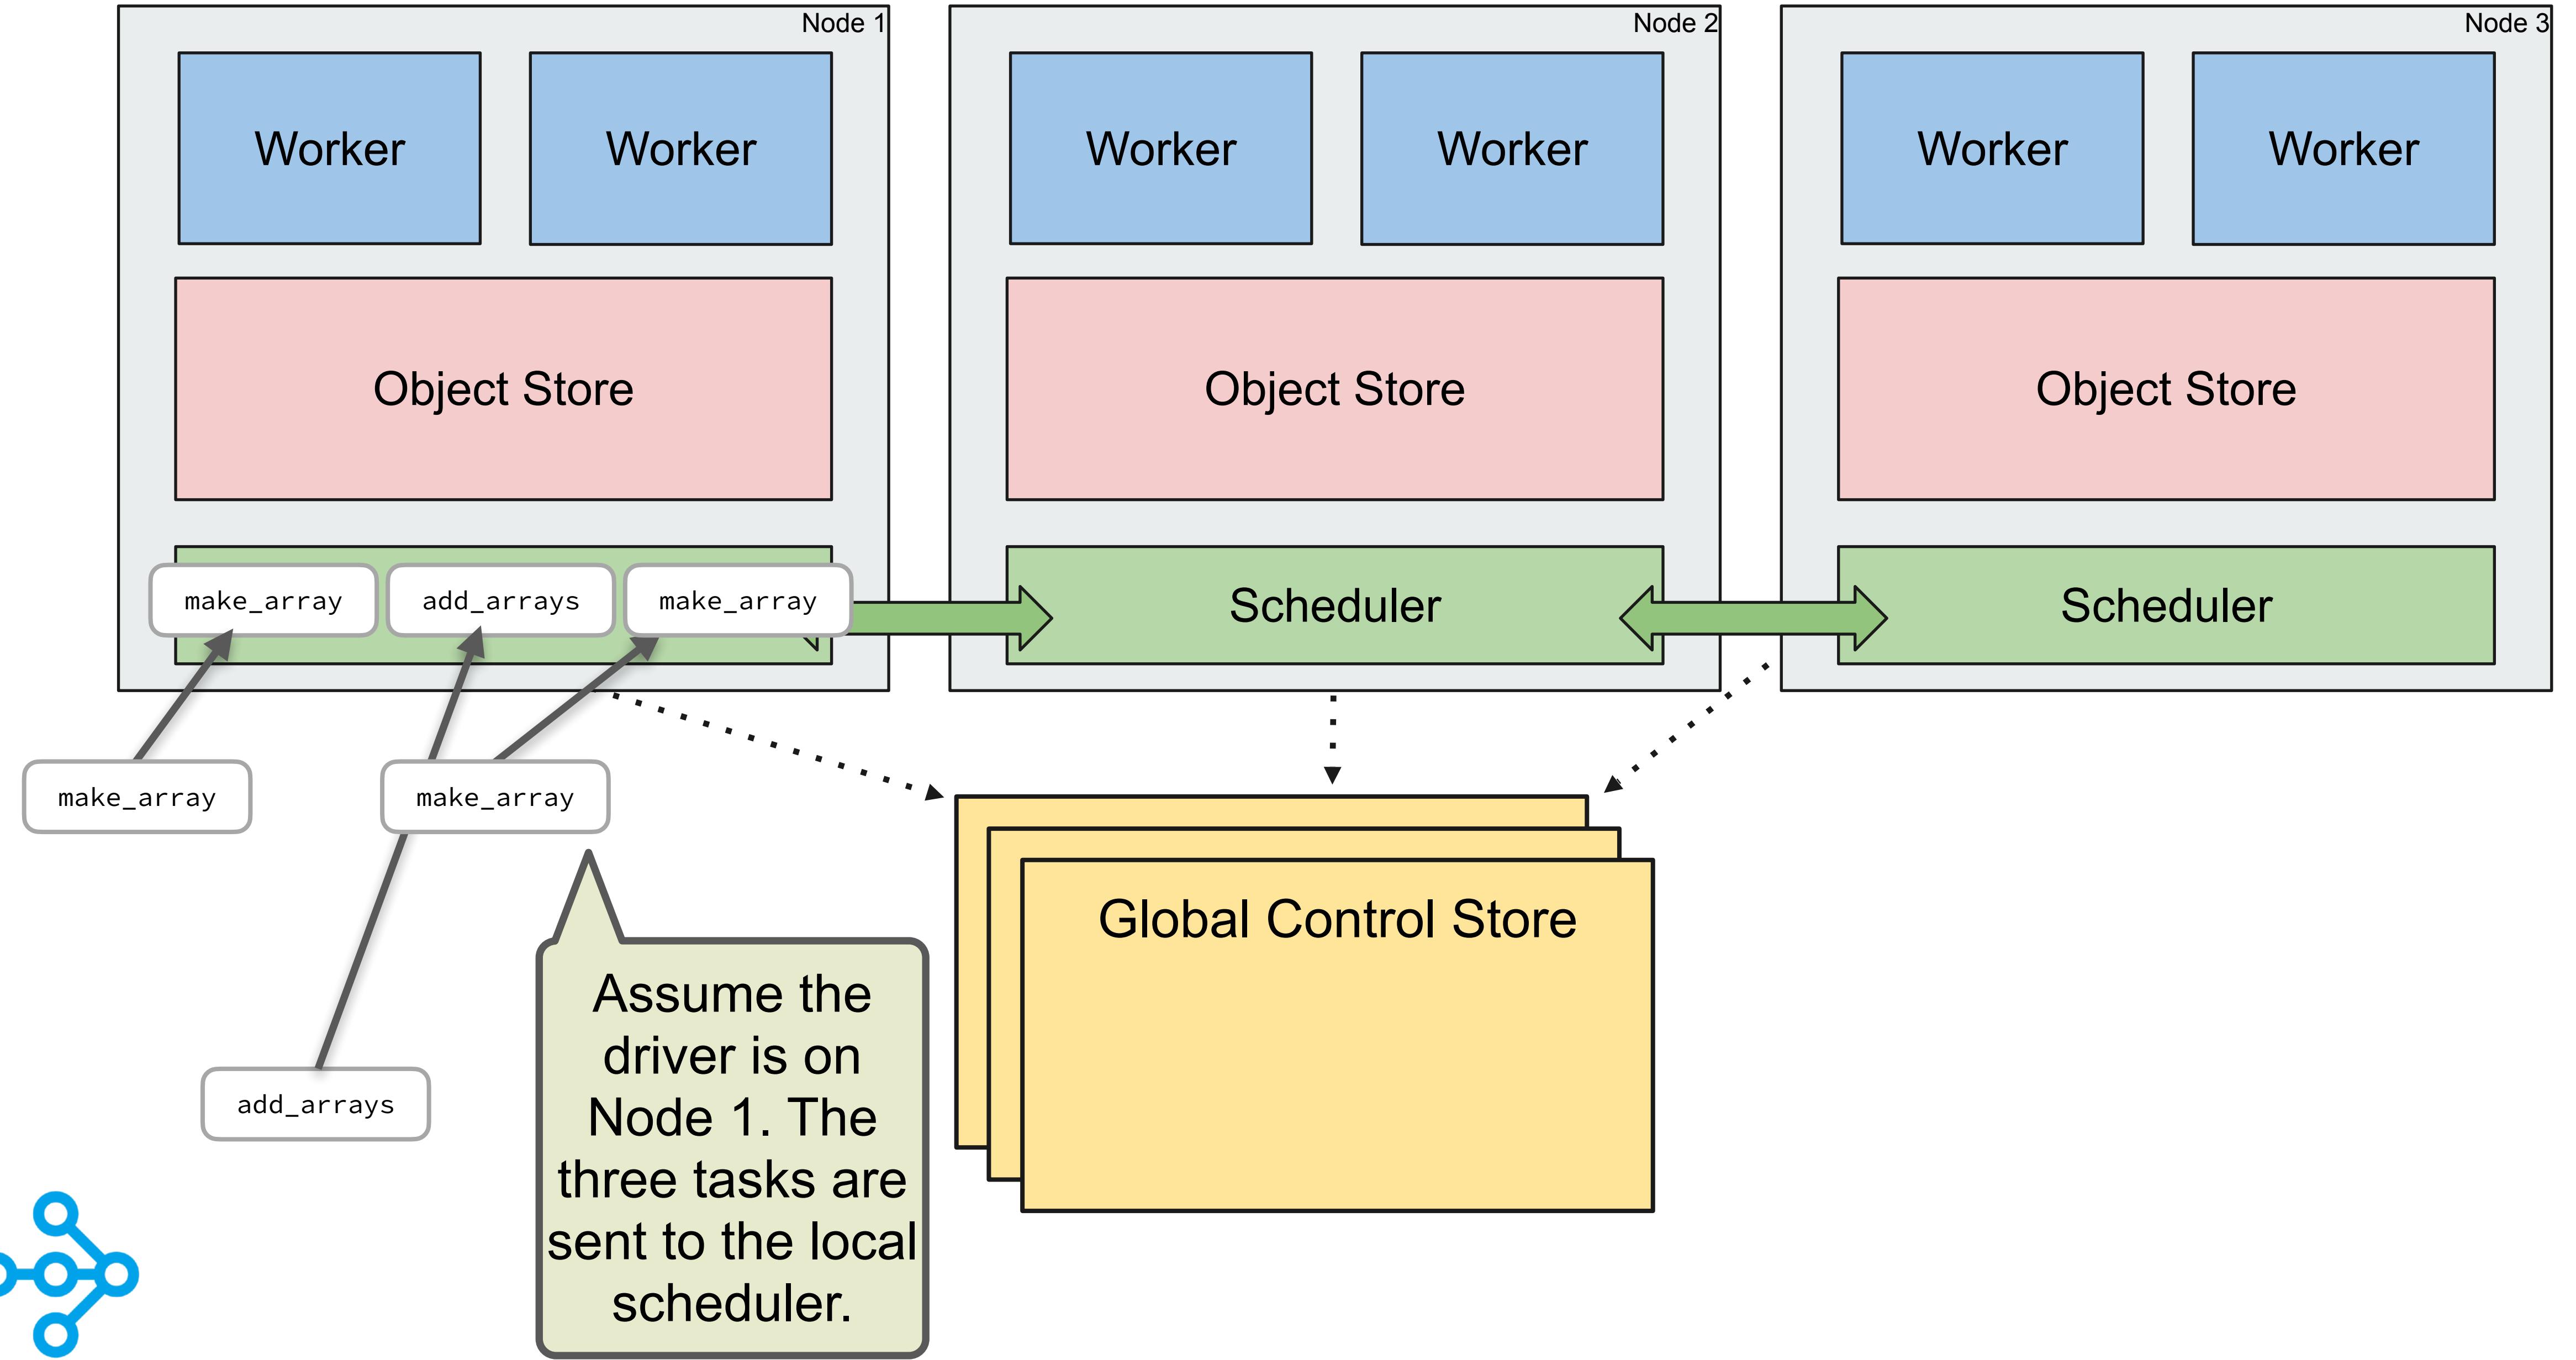
c = Counter.remote()
id4 = c.increment.remote()
id5 = c.increment.remote()
ray.get([id4, id5]) # [1, 2]
```

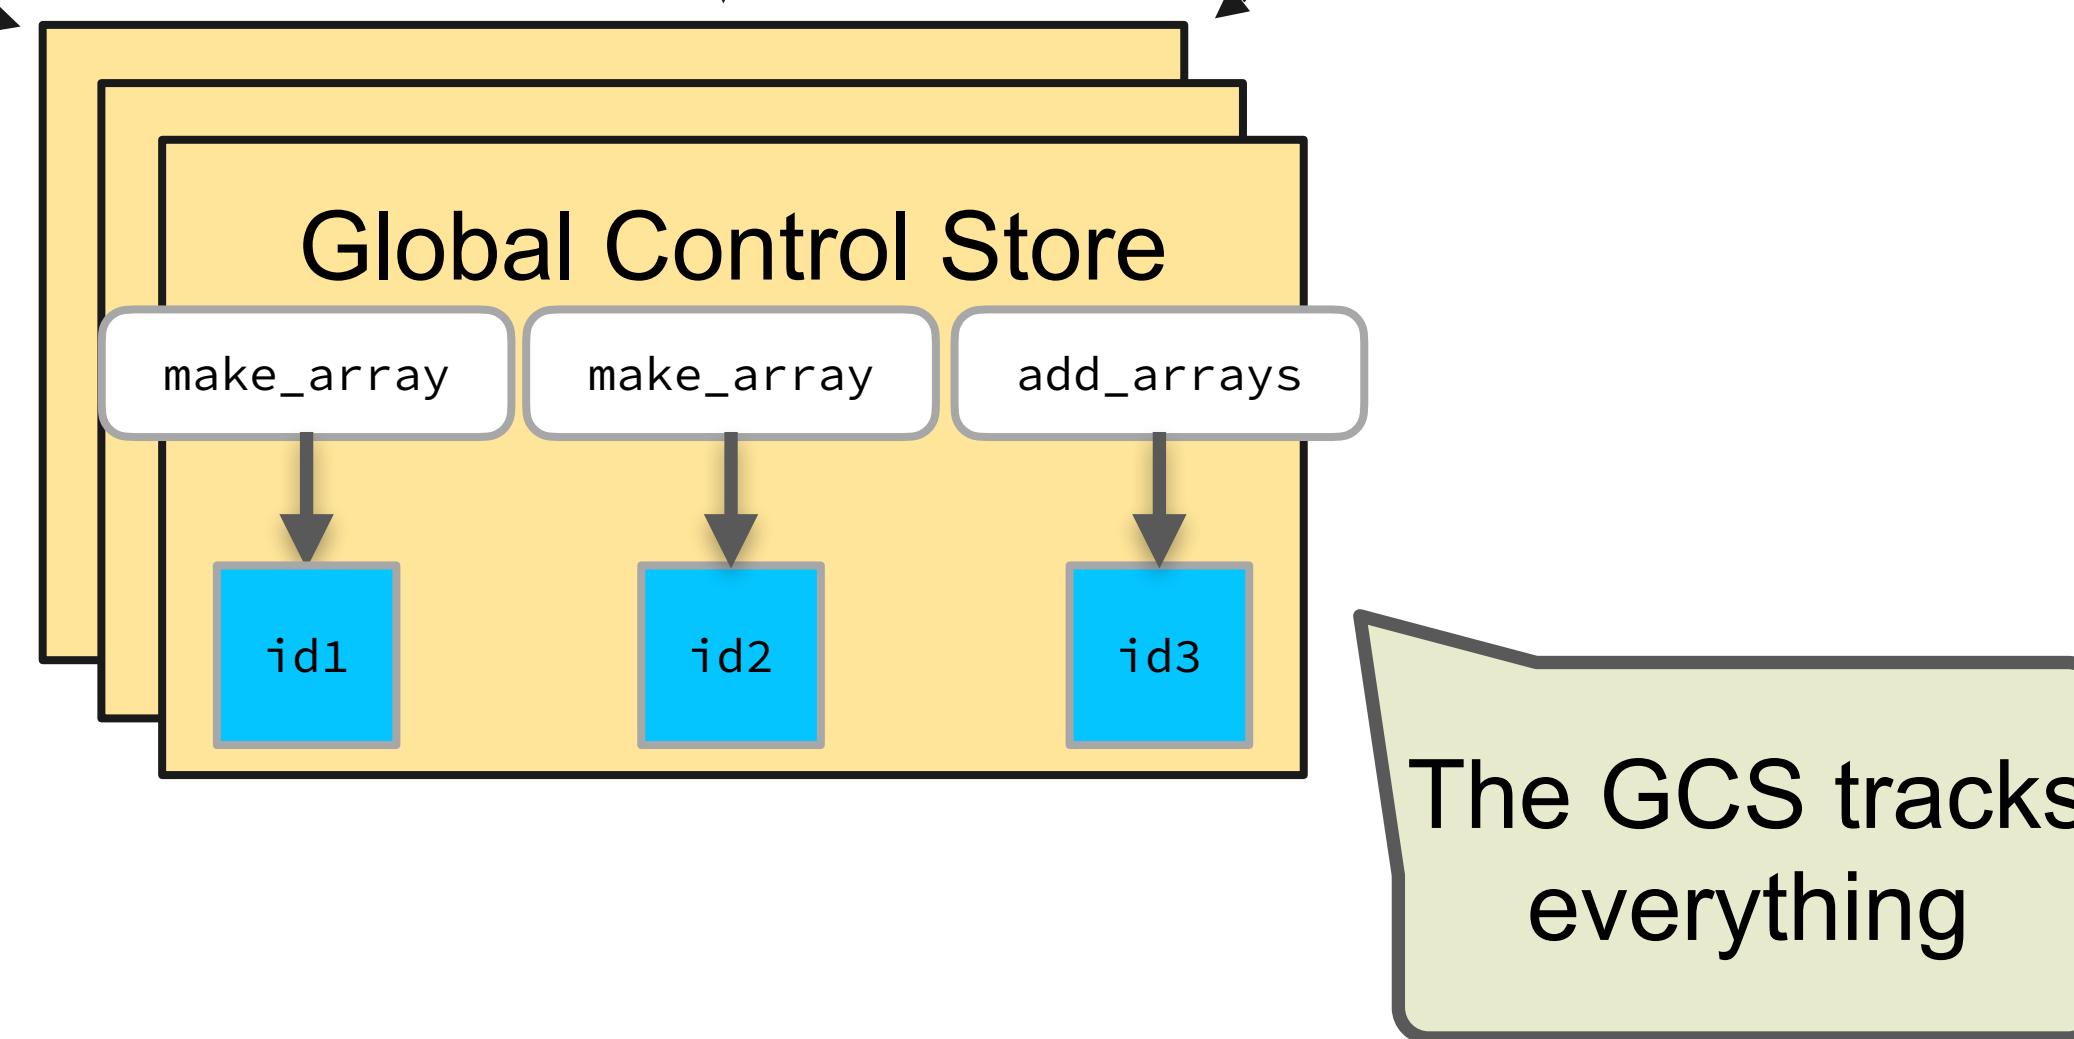
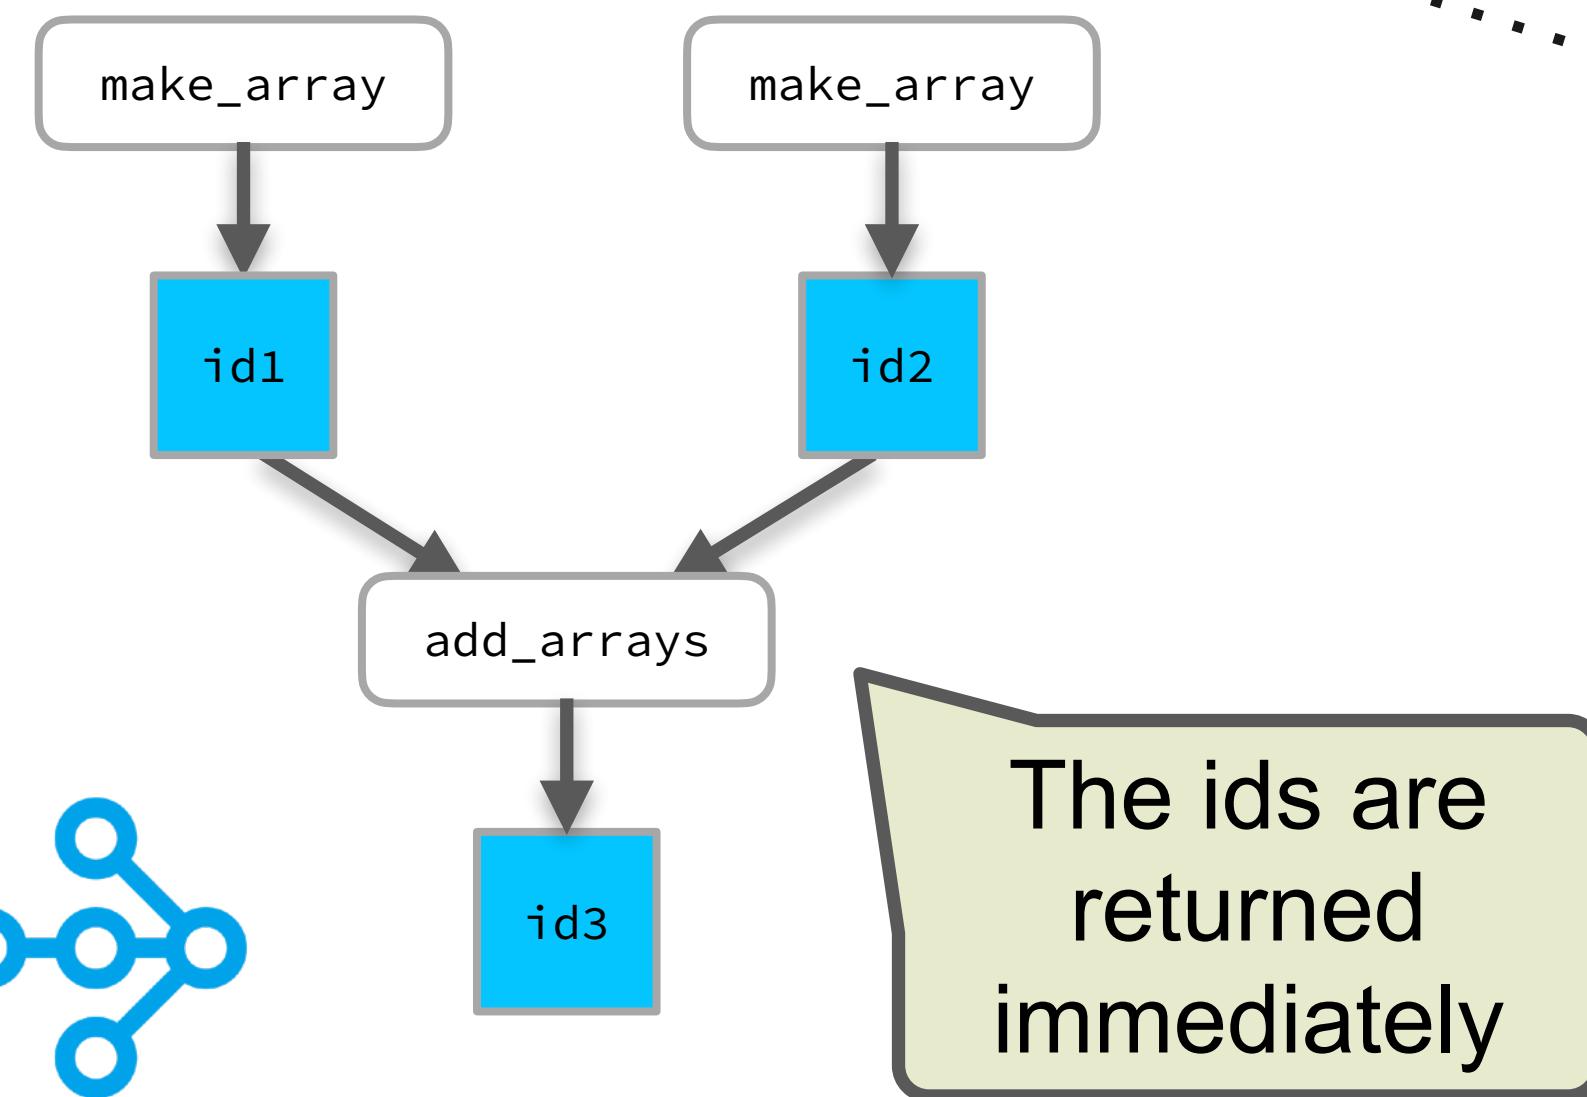
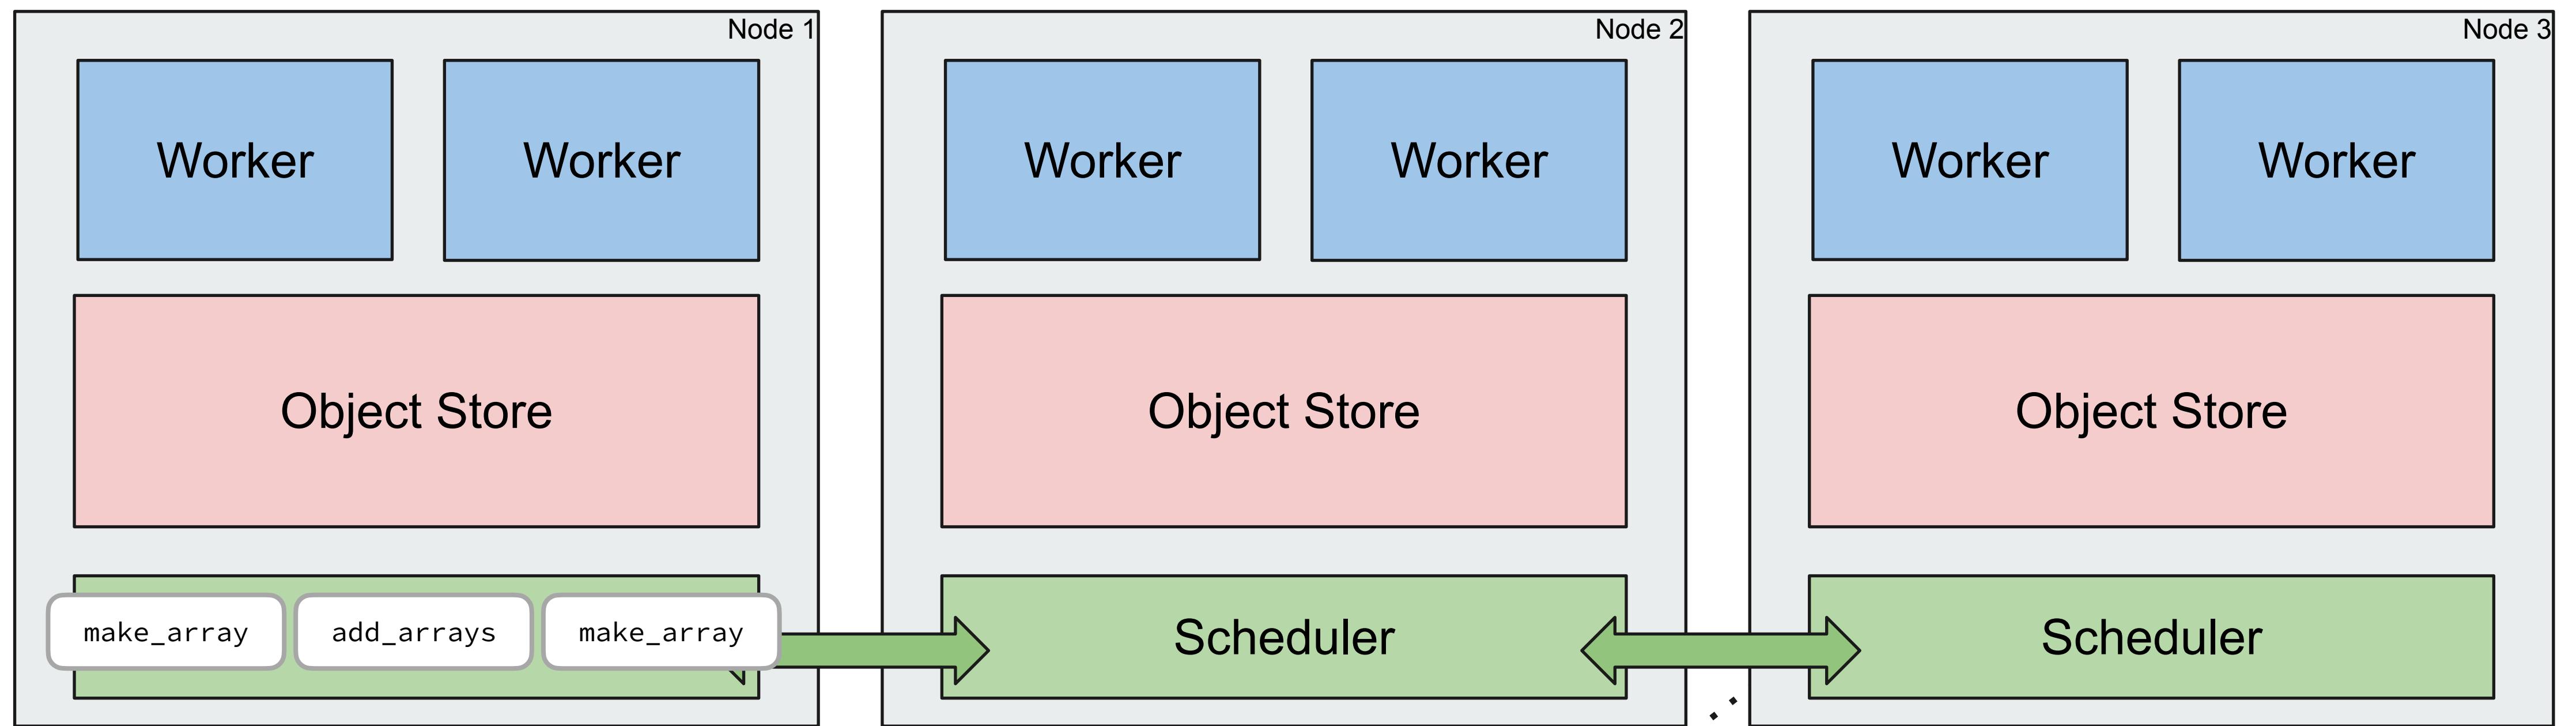


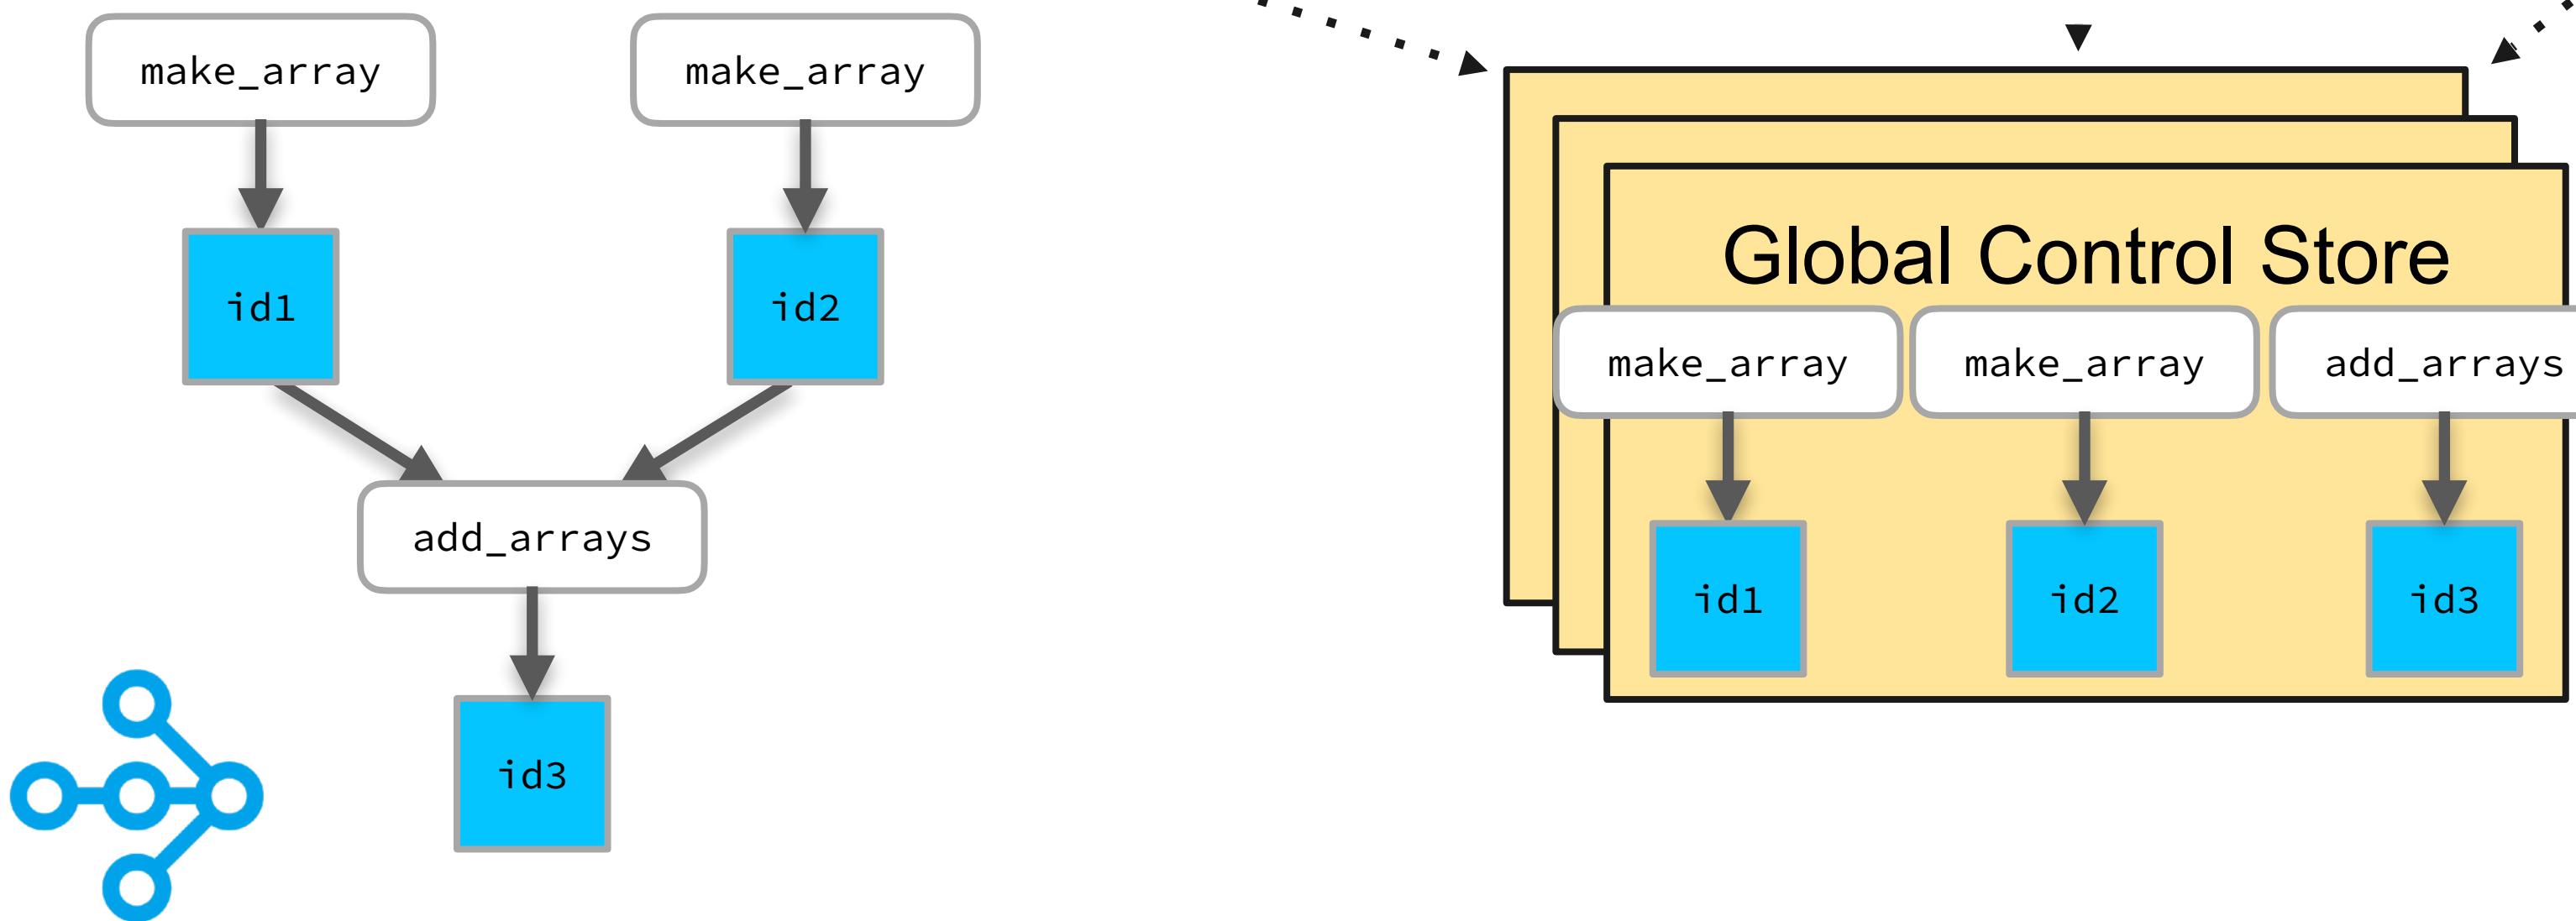
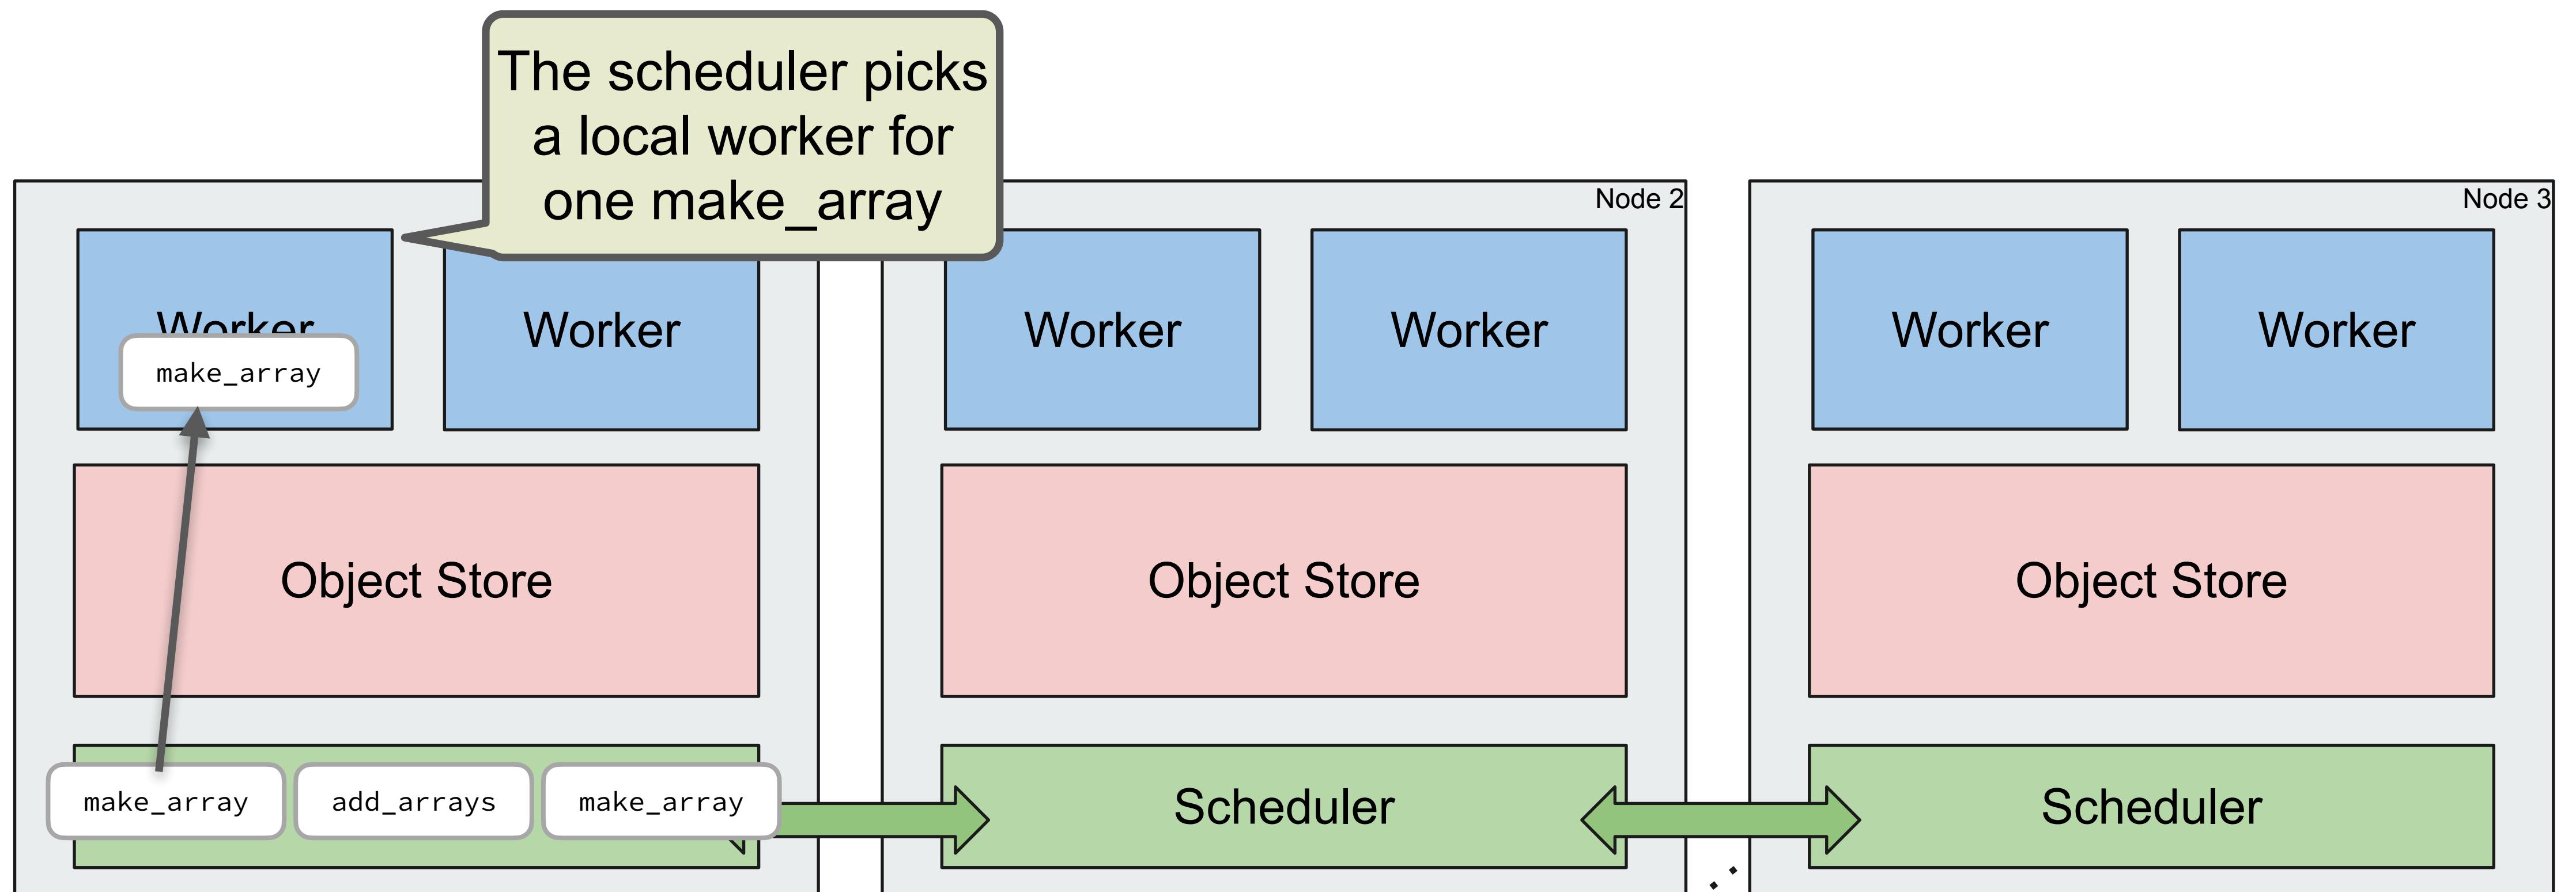
How does this work behind the scenes?

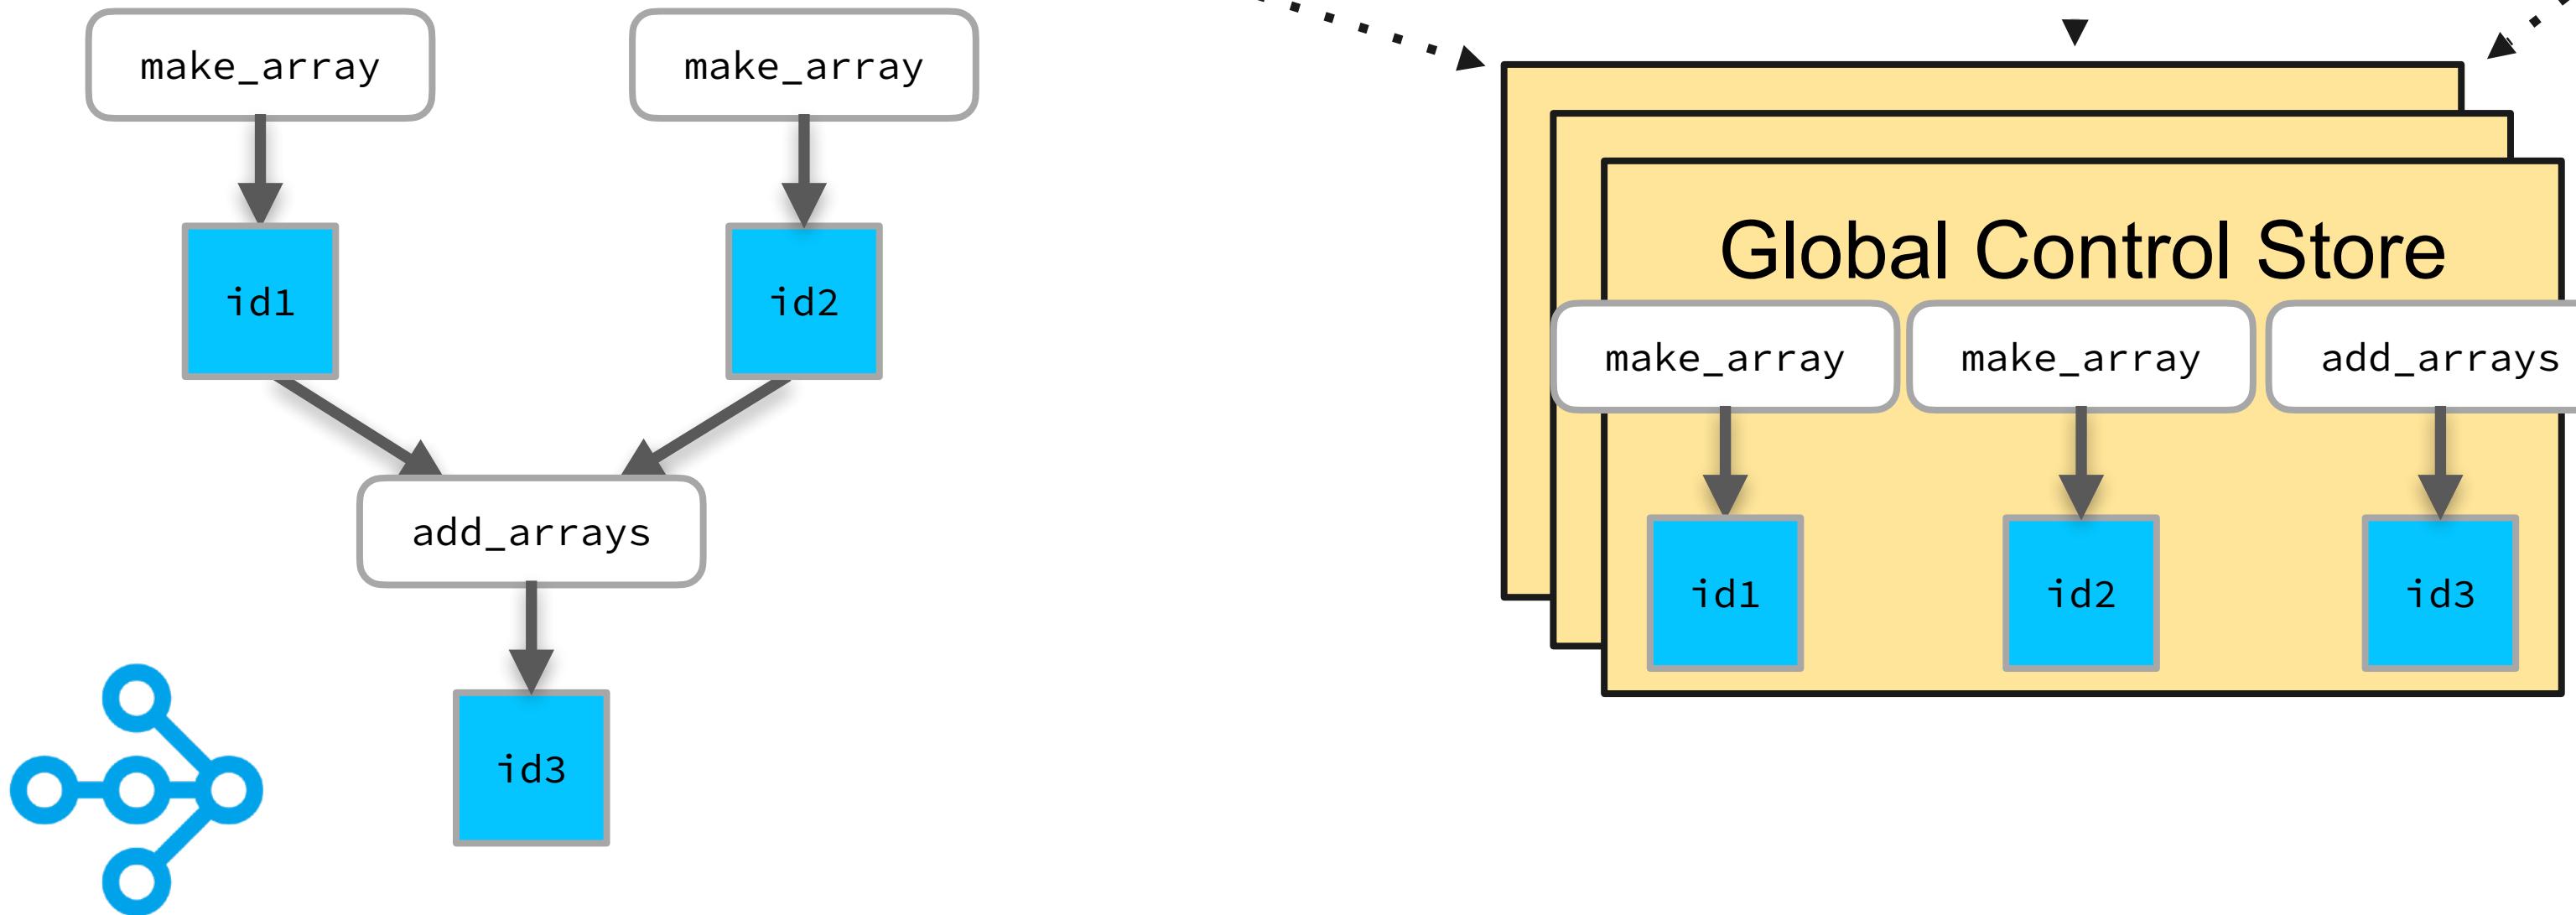
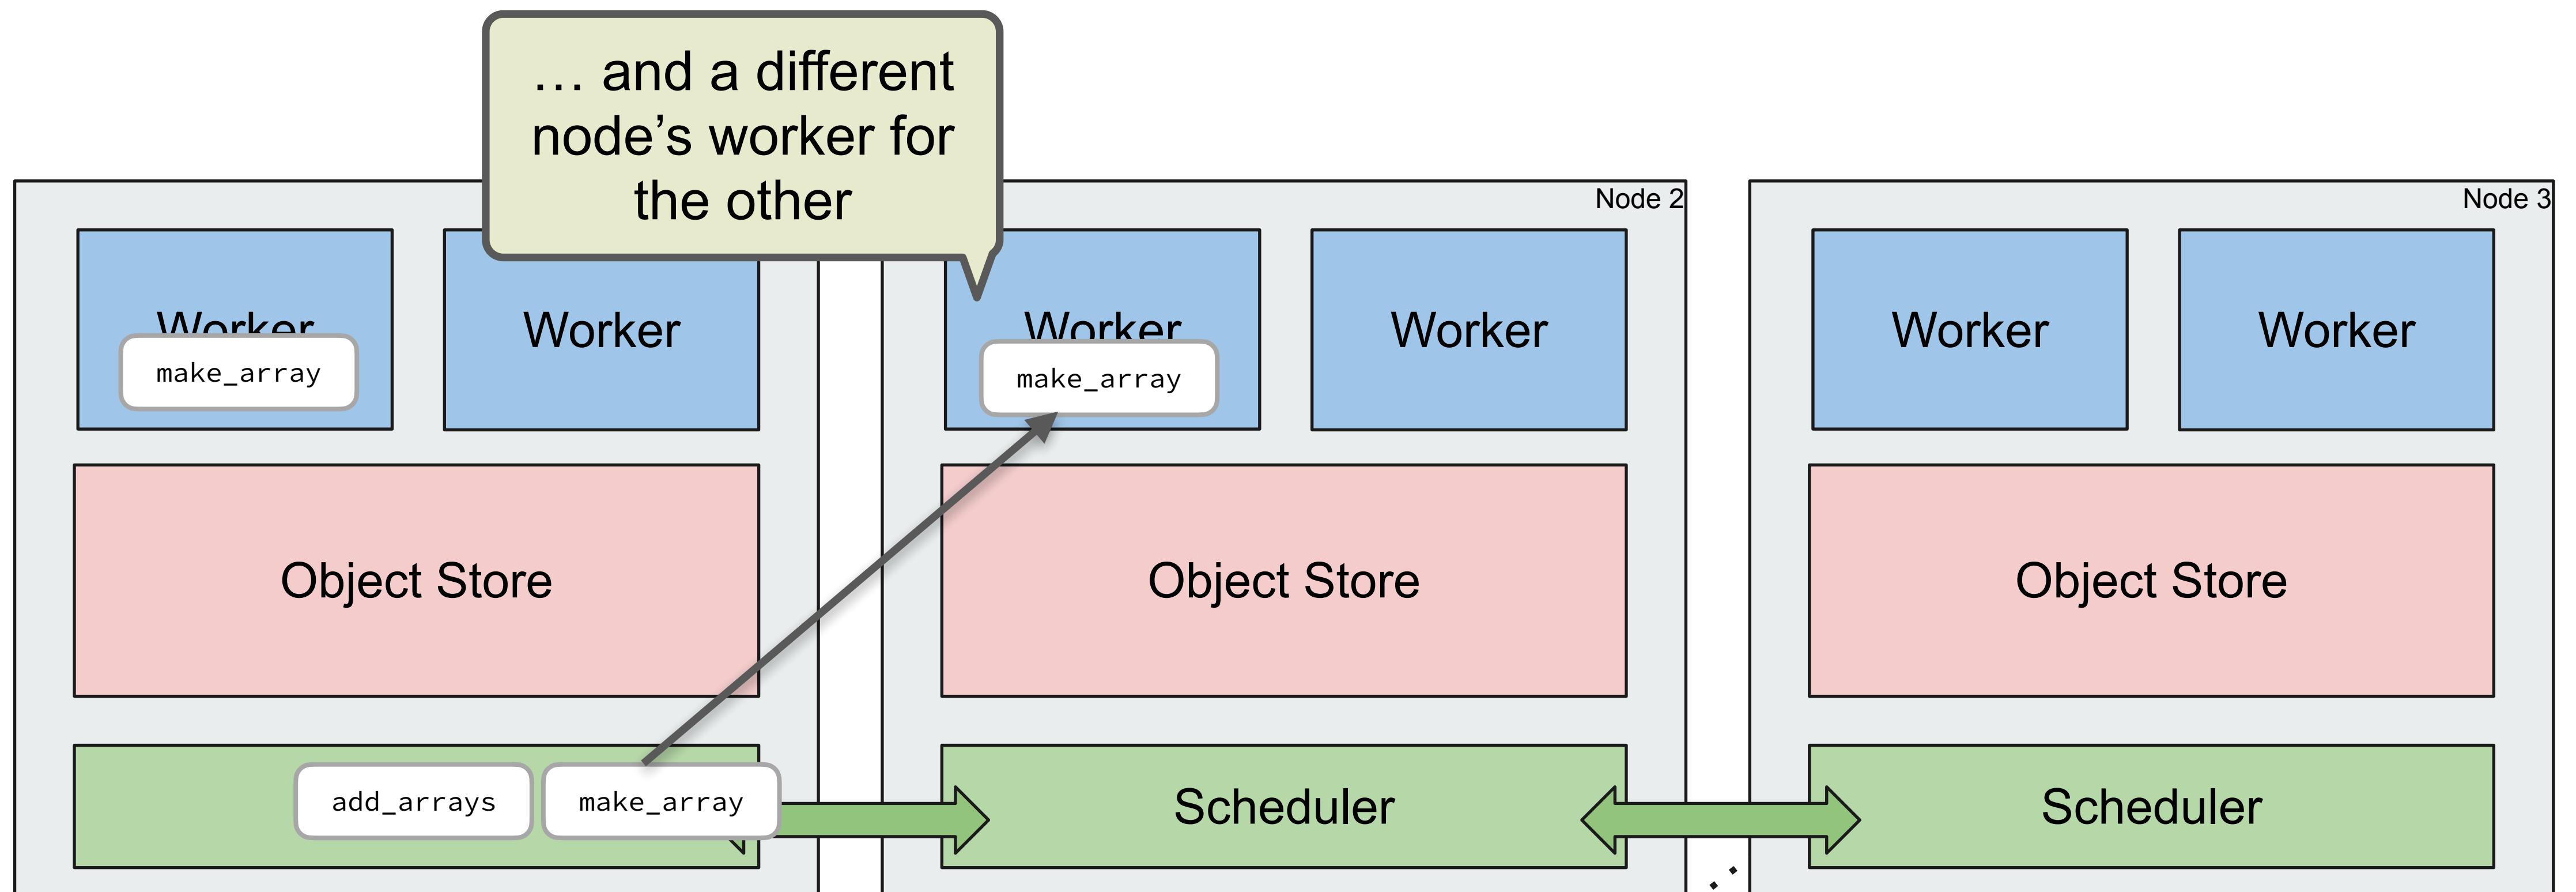


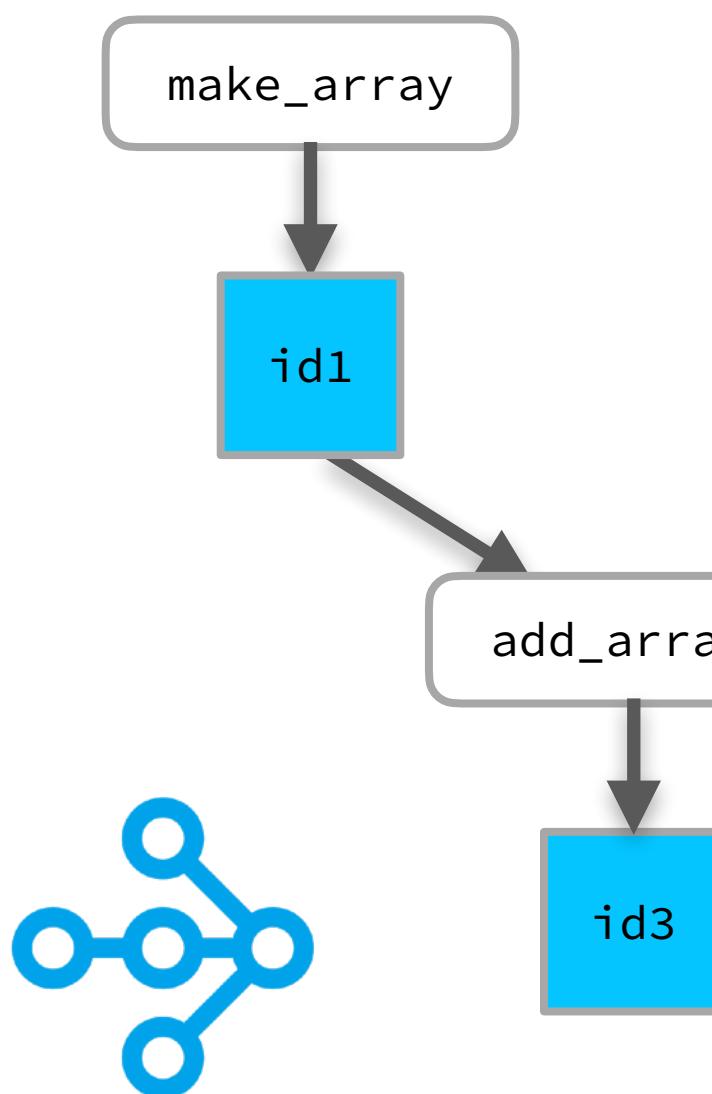
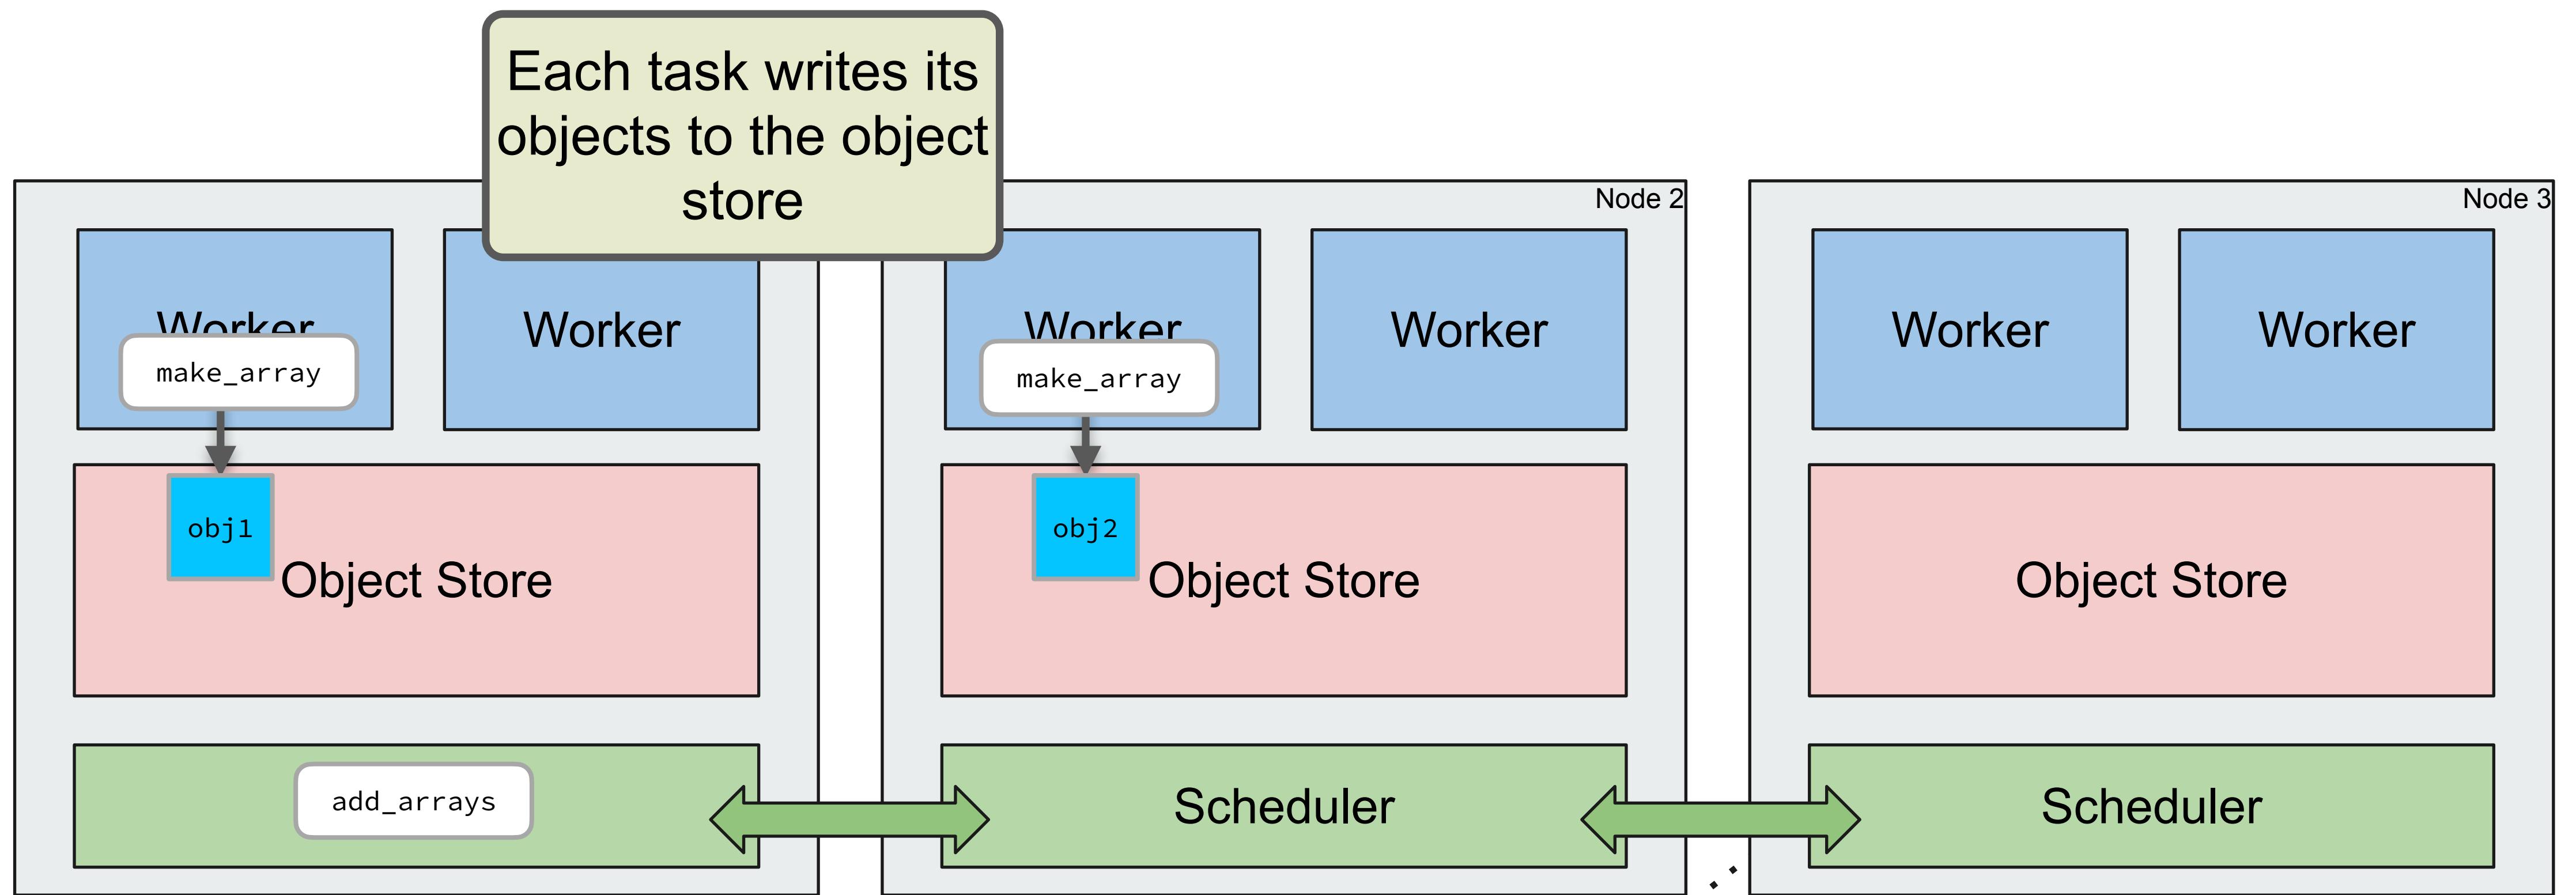


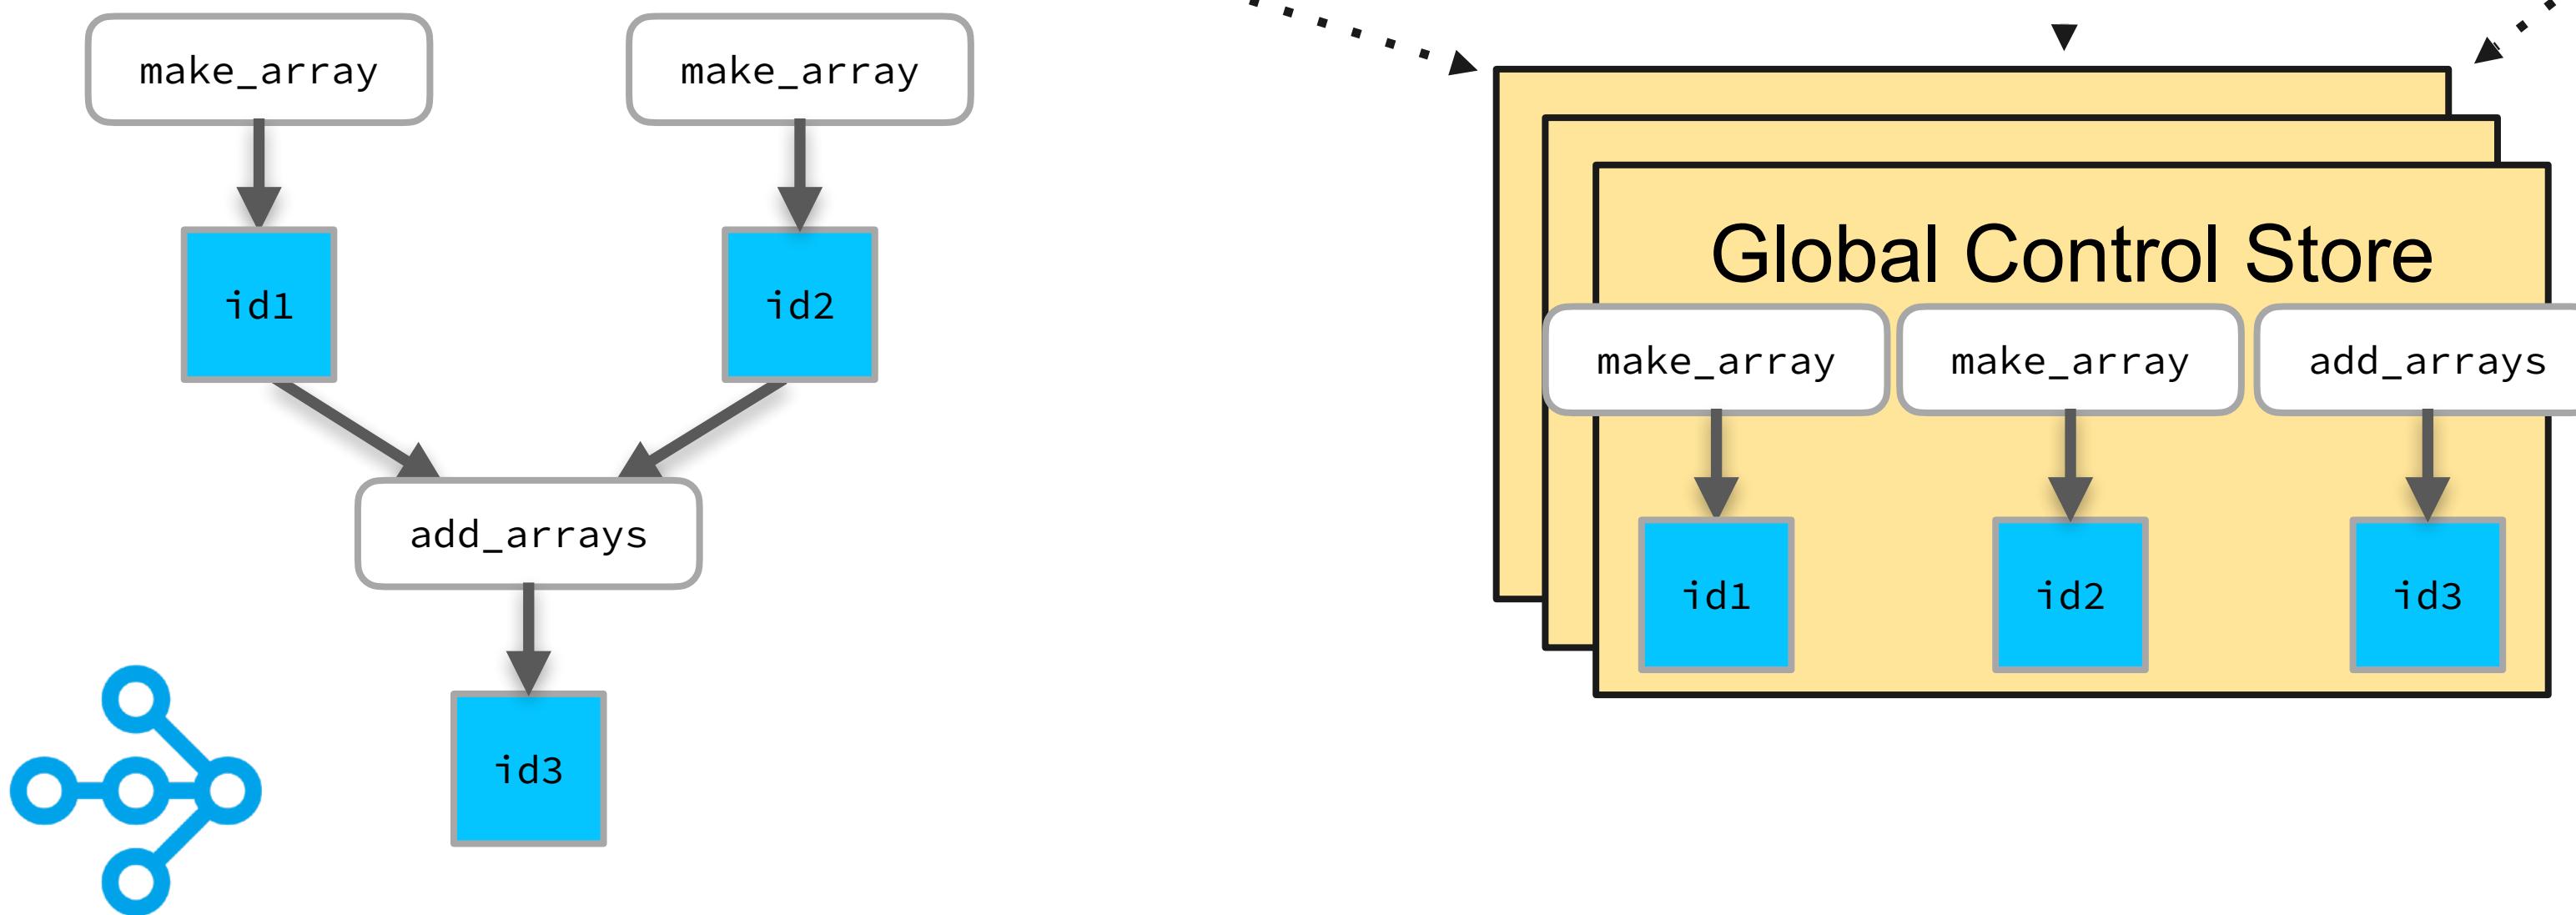
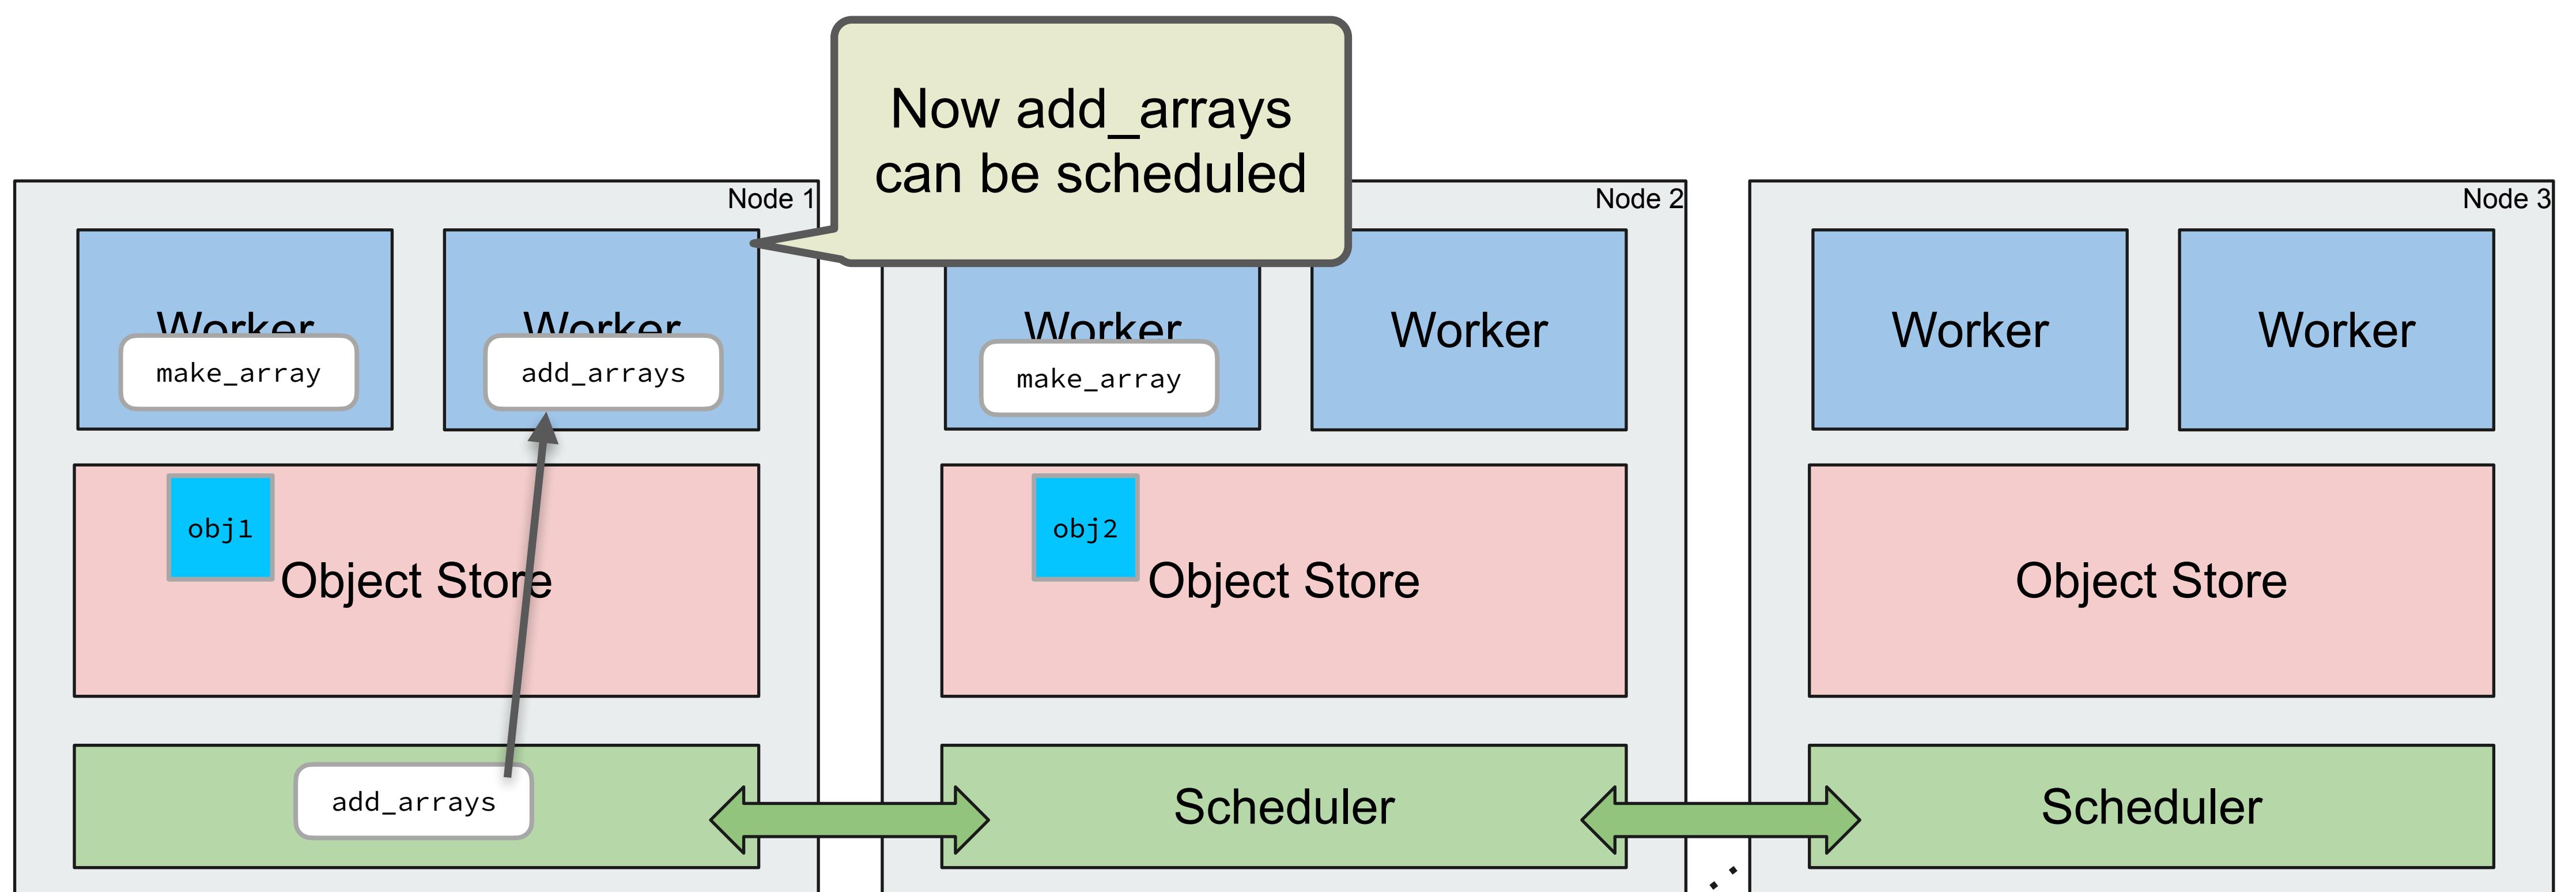






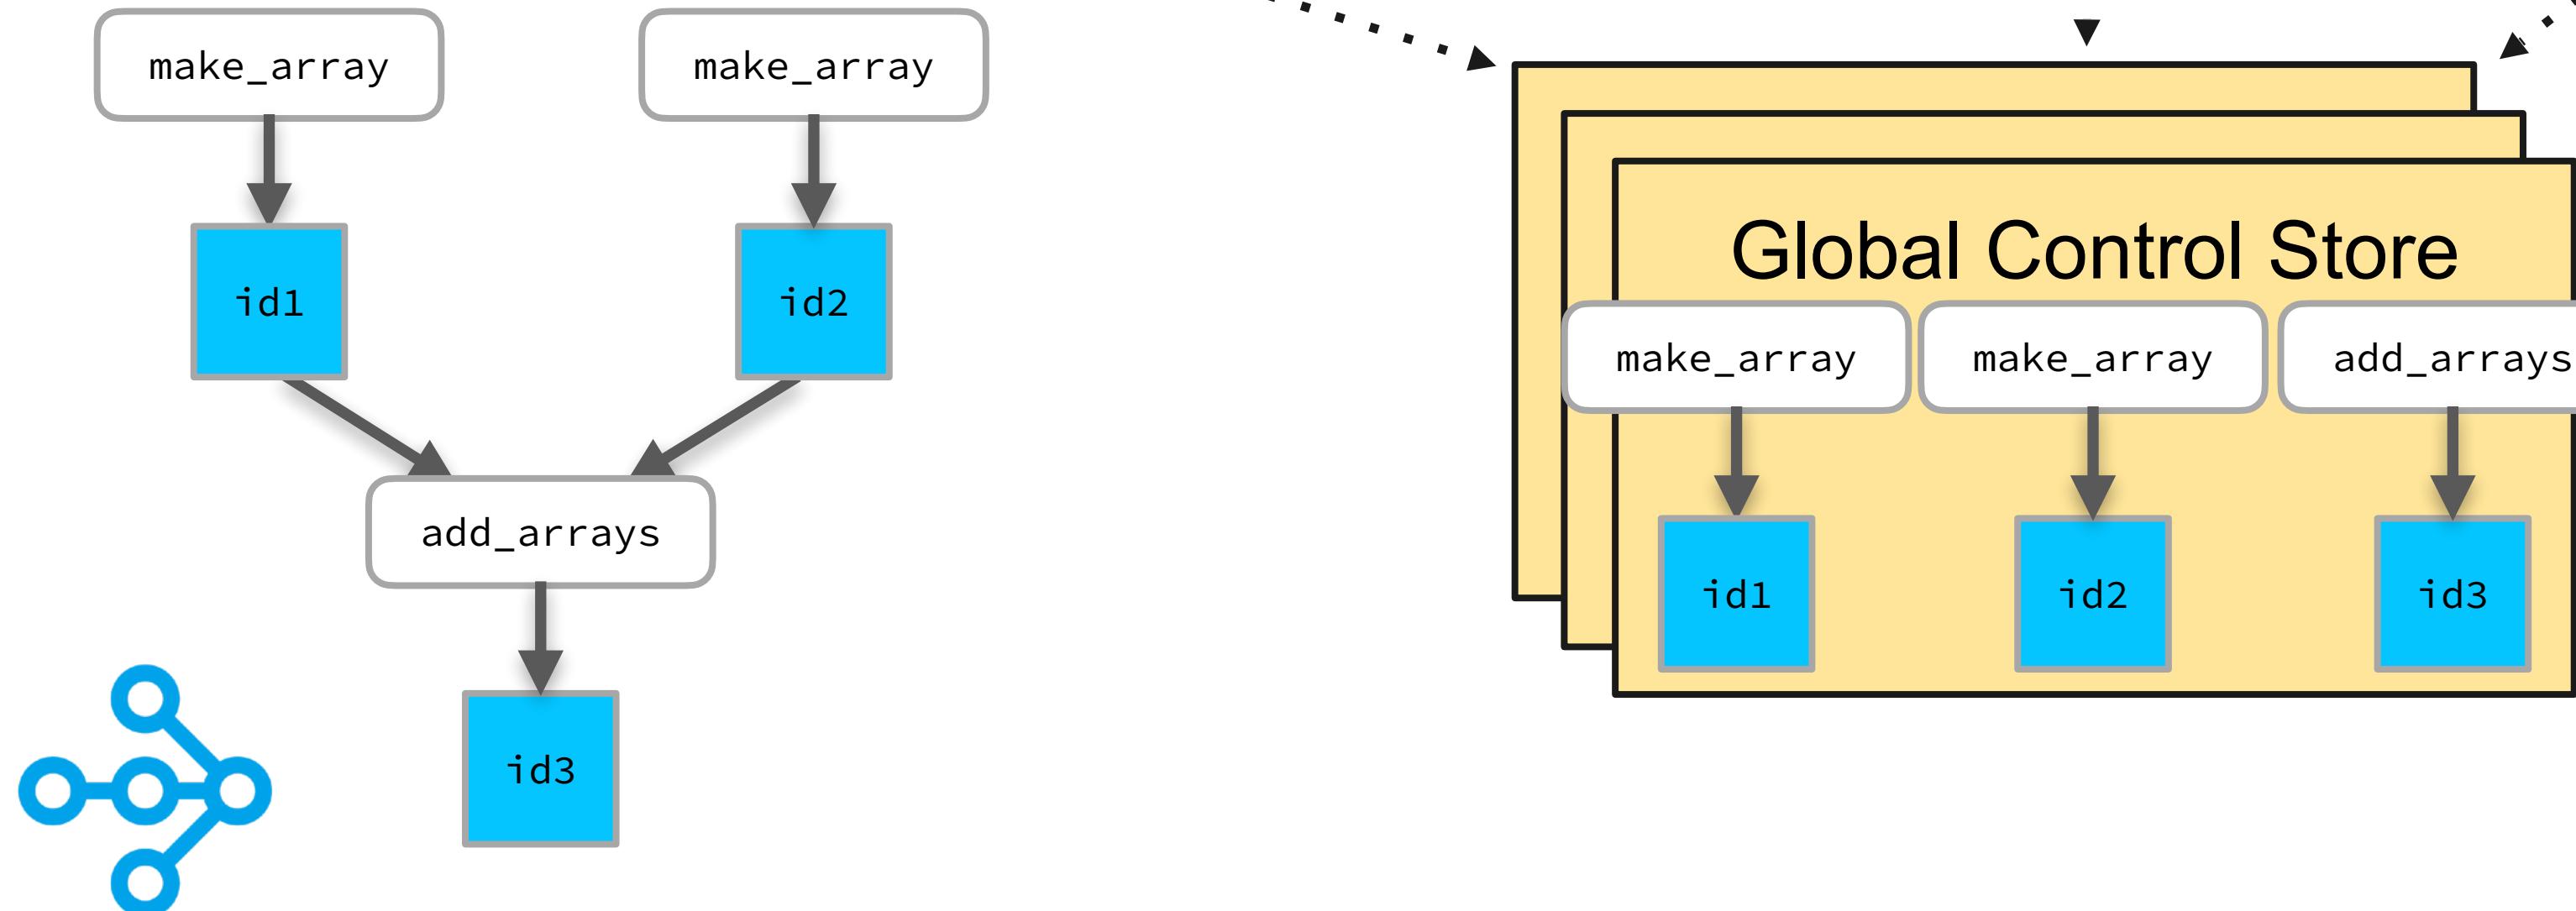
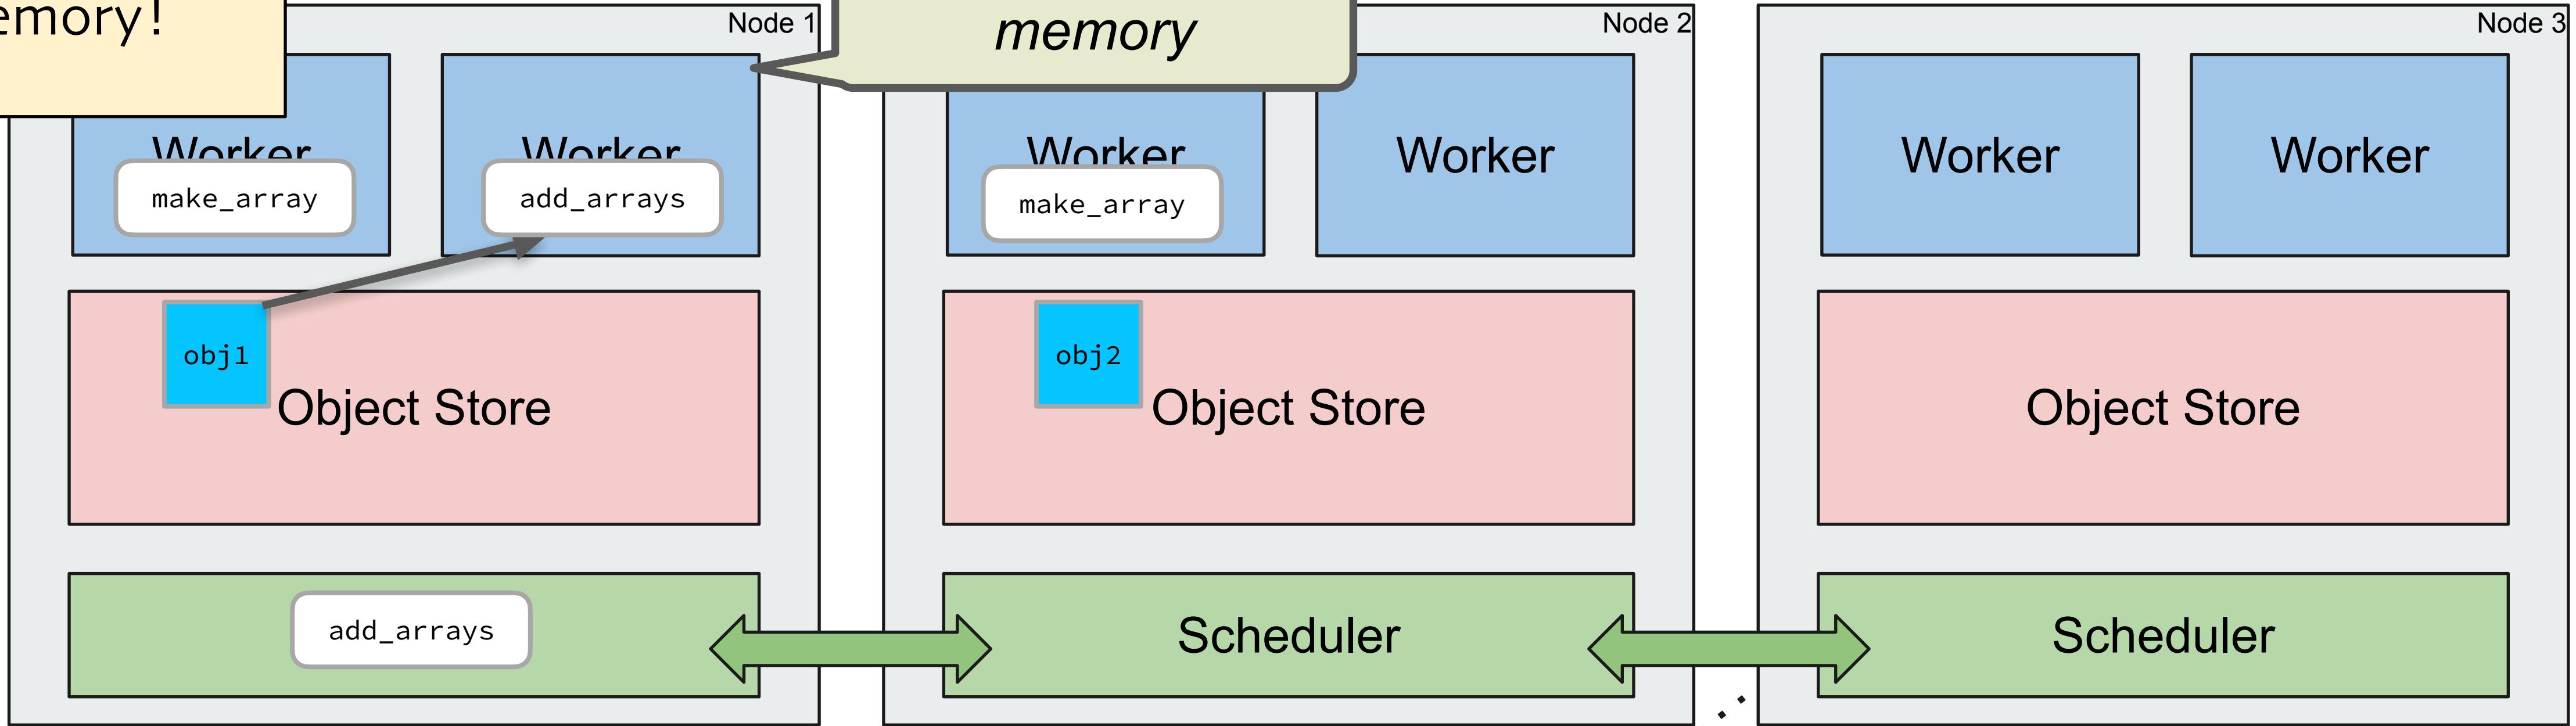


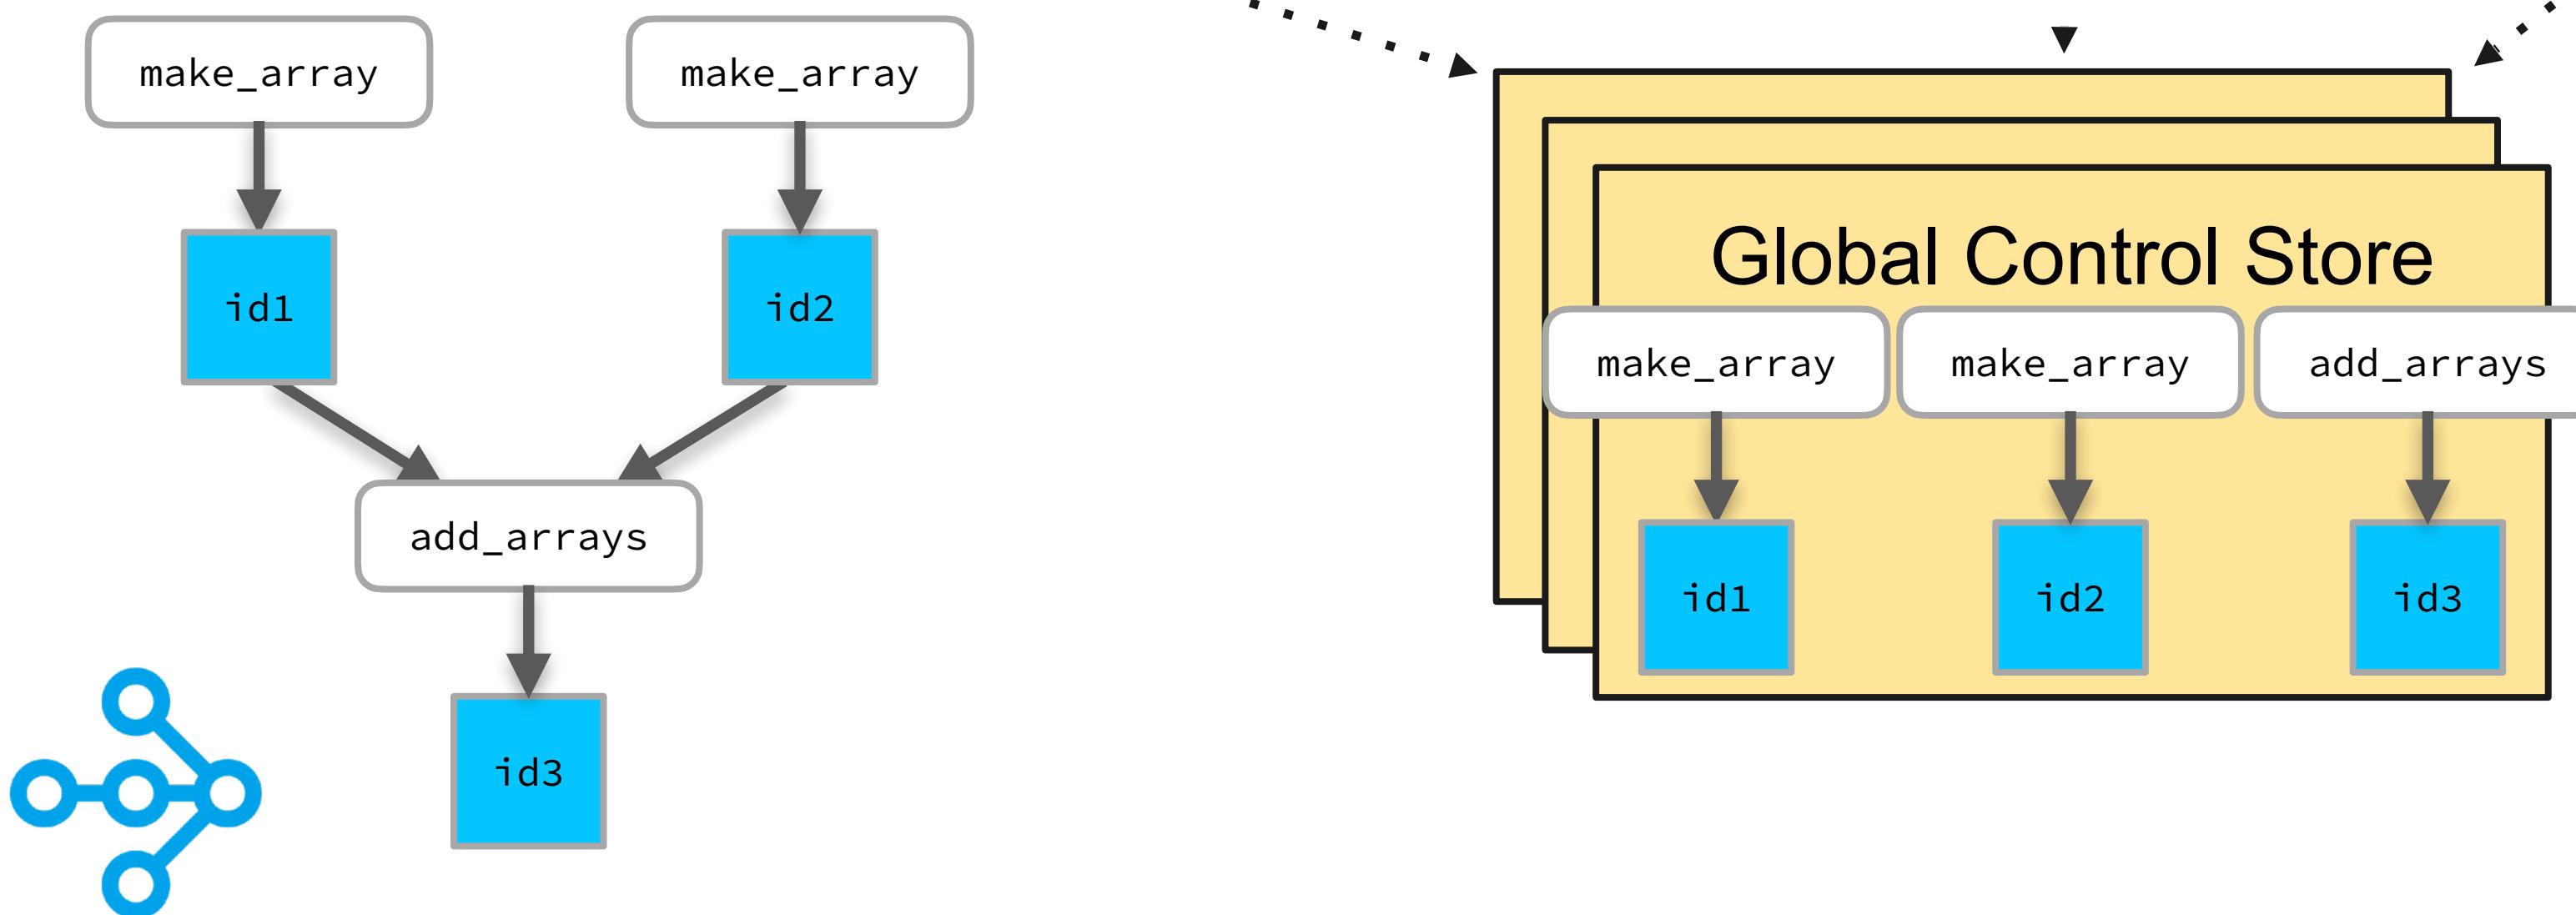
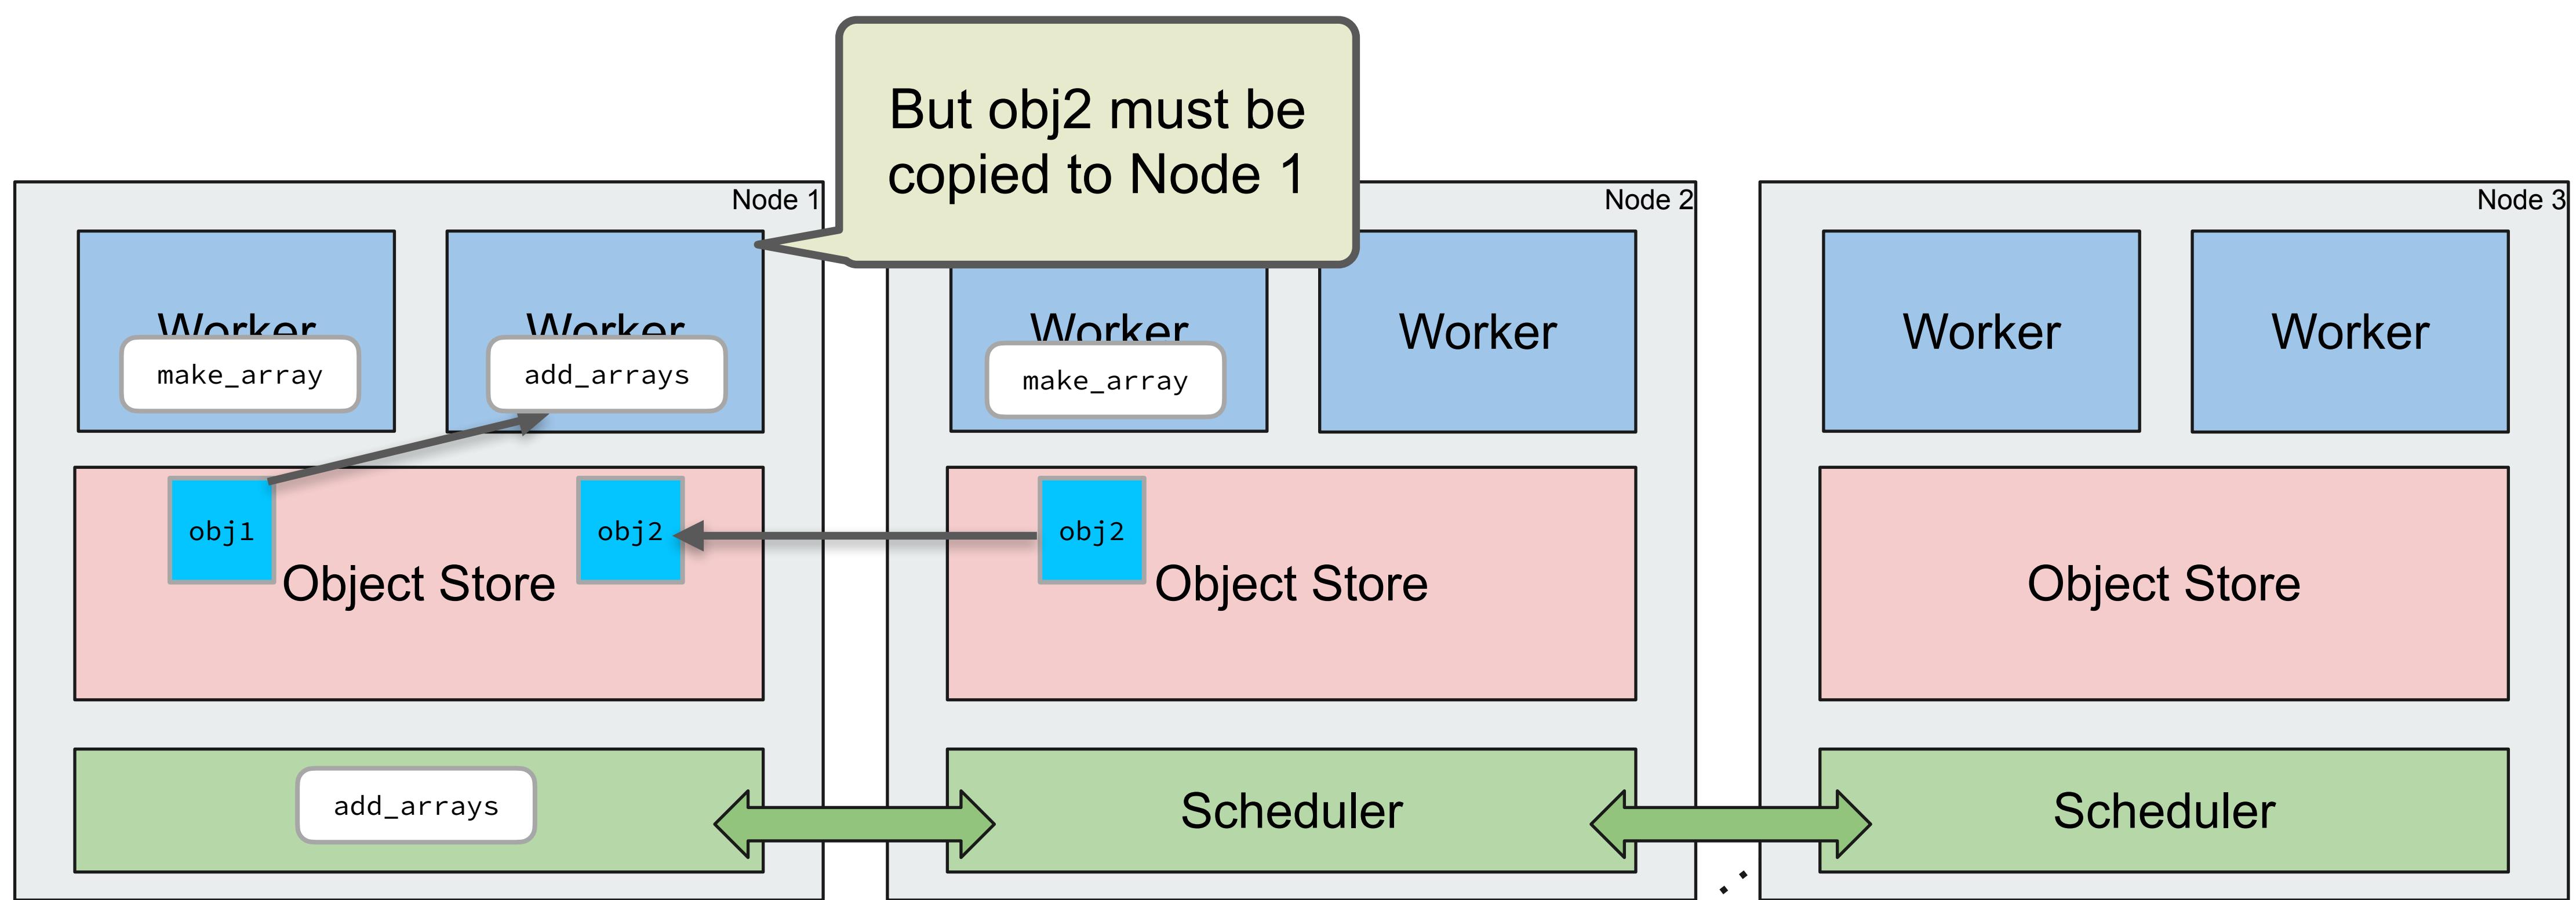


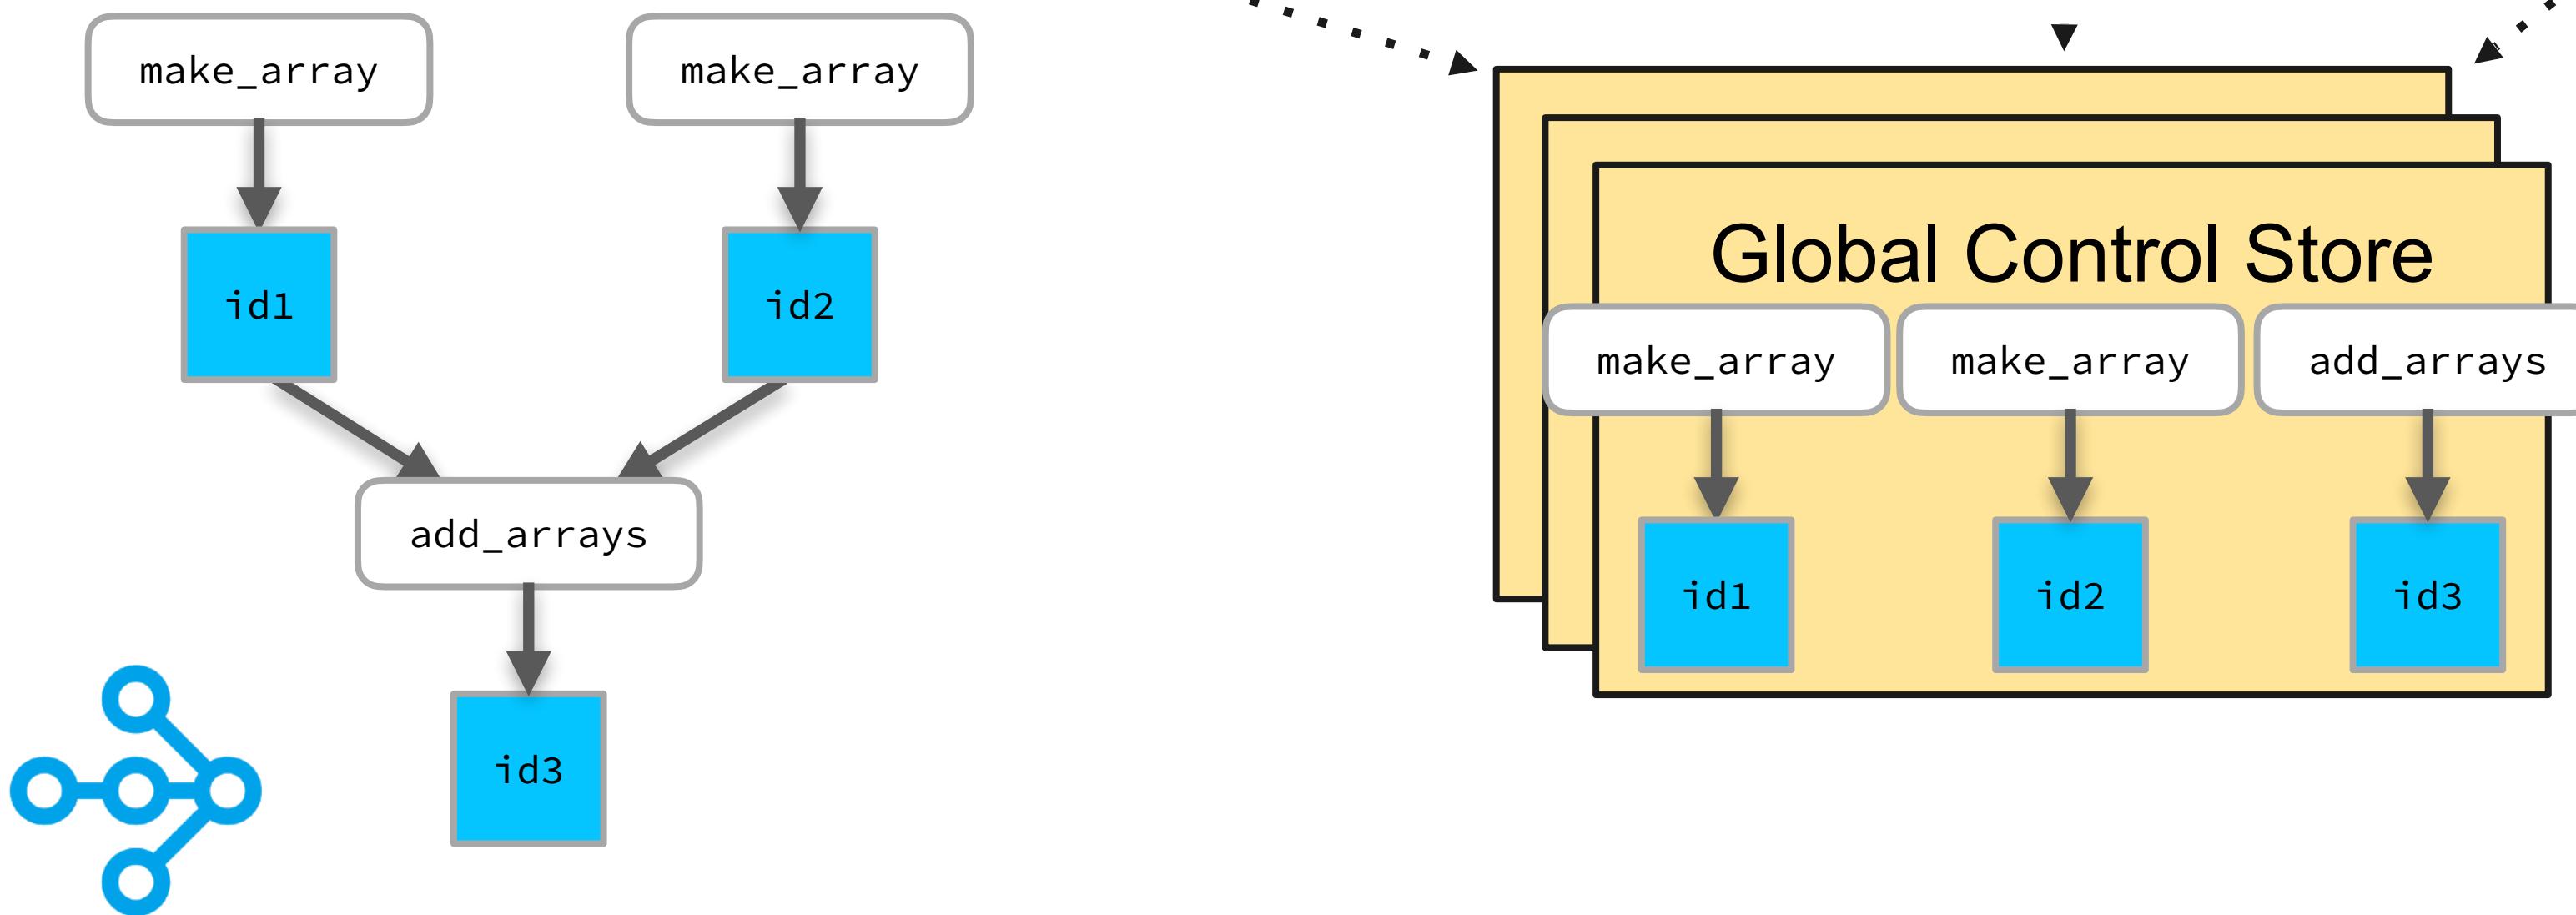
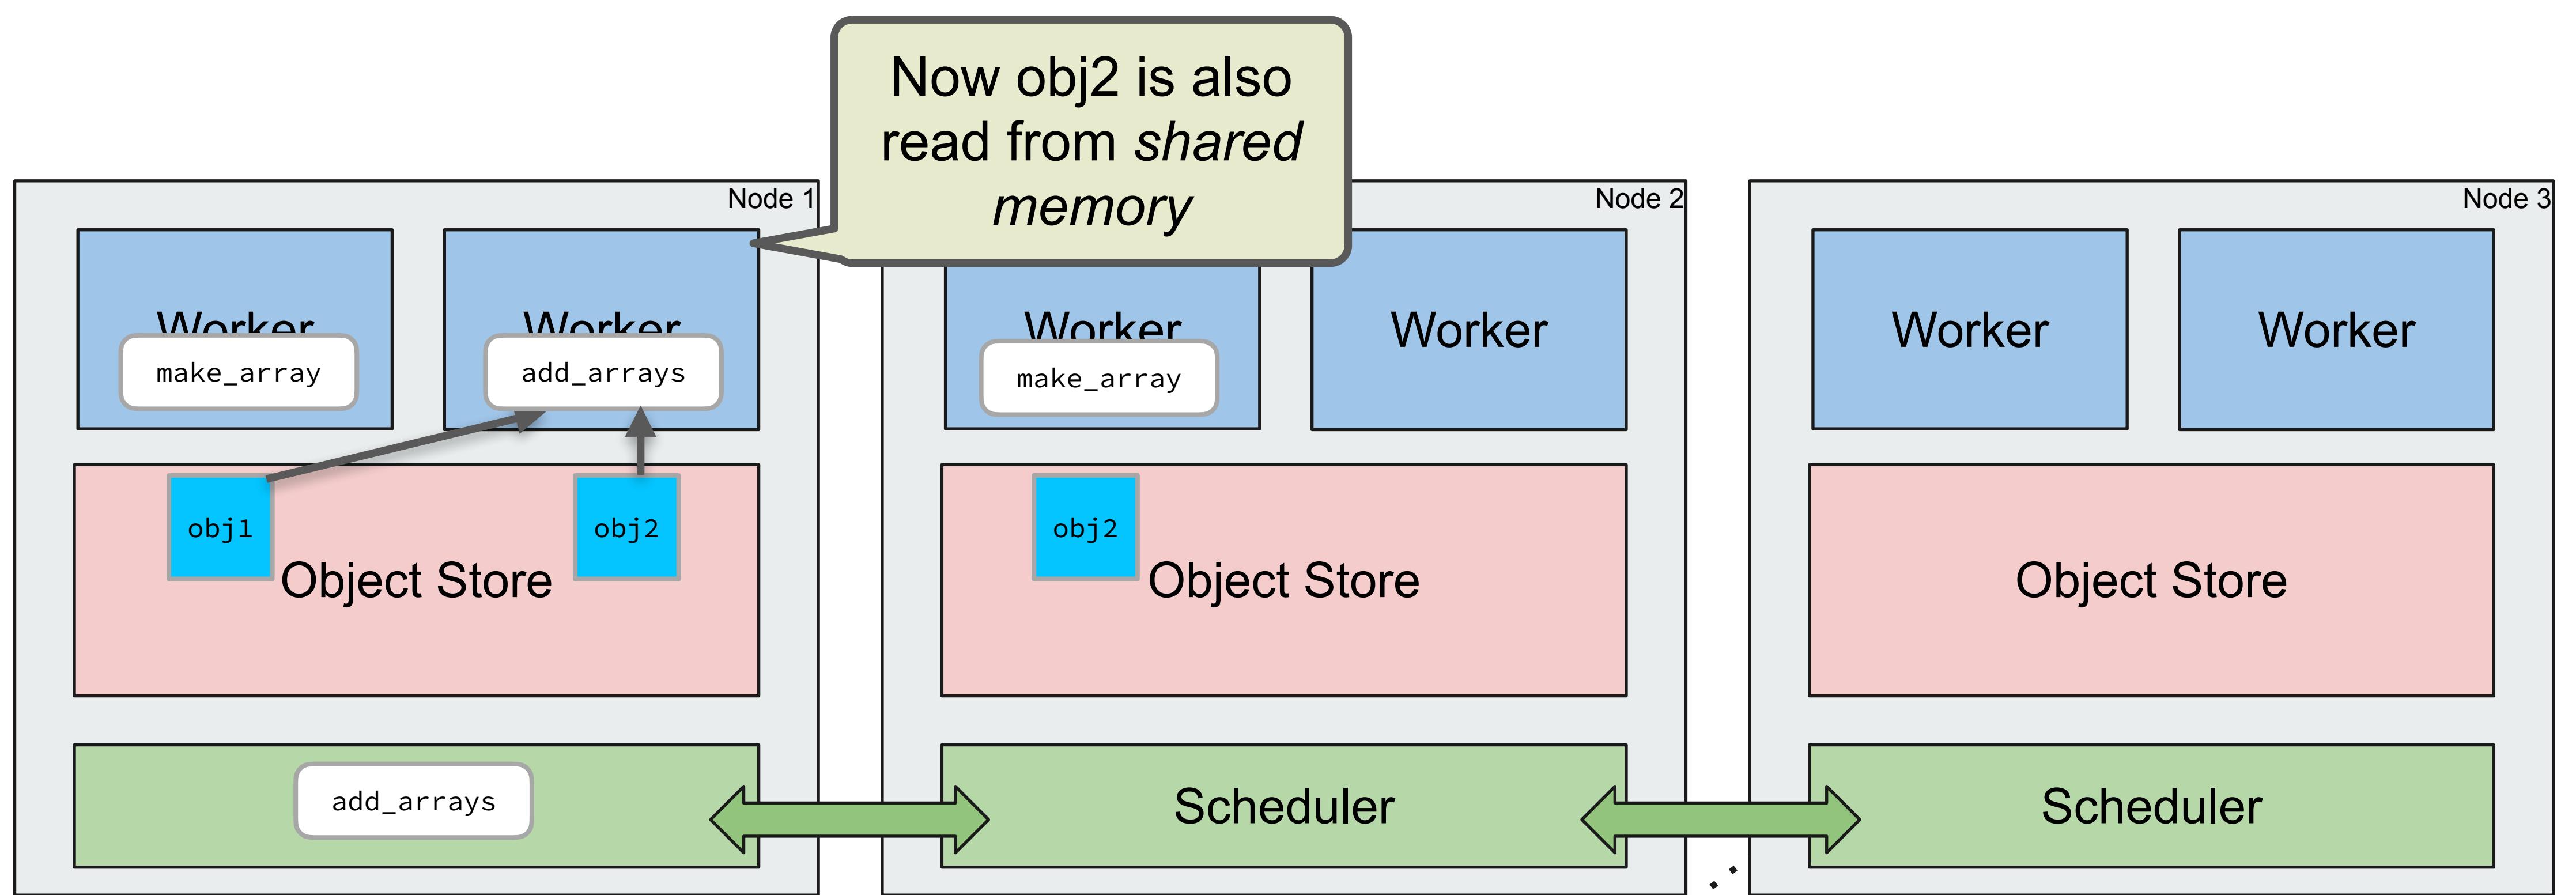


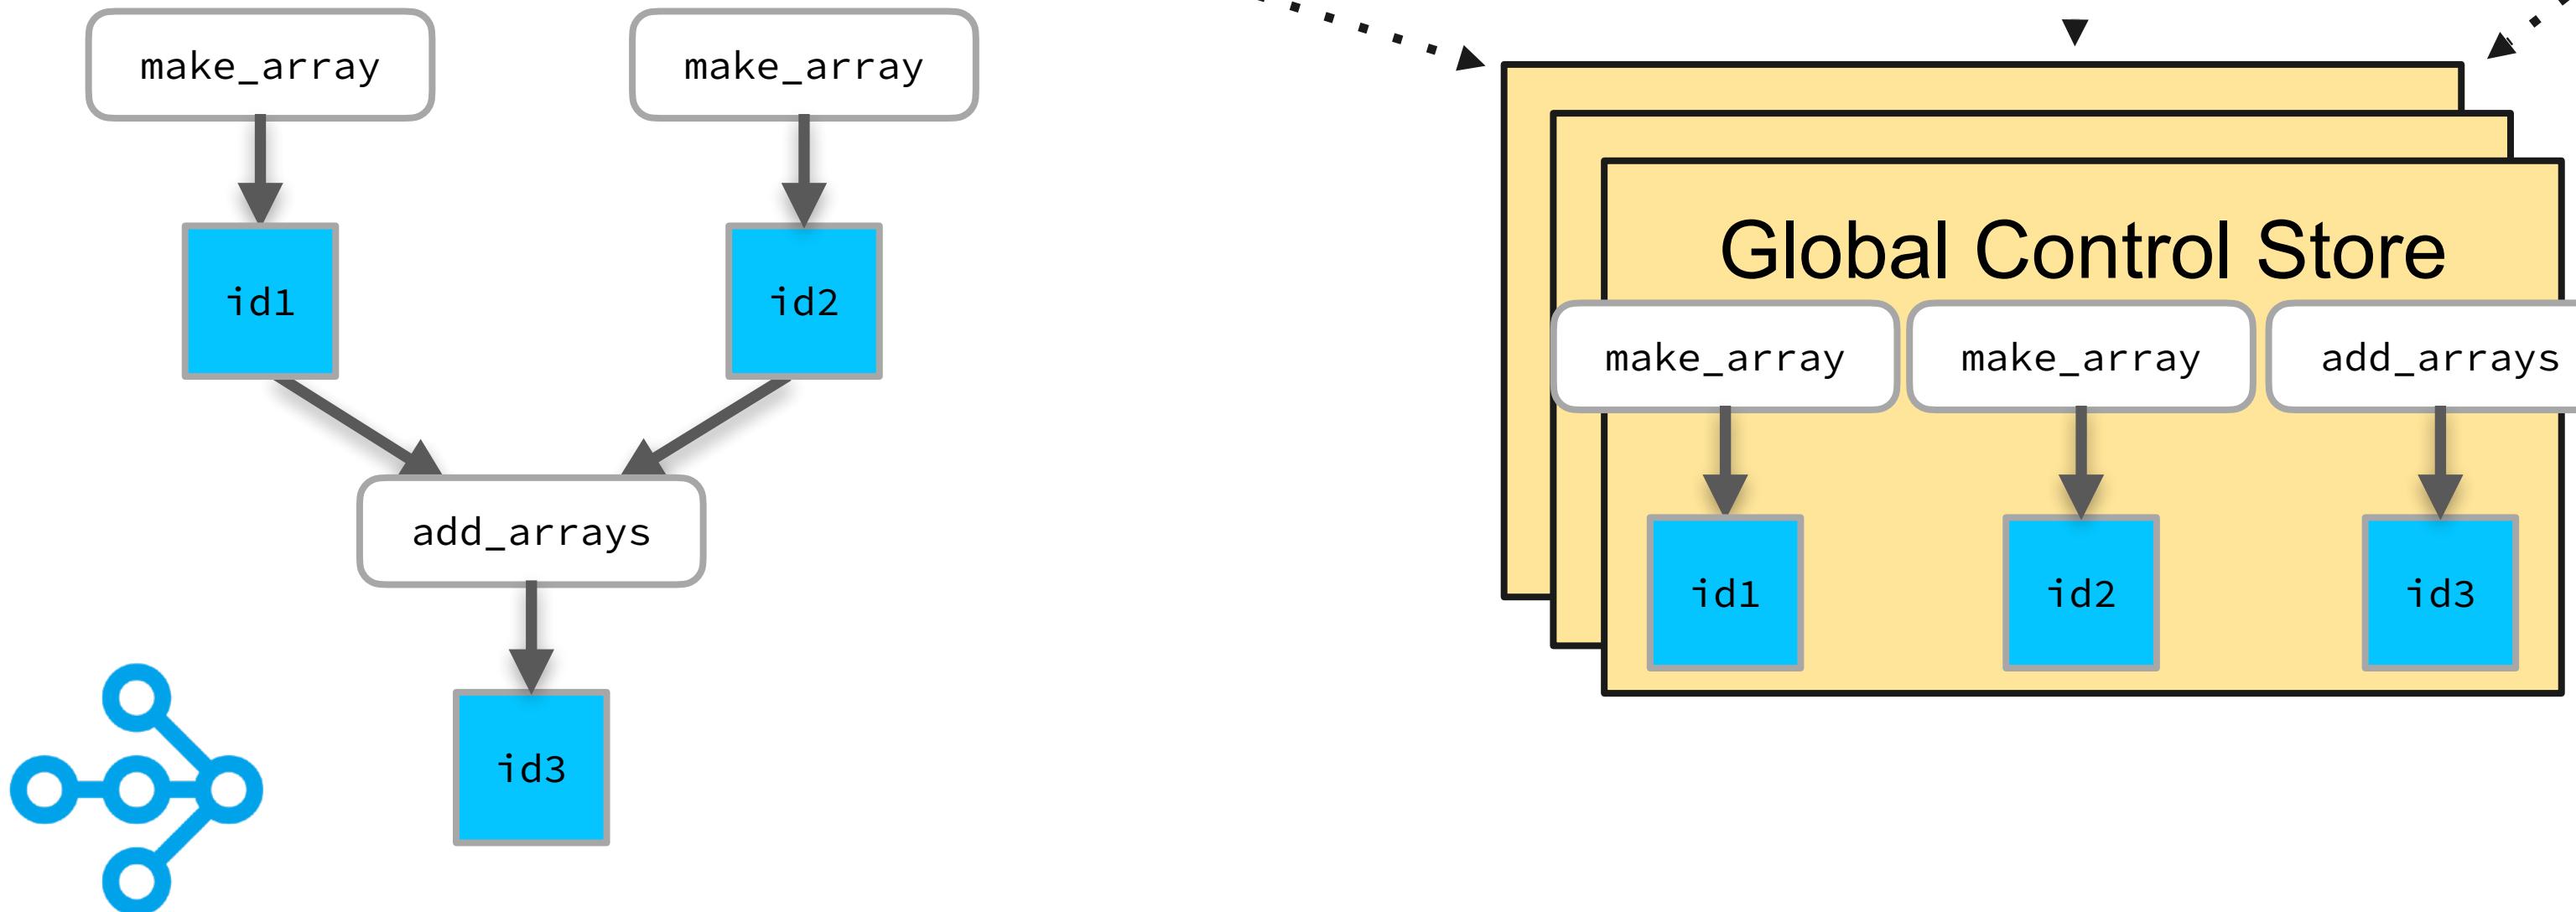
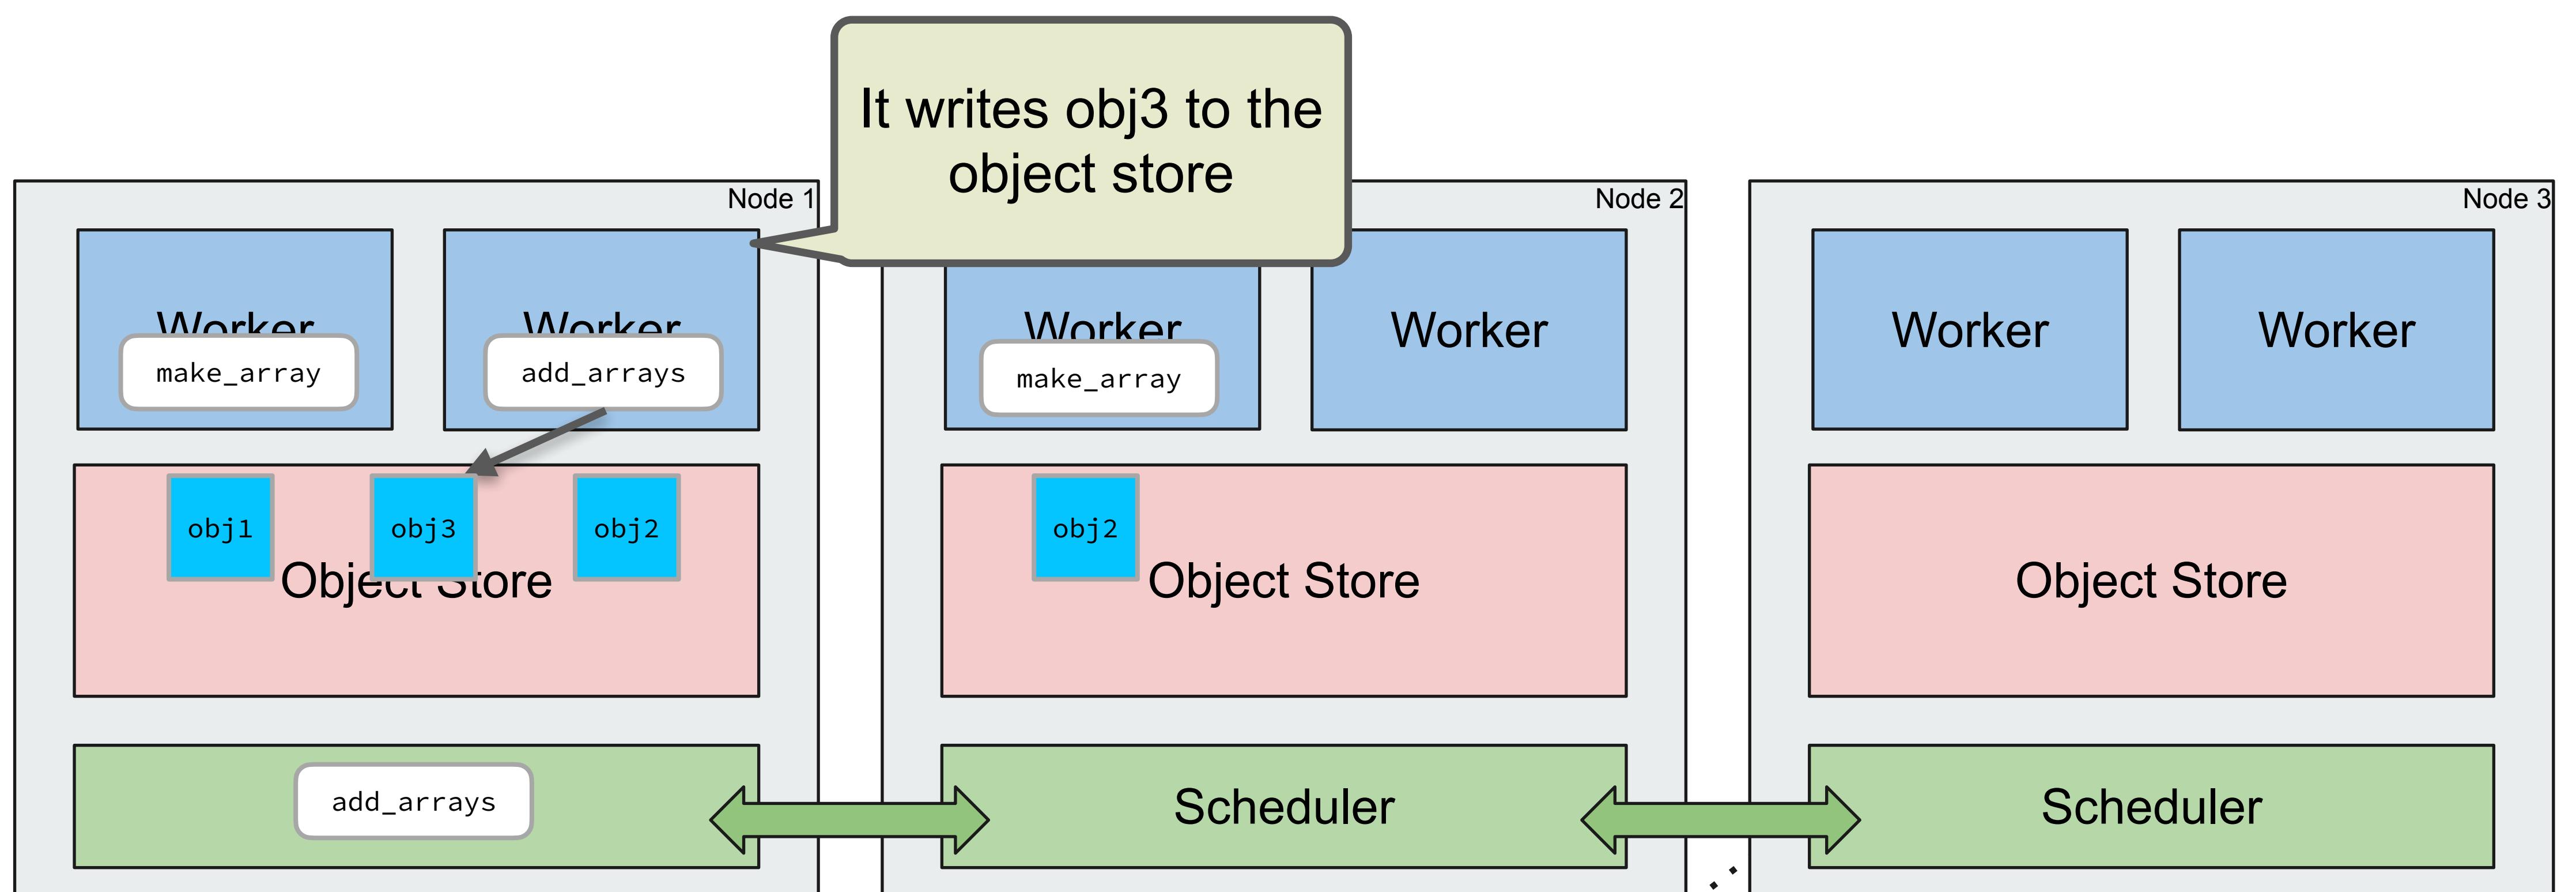
No need to copy
to the worker's
memory!

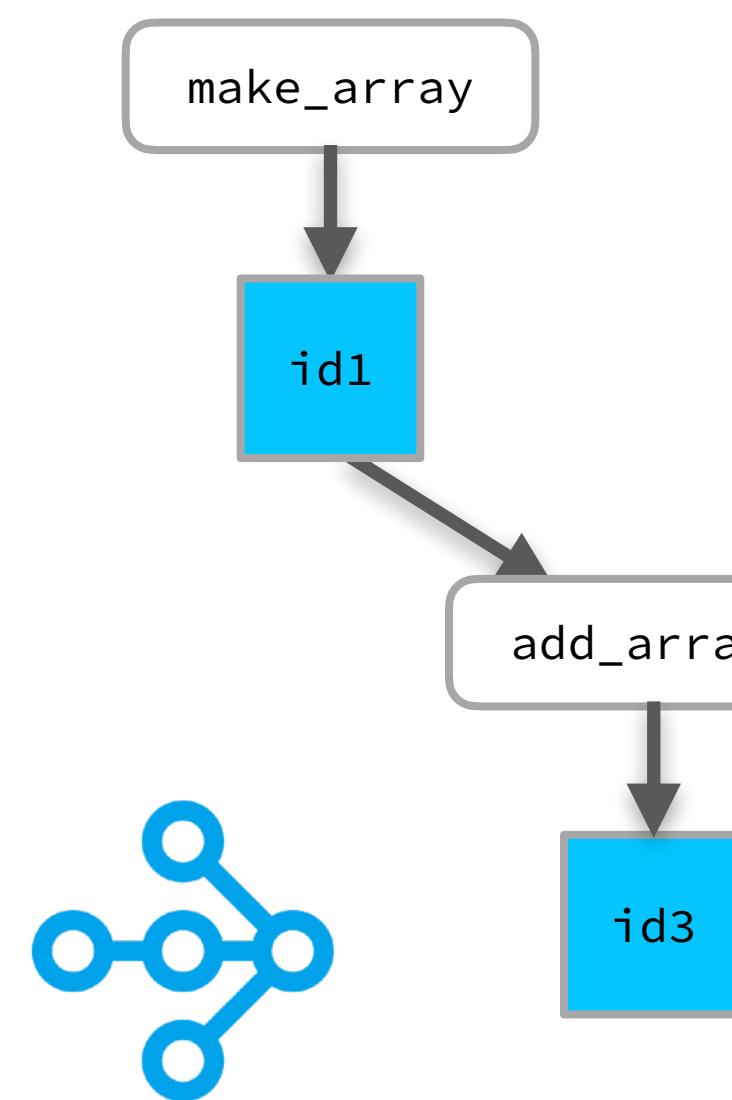
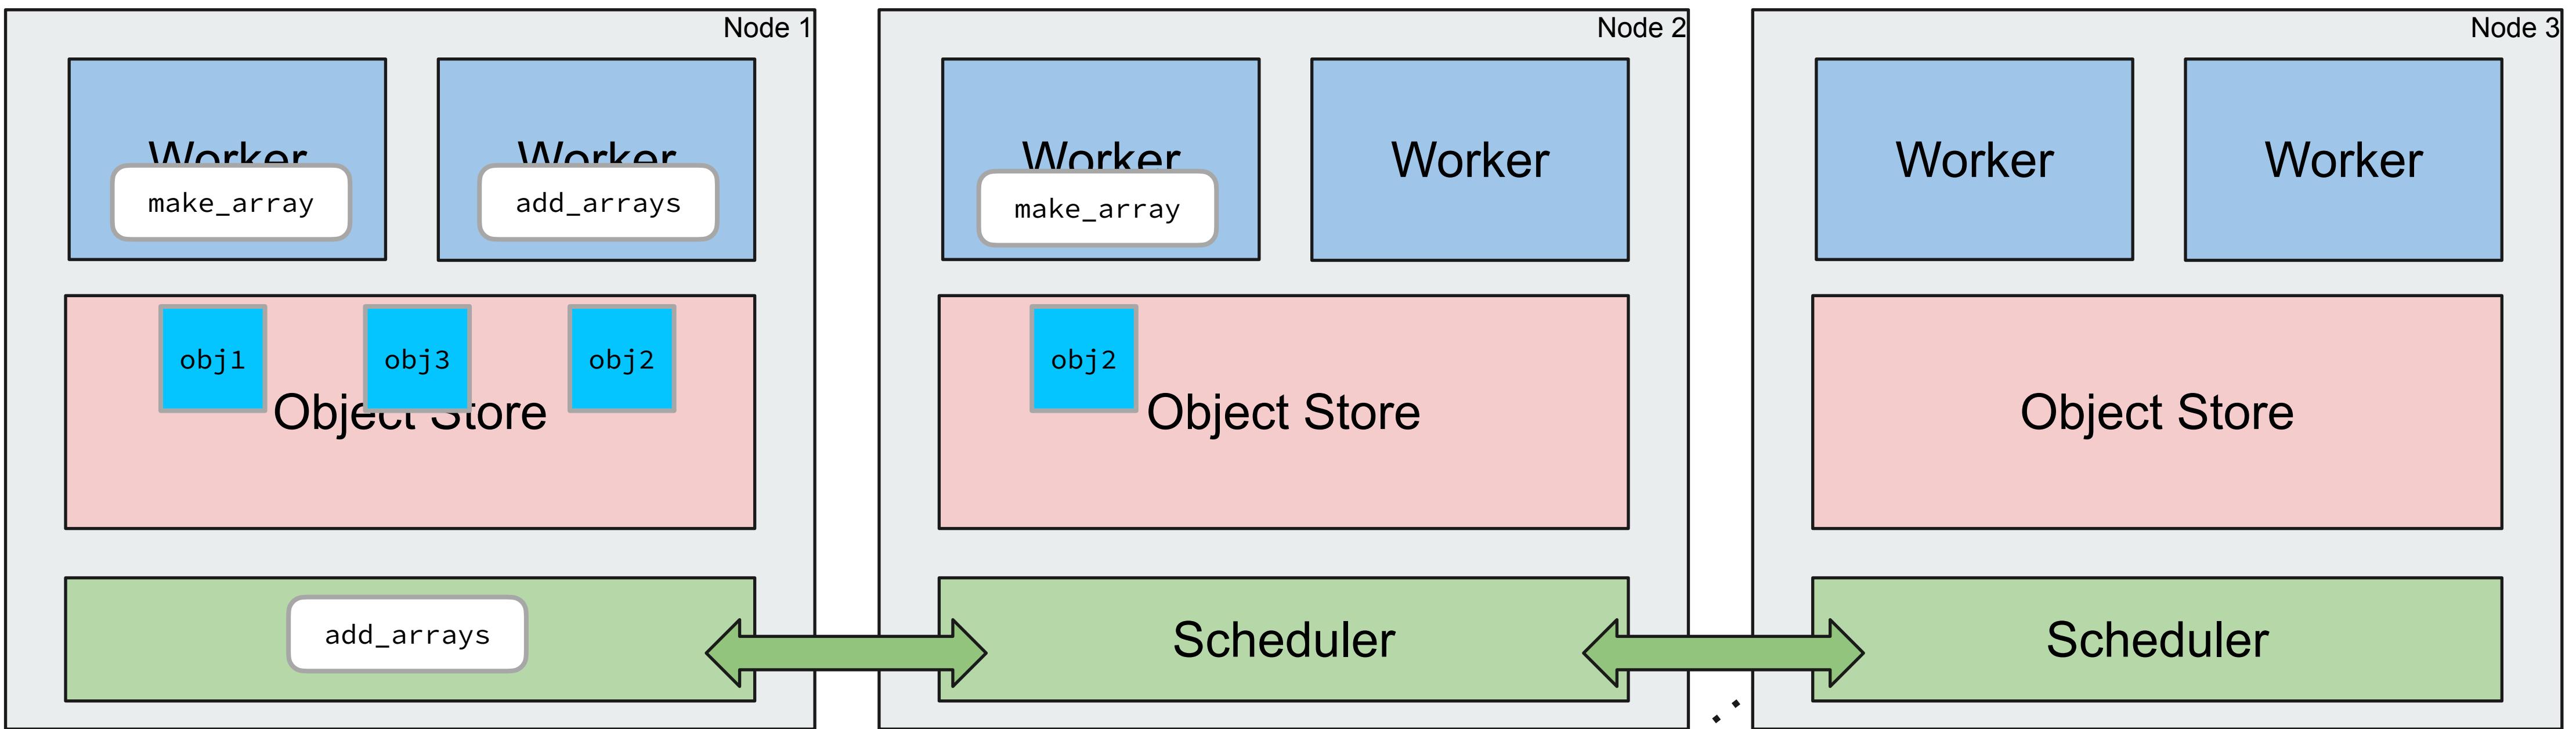
It can read obj1
from *shared
memory*



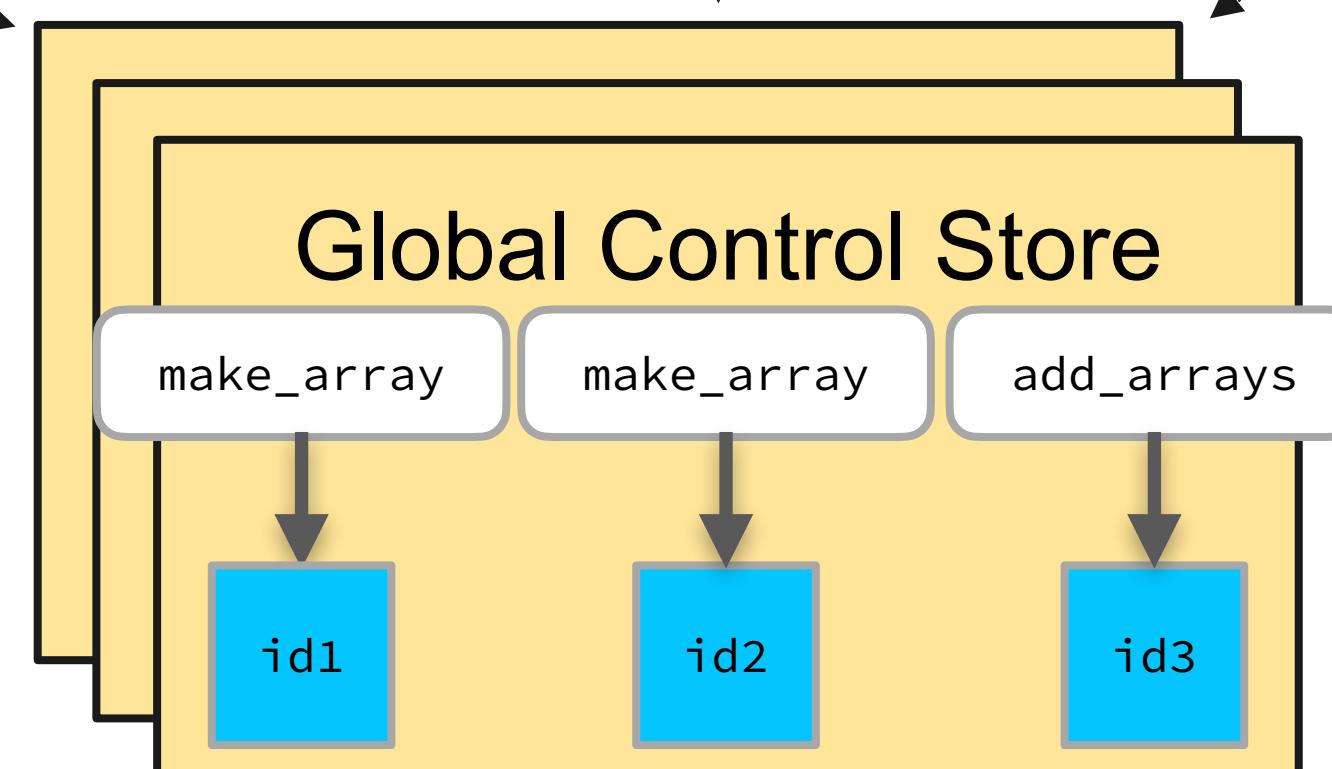








ray.get(id3)
returns obj3



And we're done!



Migrating to Ray

If you're already using...

- `asyncio`
- `joblib`
- `multiprocessing.Pool`

For example, from this:

```
from multiprocessing.pool import Pool
```

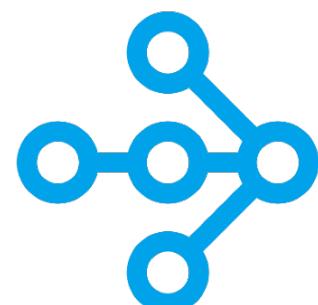
To this:

```
from ray.util.multiprocessing.pool import Pool
```

- Use Ray's implementations
 - Drop-in replacements
 - Change import statements
 - Break the one-node limitation!

See this blog post:

<https://medium.com/distributed-computing-with-ray/how-to-scale-python-multiprocessing-to-a-cluster-with-one-line-of-code-d19f242f60ff>

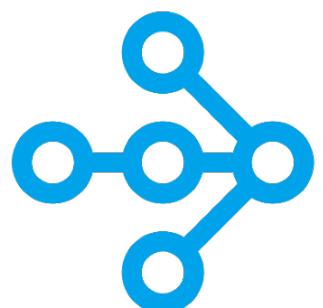
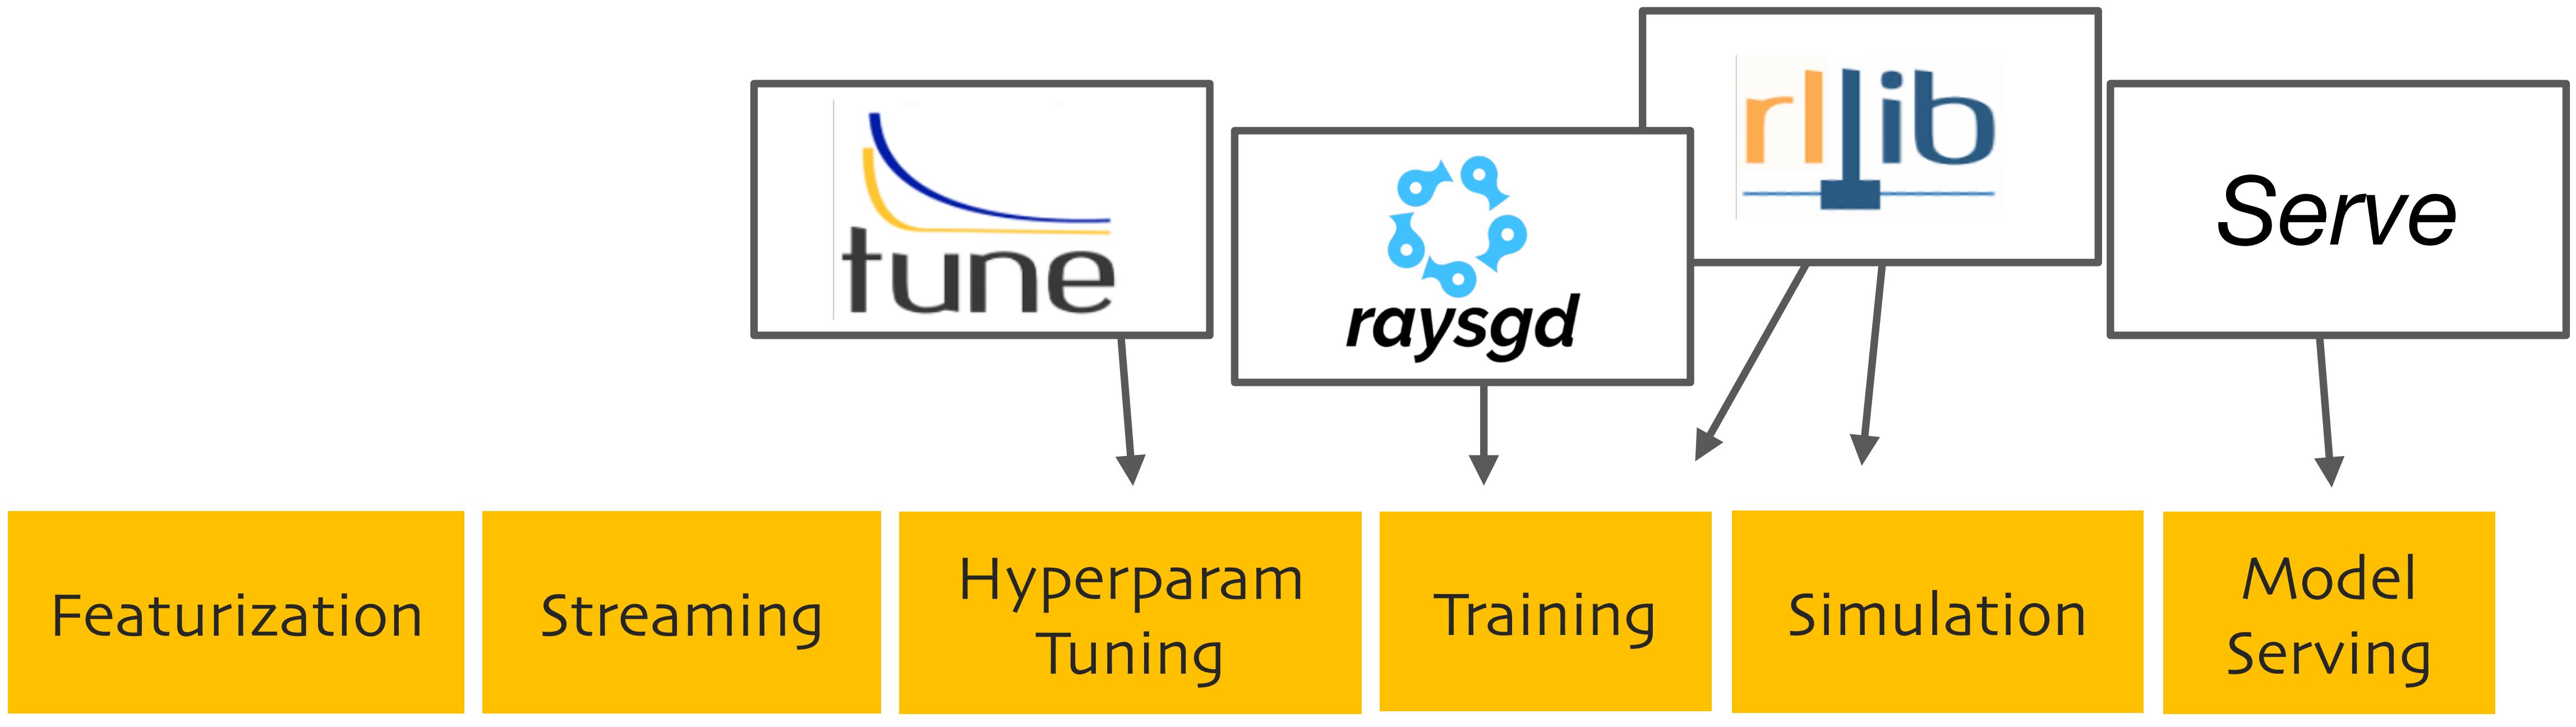


@deanwampler



Machine Learning with Ray-based Libraries

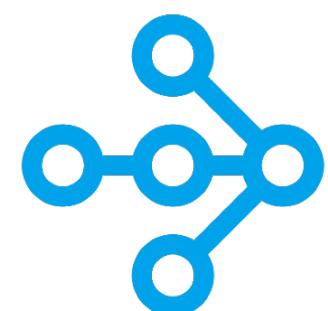
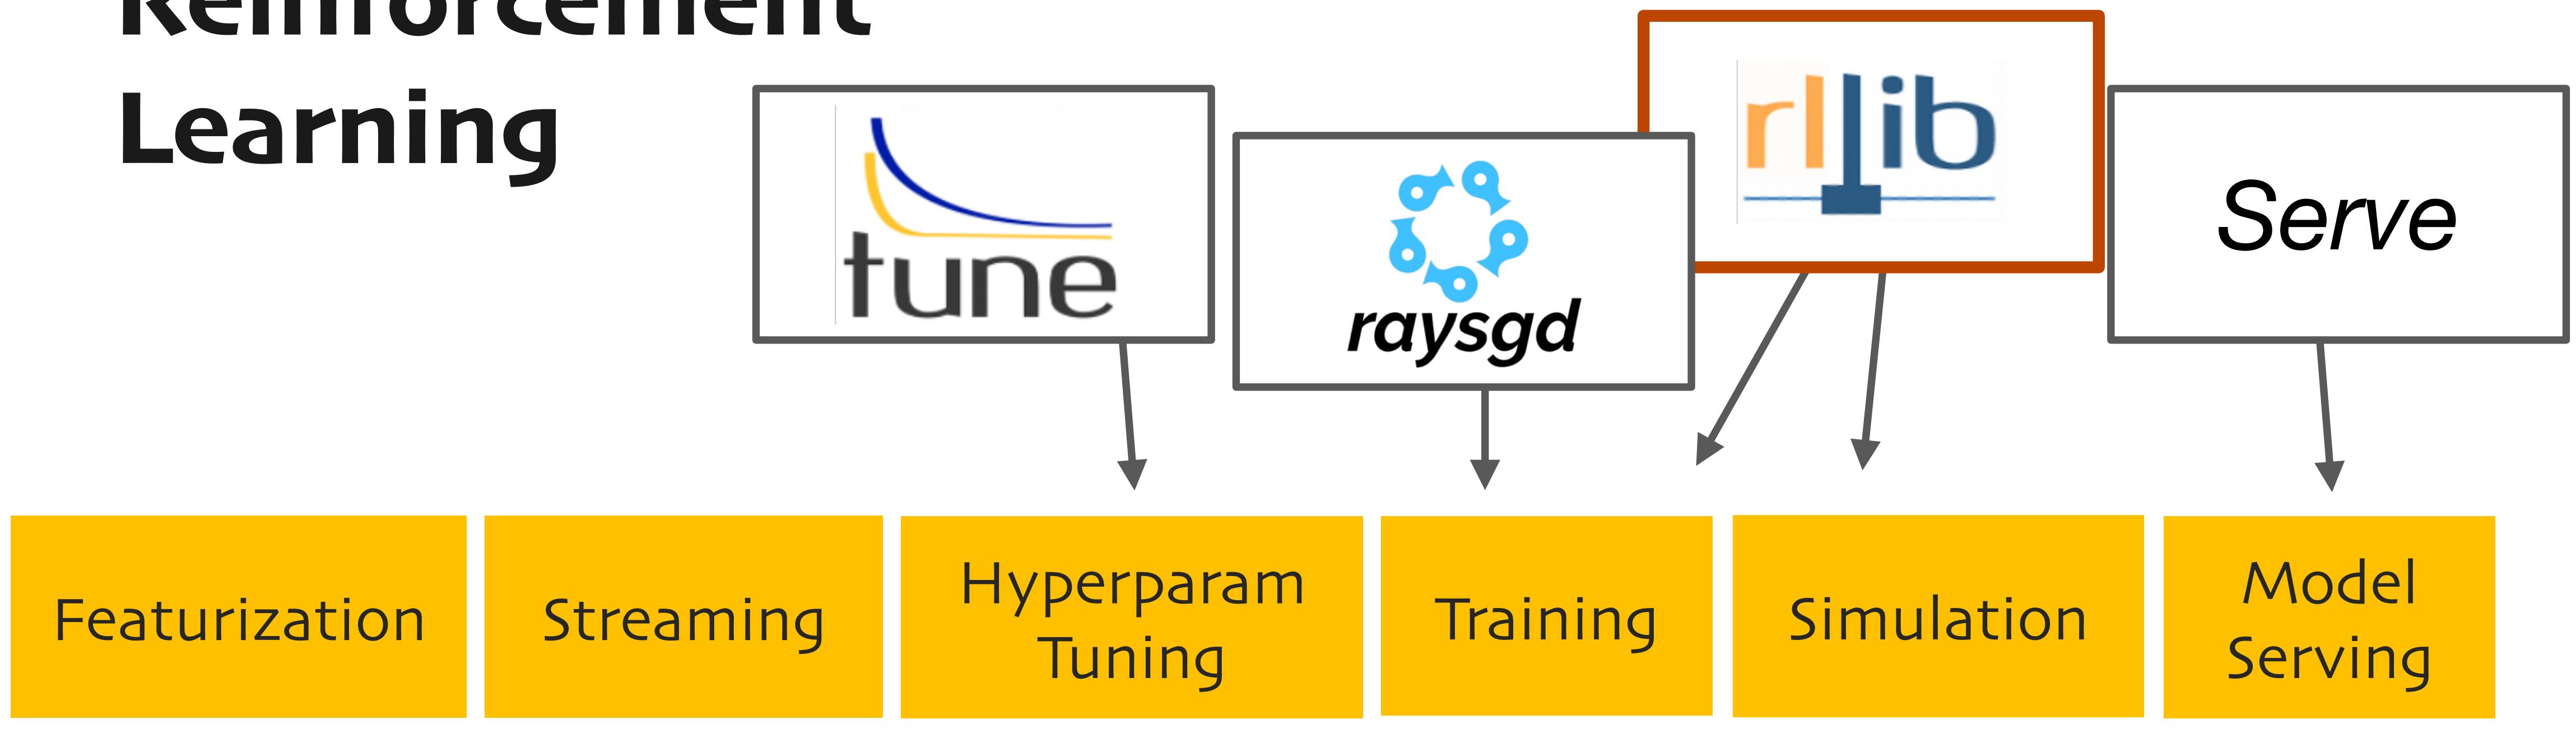
Ray Libraries



@deanwampler

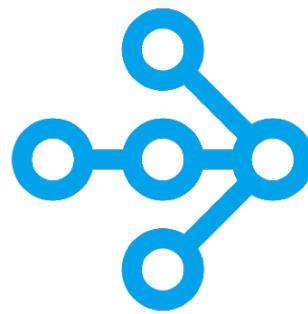
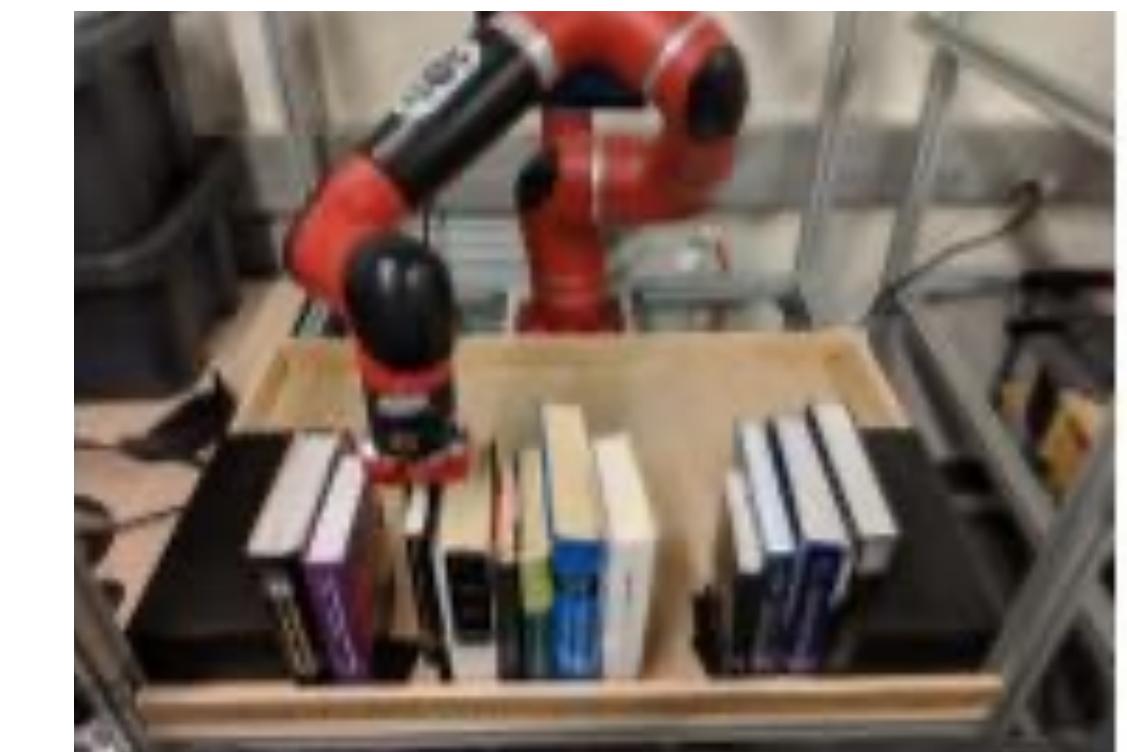
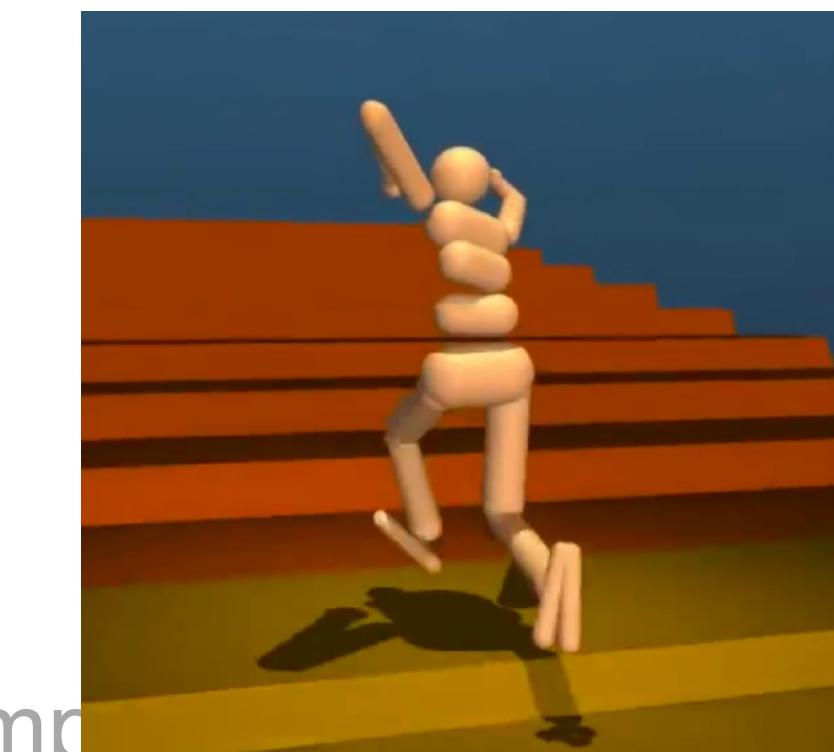
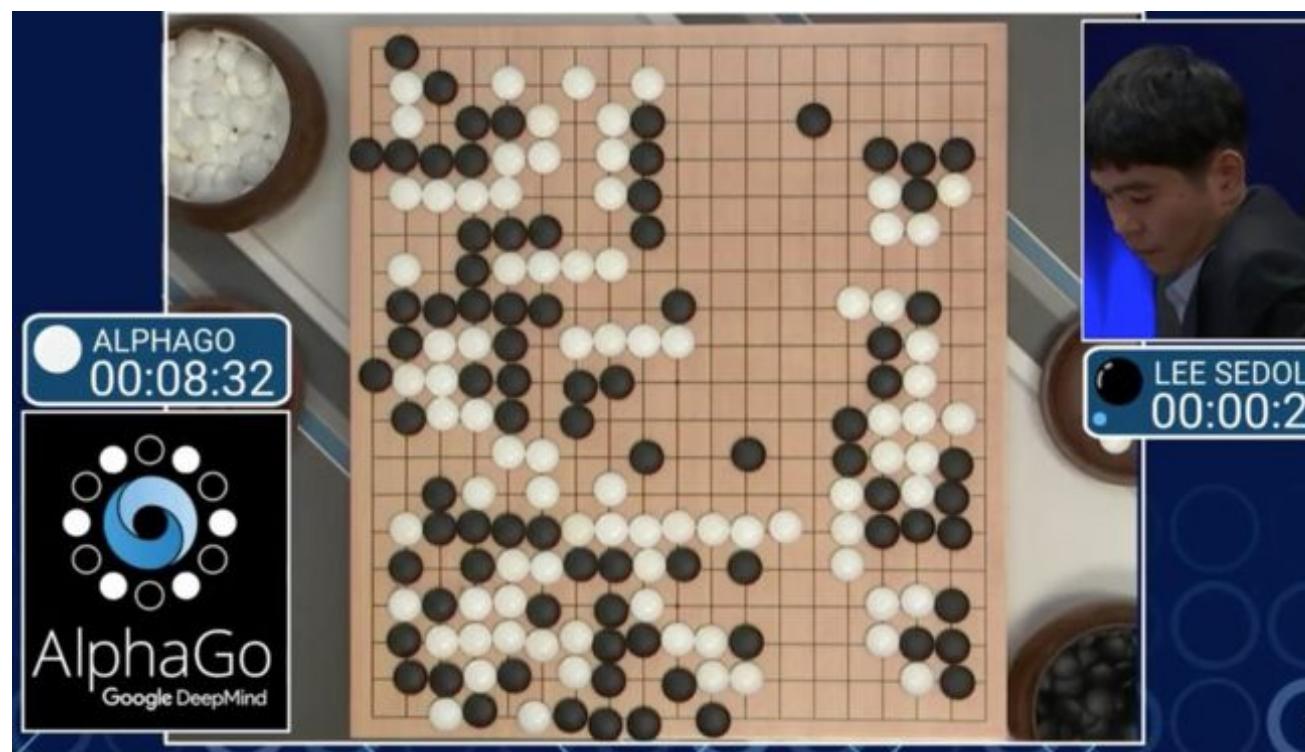
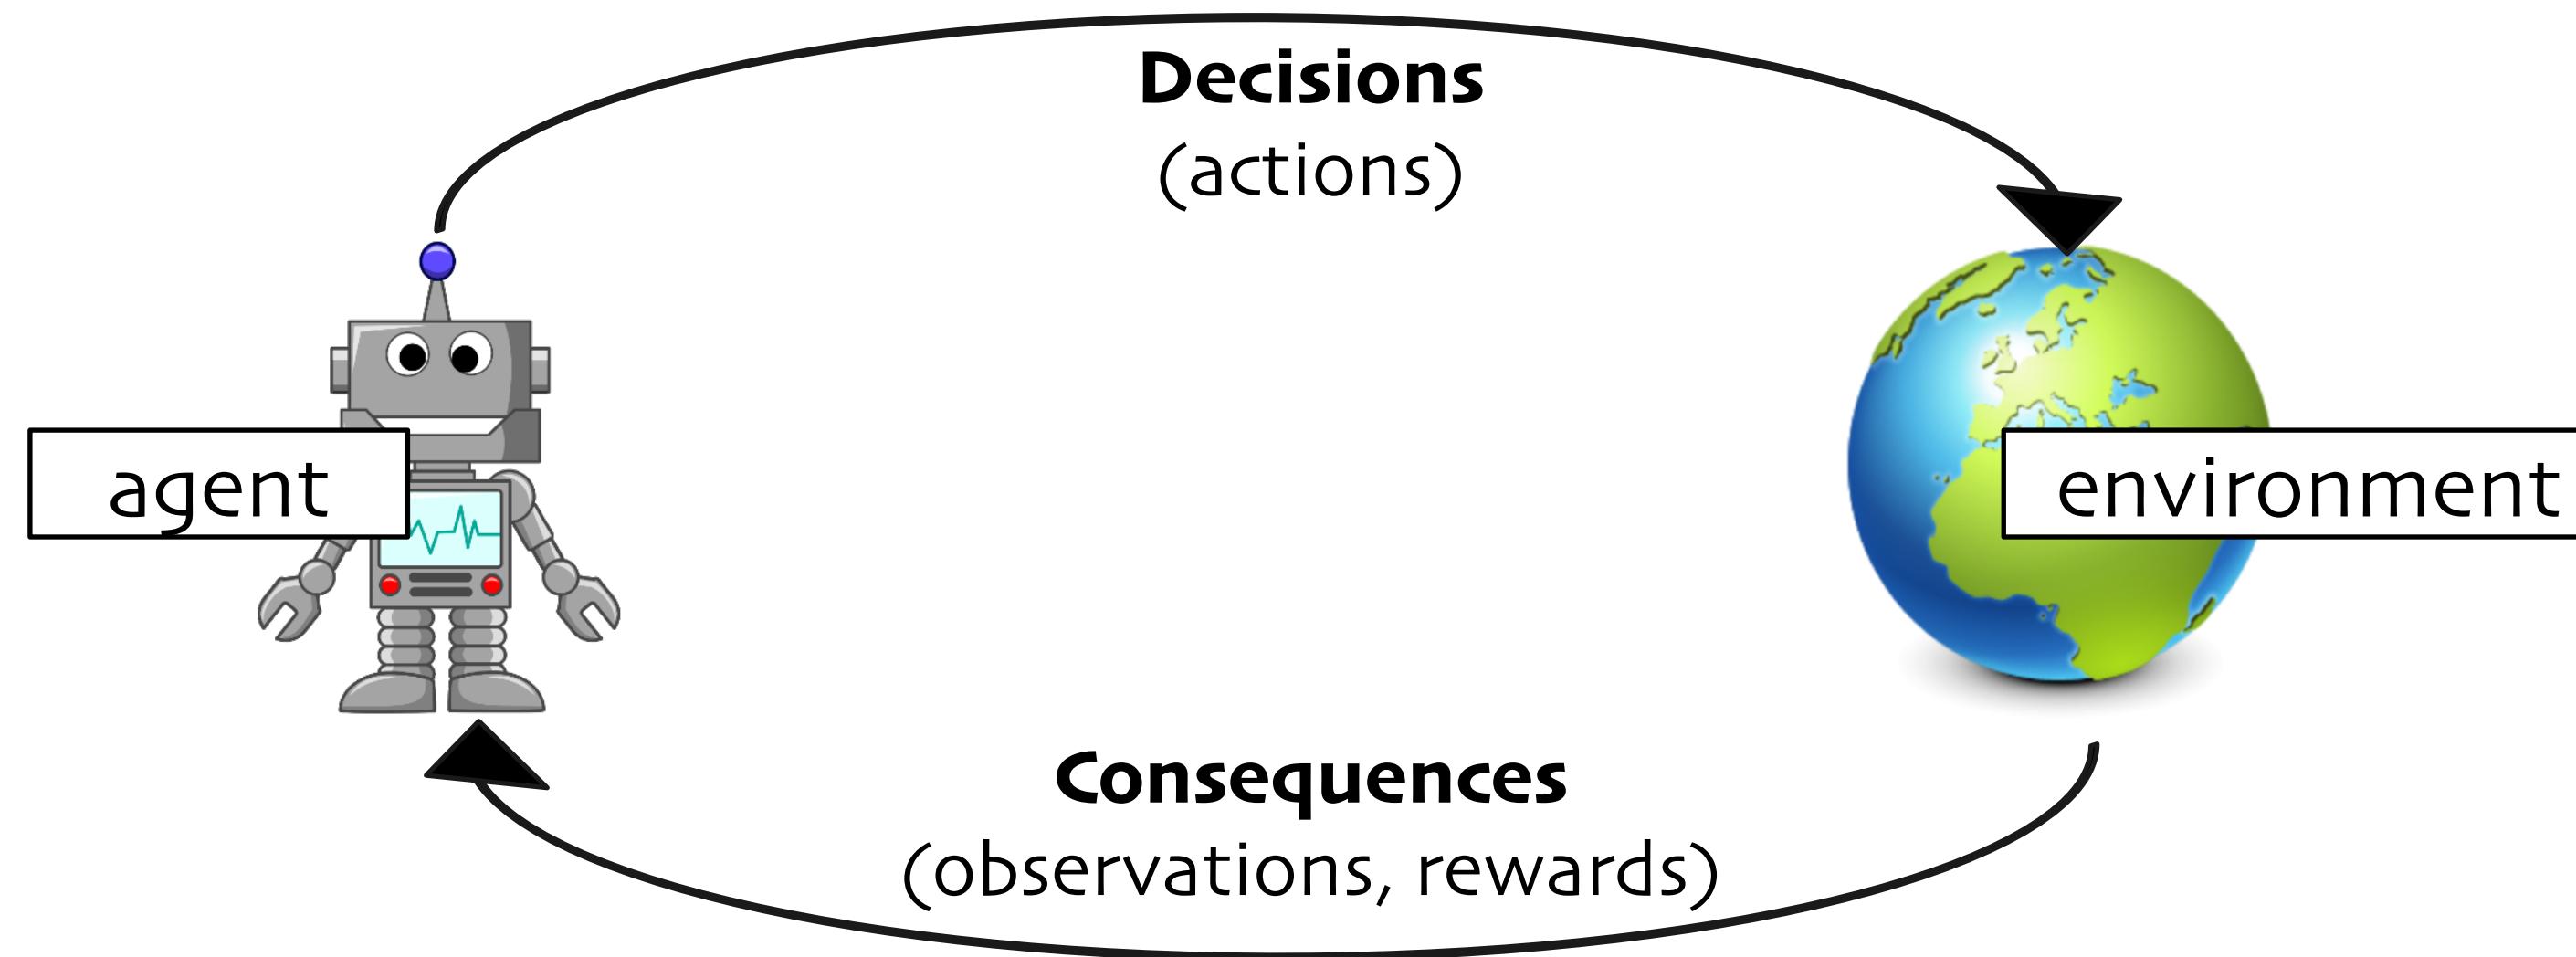
More libraries
coming soon

Reinforcement Learning



@deanwampler

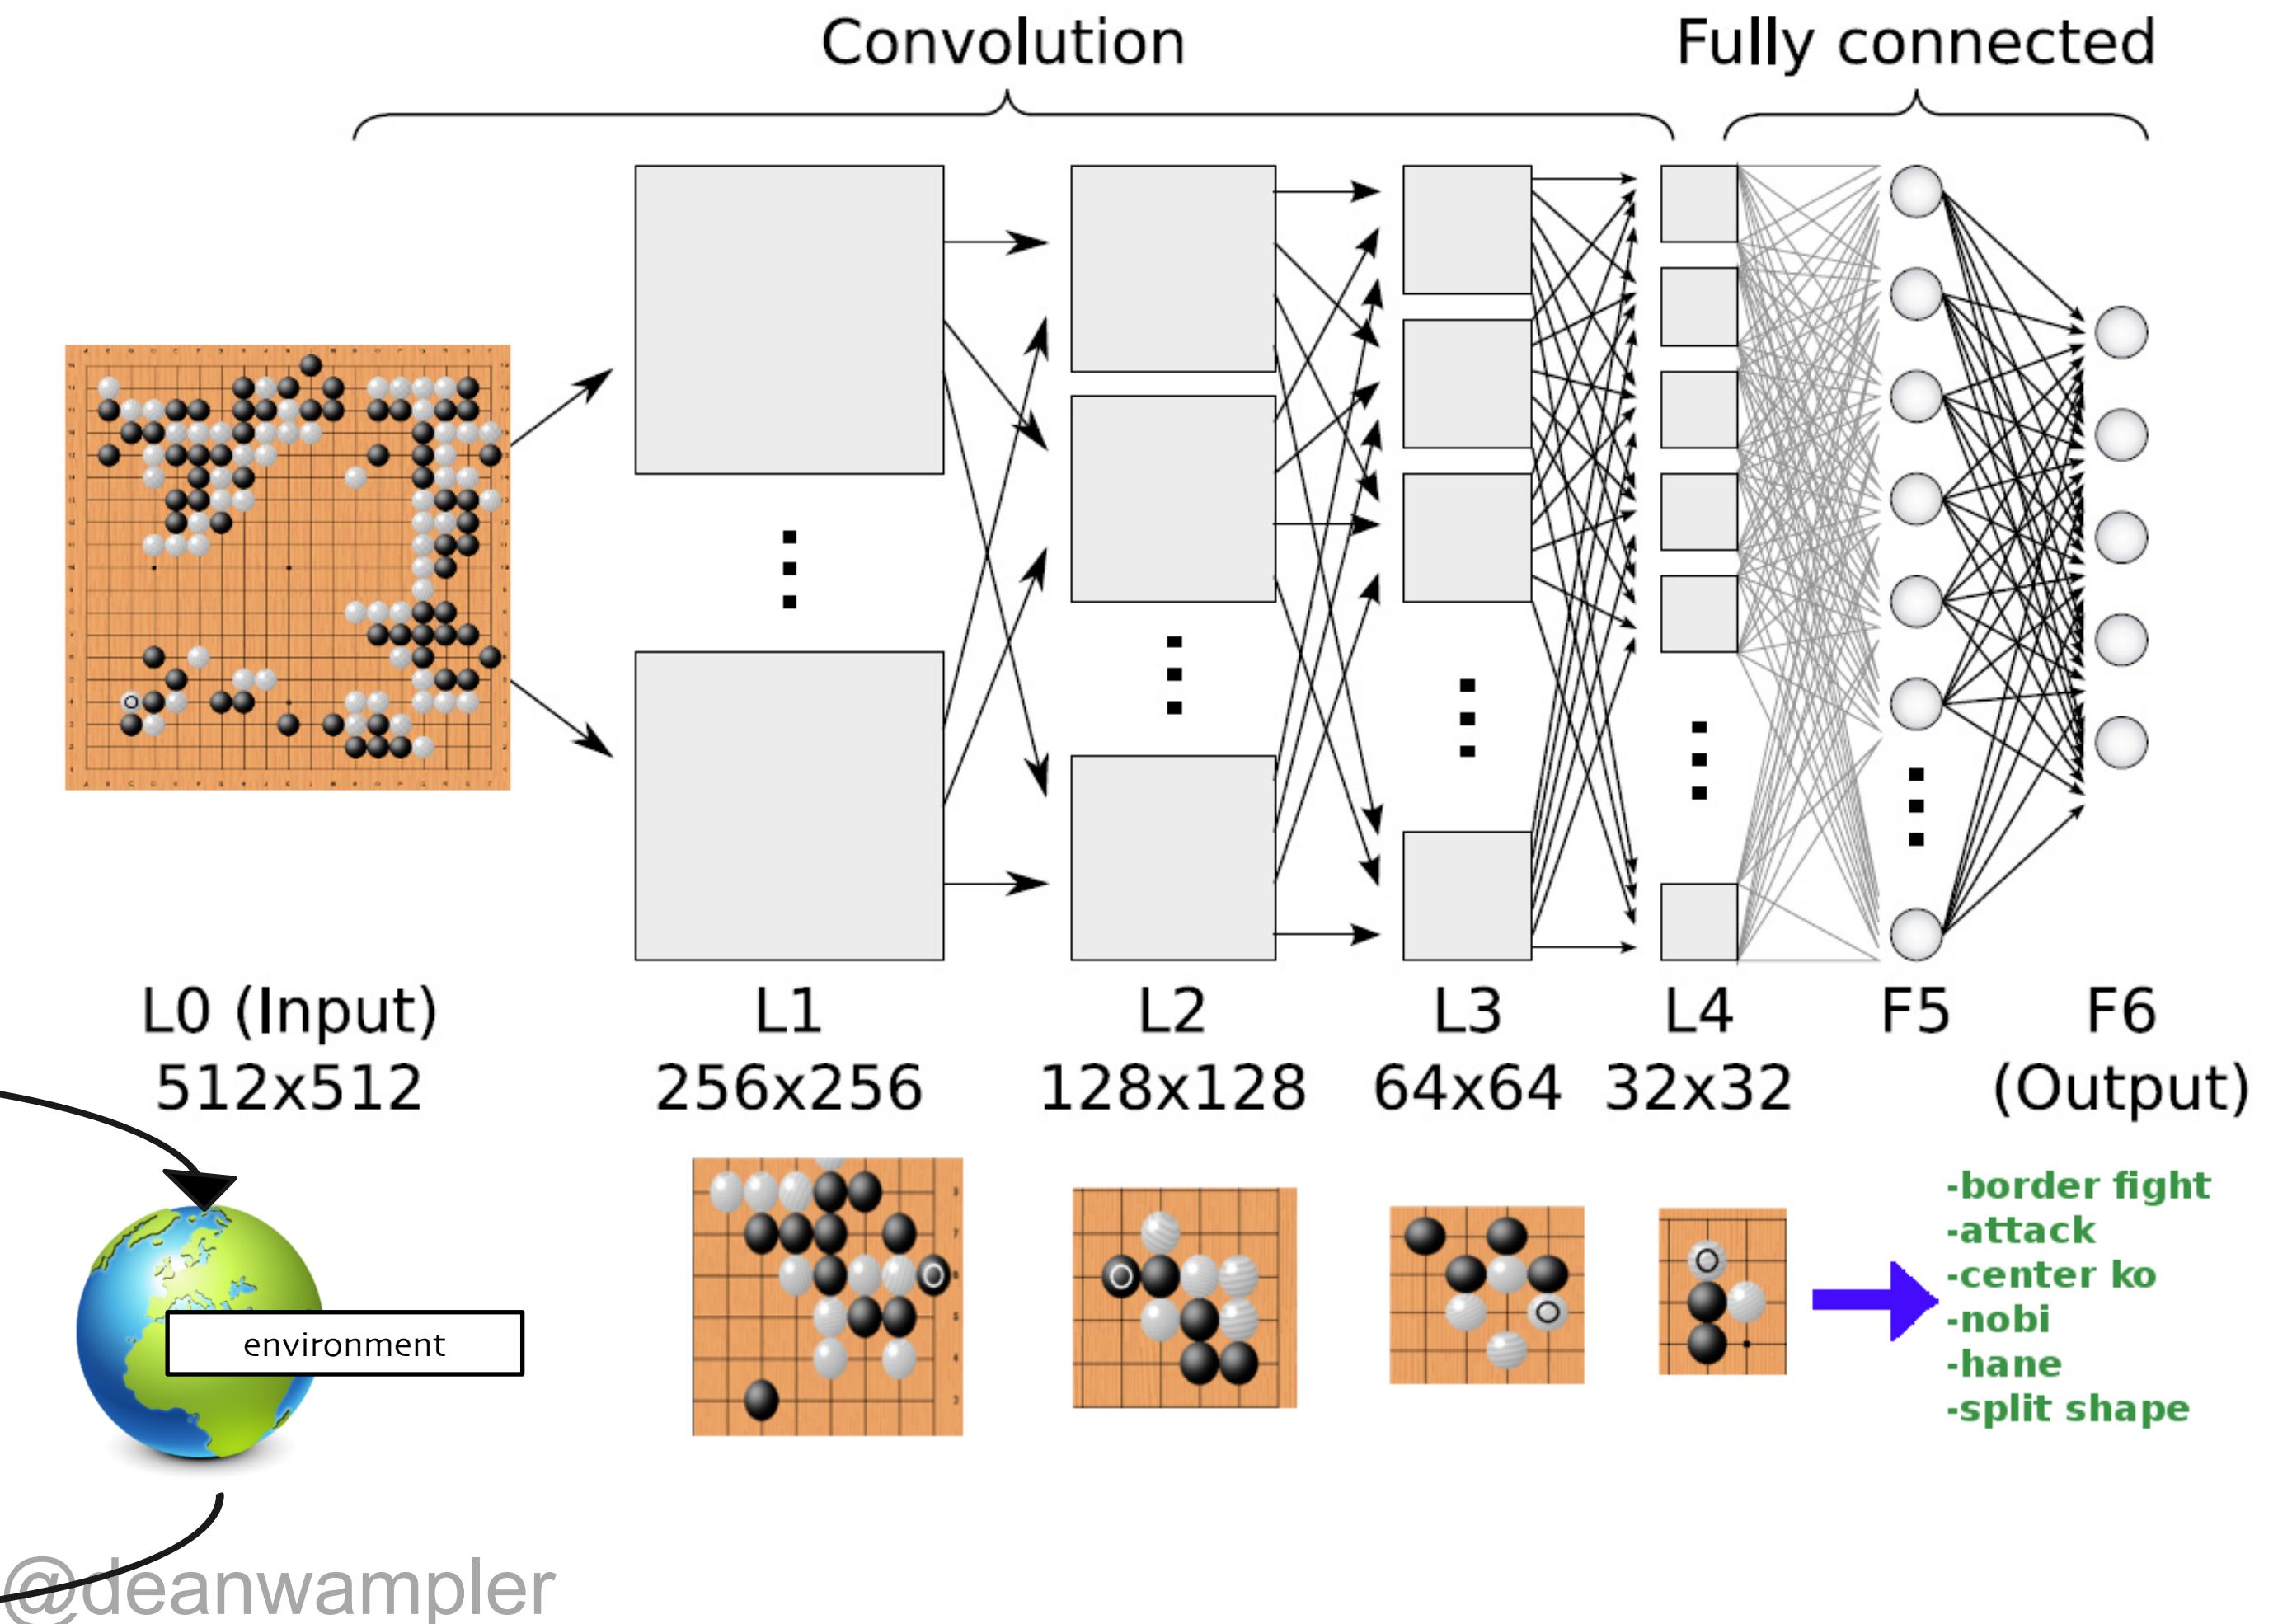
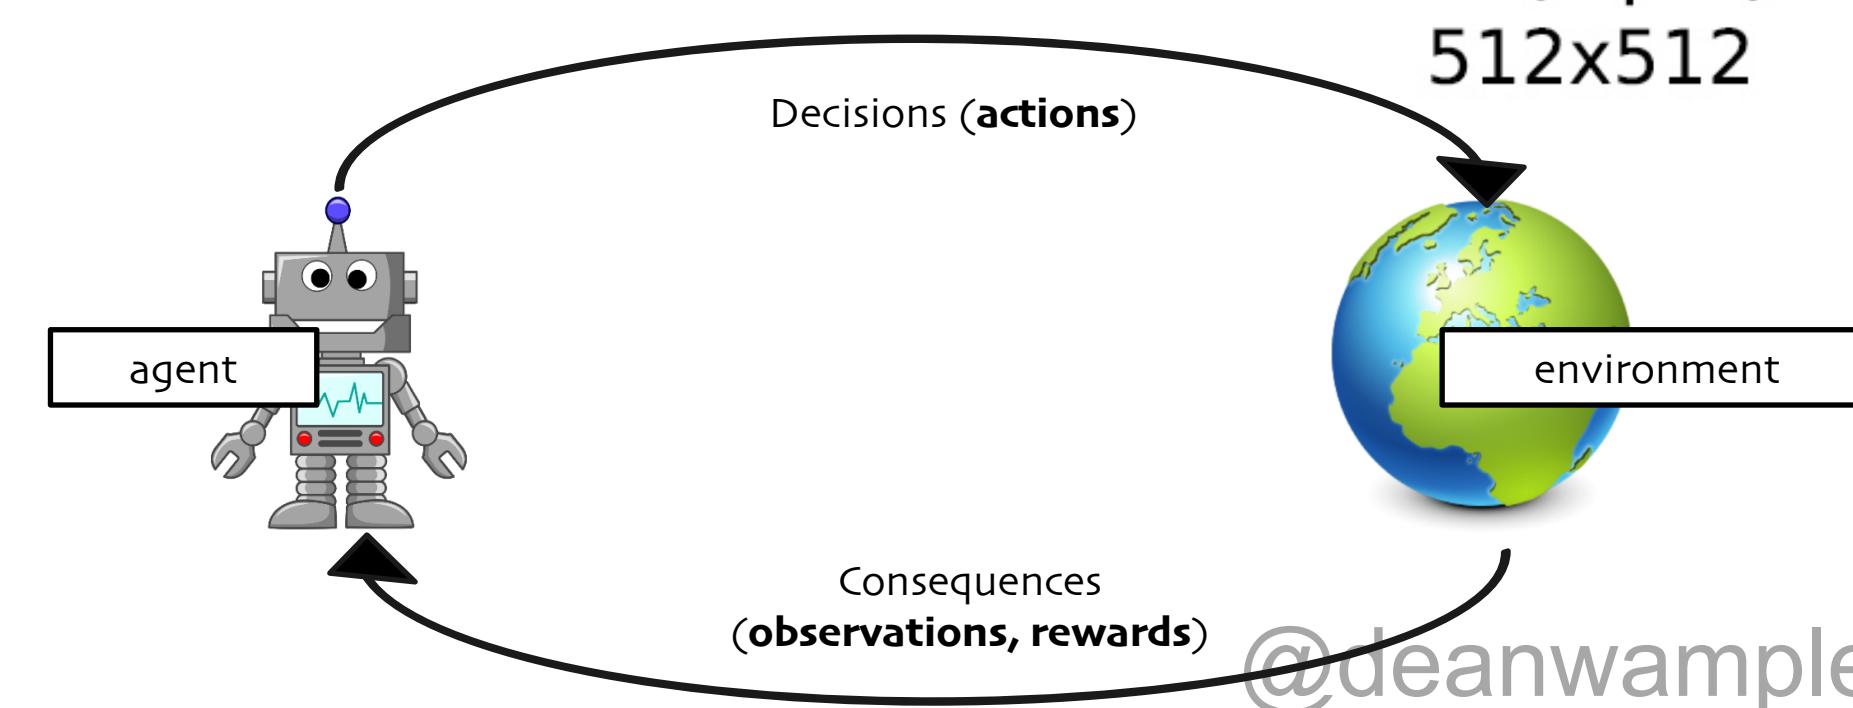
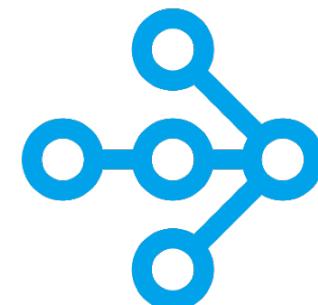
Background: Reinforcement Learning



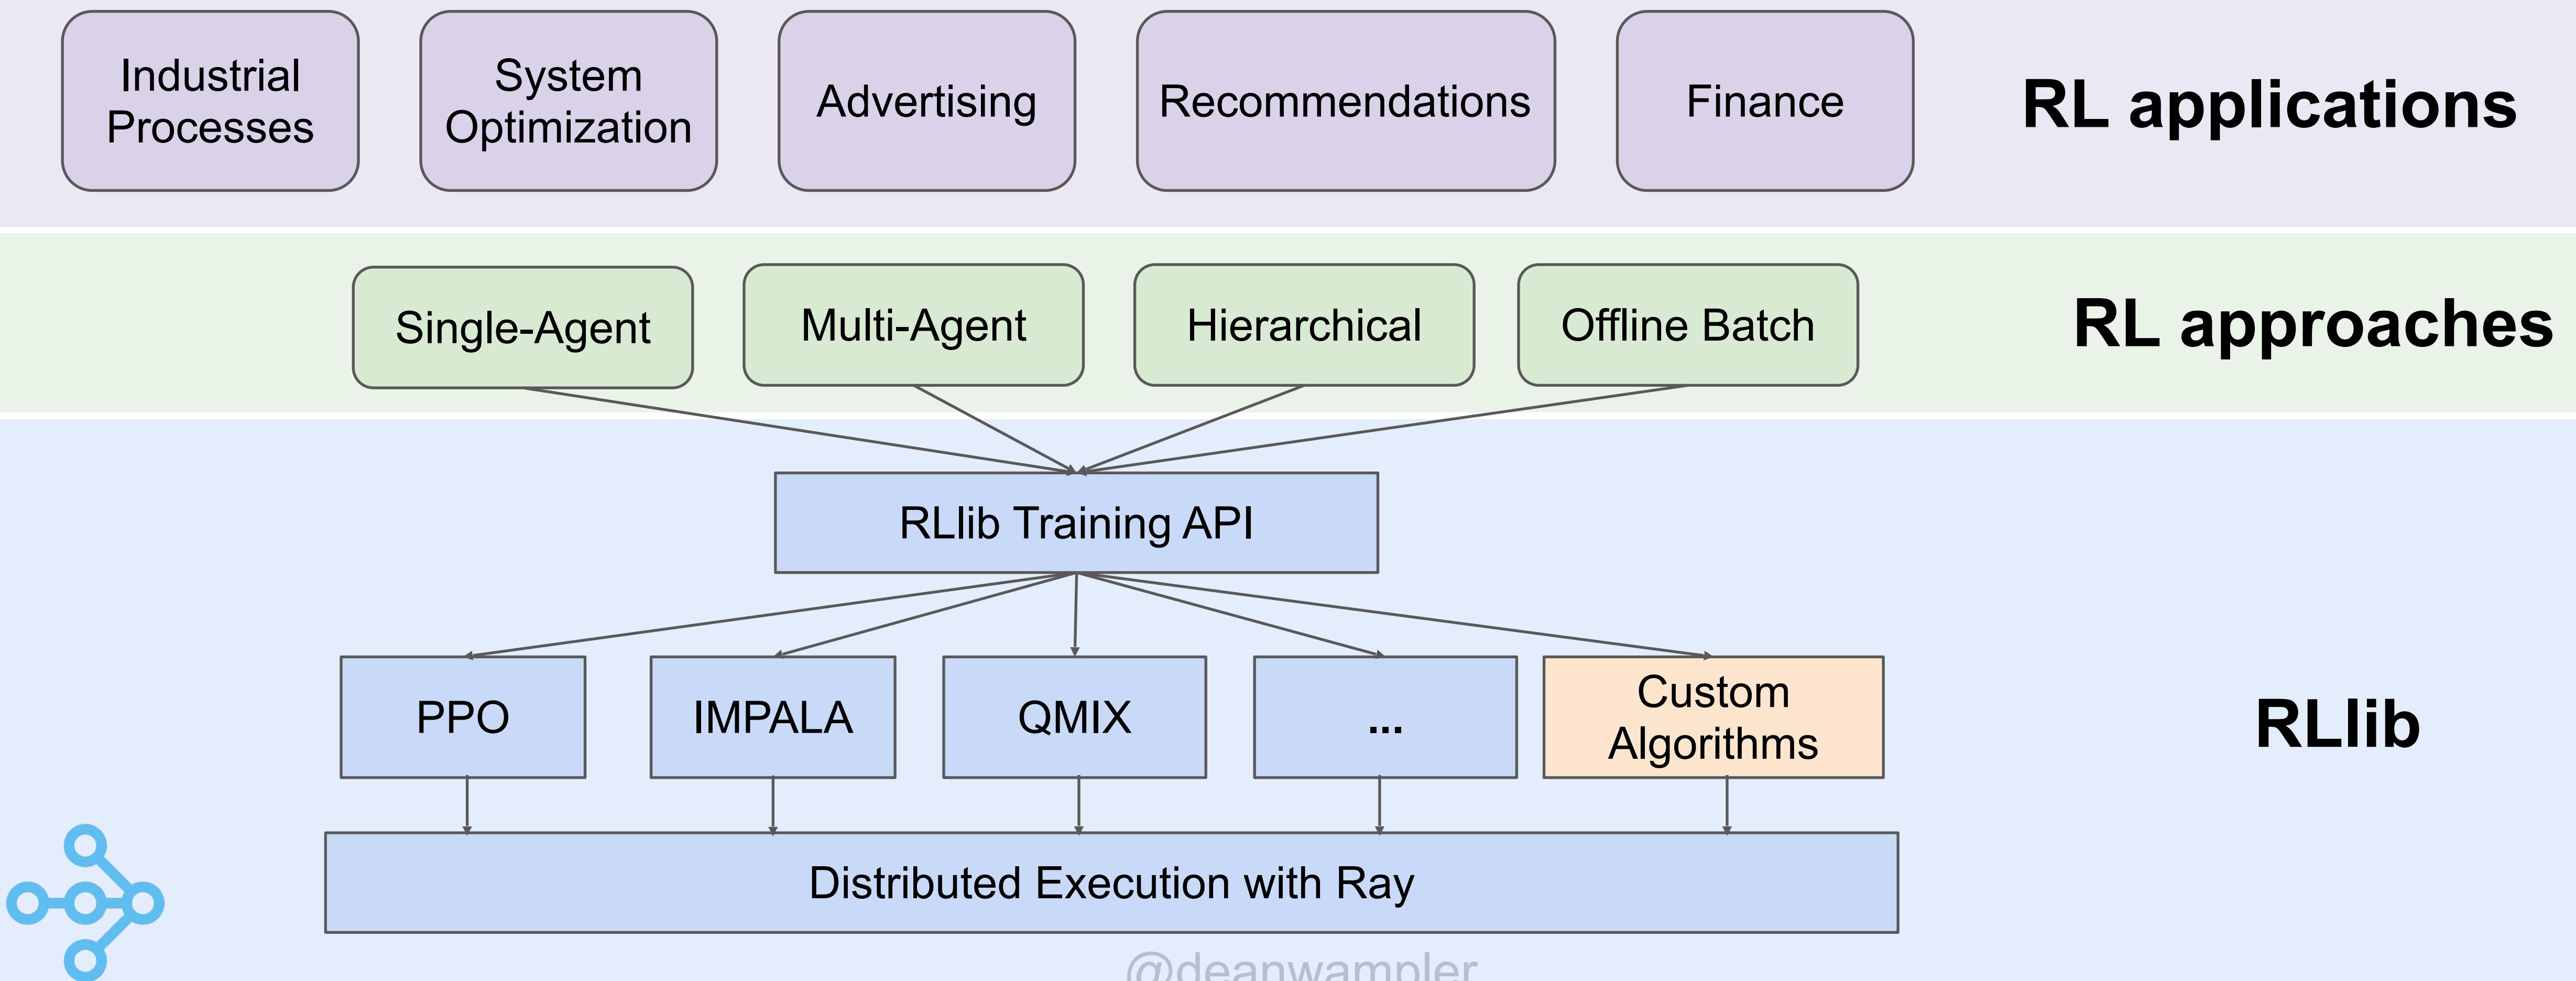
Go as a Reinforcement Learning Problem

AlphaGo (Silver et al. 2016)

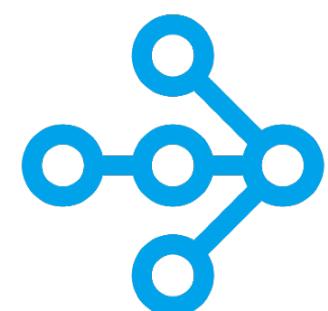
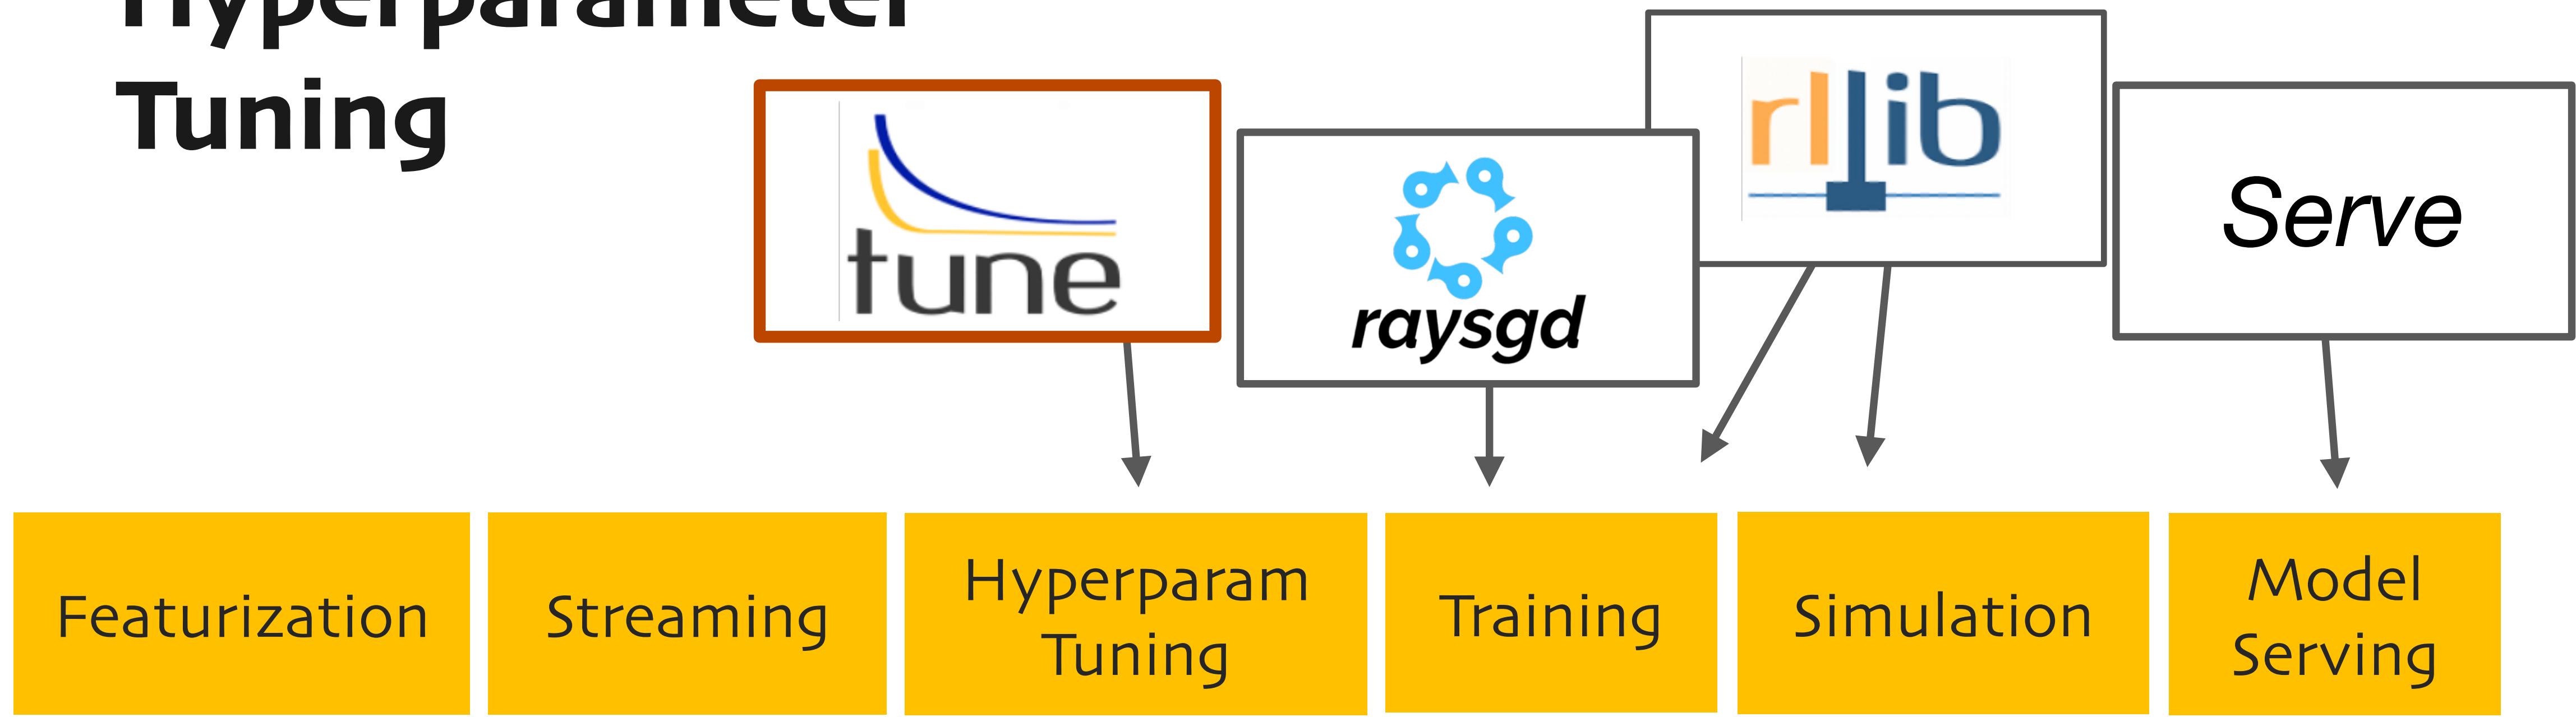
- **Observations:**
 - board state
- **Actions:**
 - where to place the stones
- **Rewards:**
 - 1 if win
 - 0 otherwise



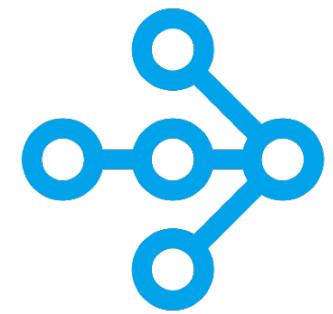
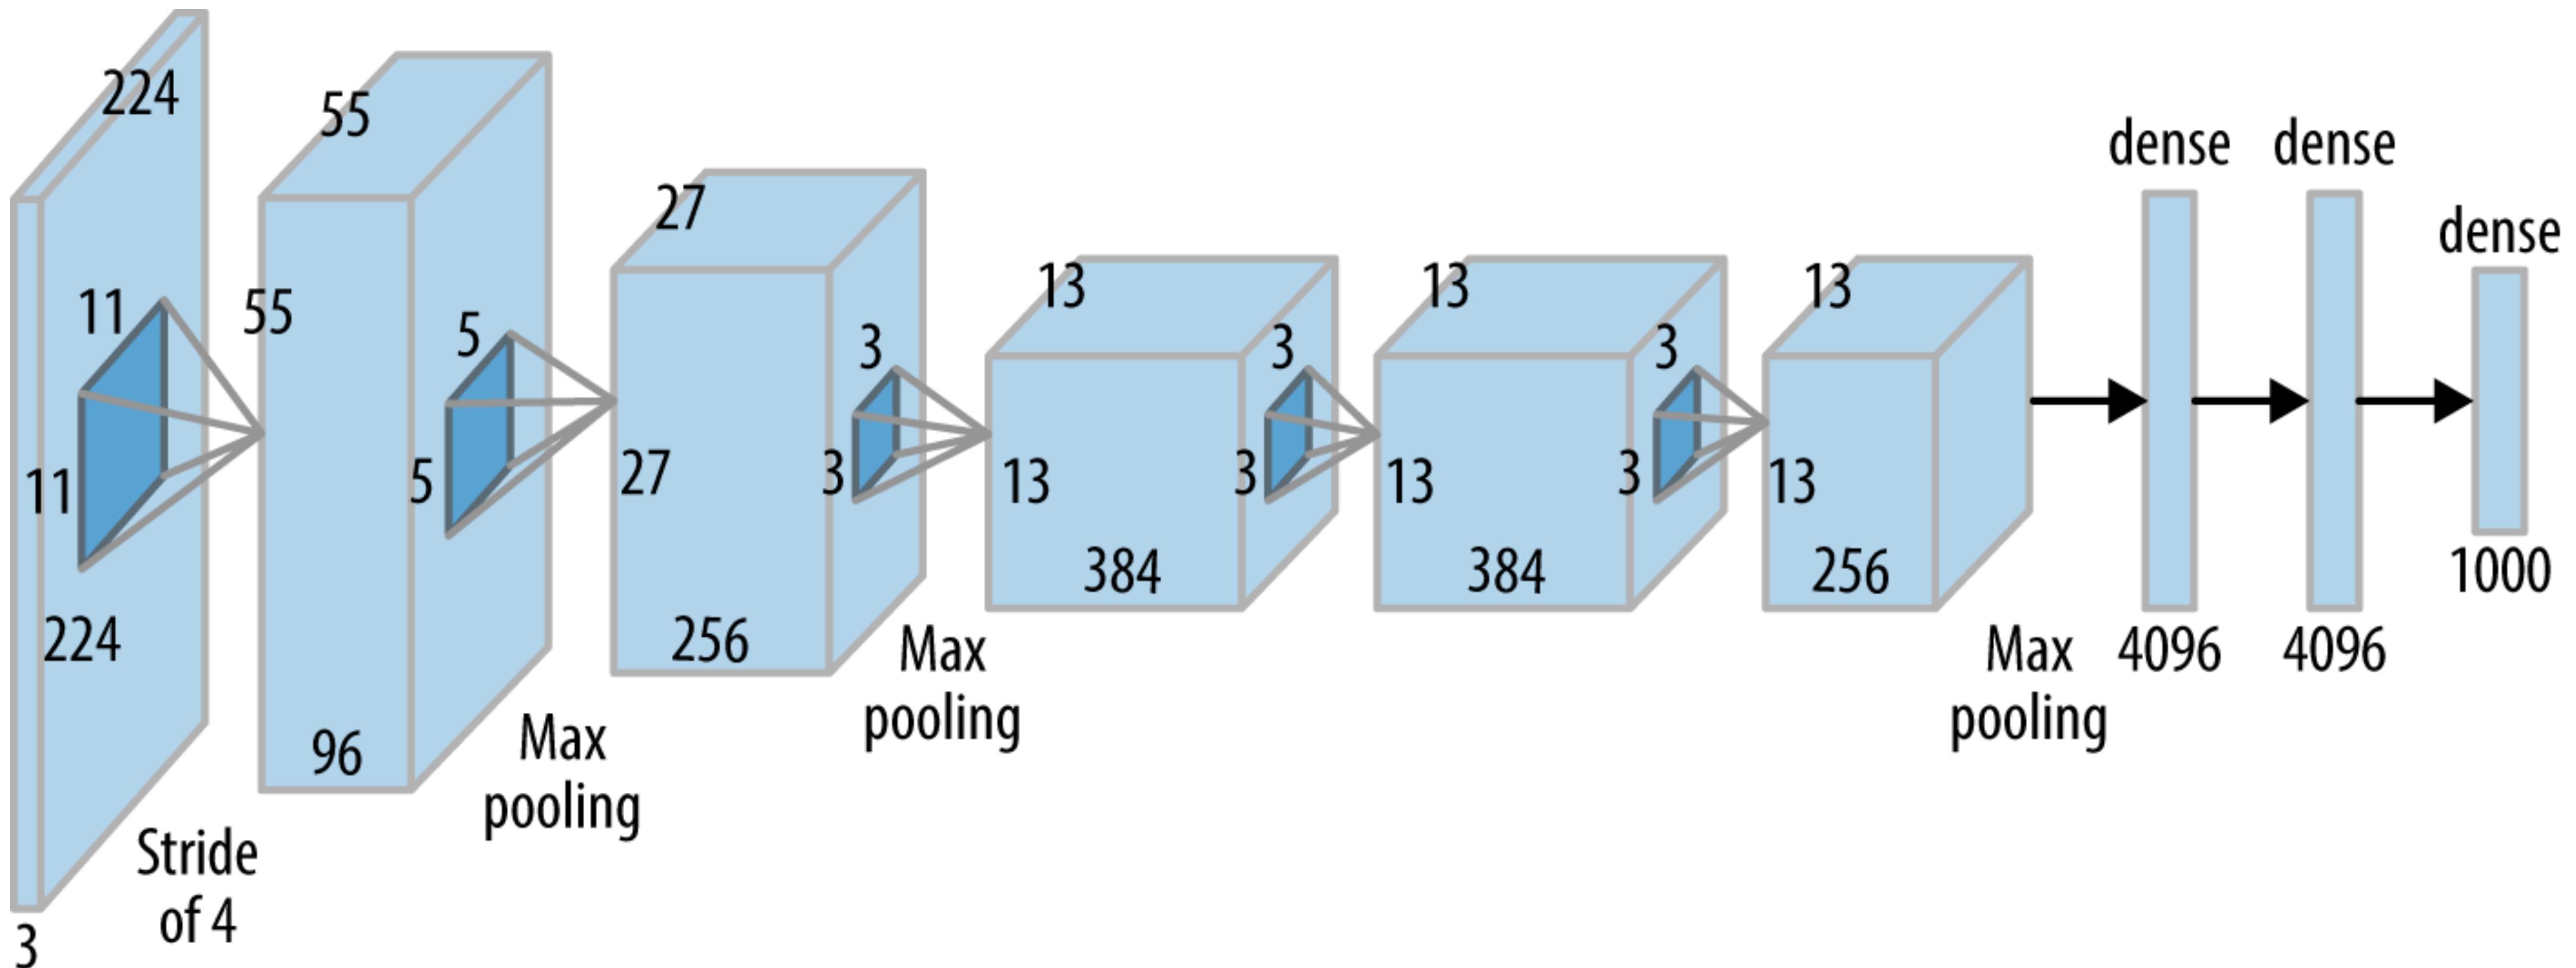
RLLib: A Scalable, Unified Library for RL



Hyperparameter Tuning

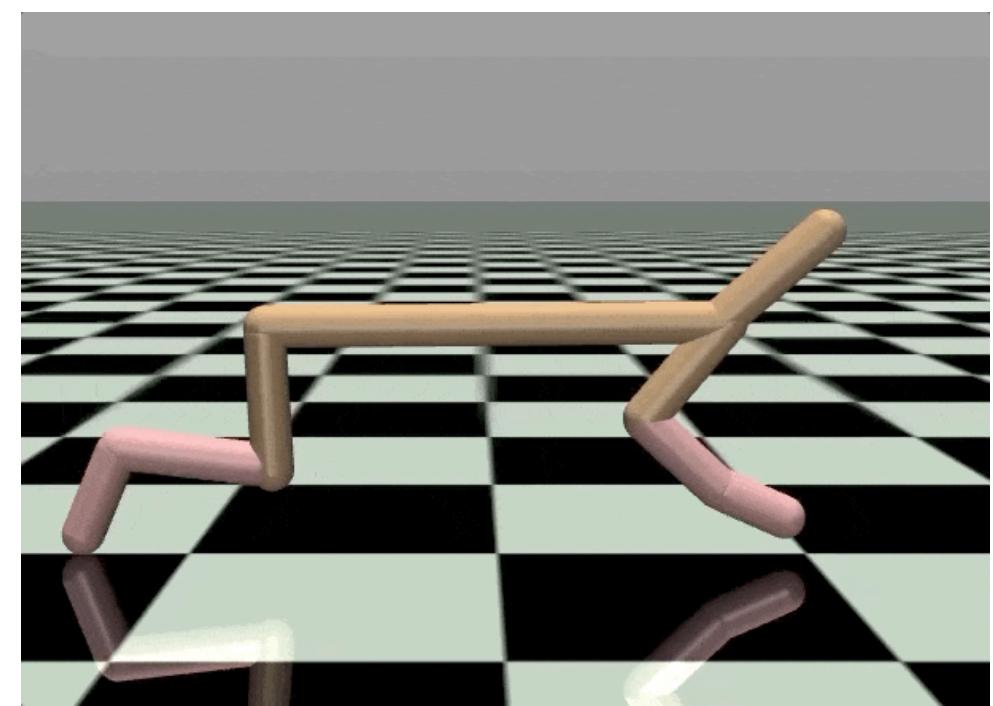
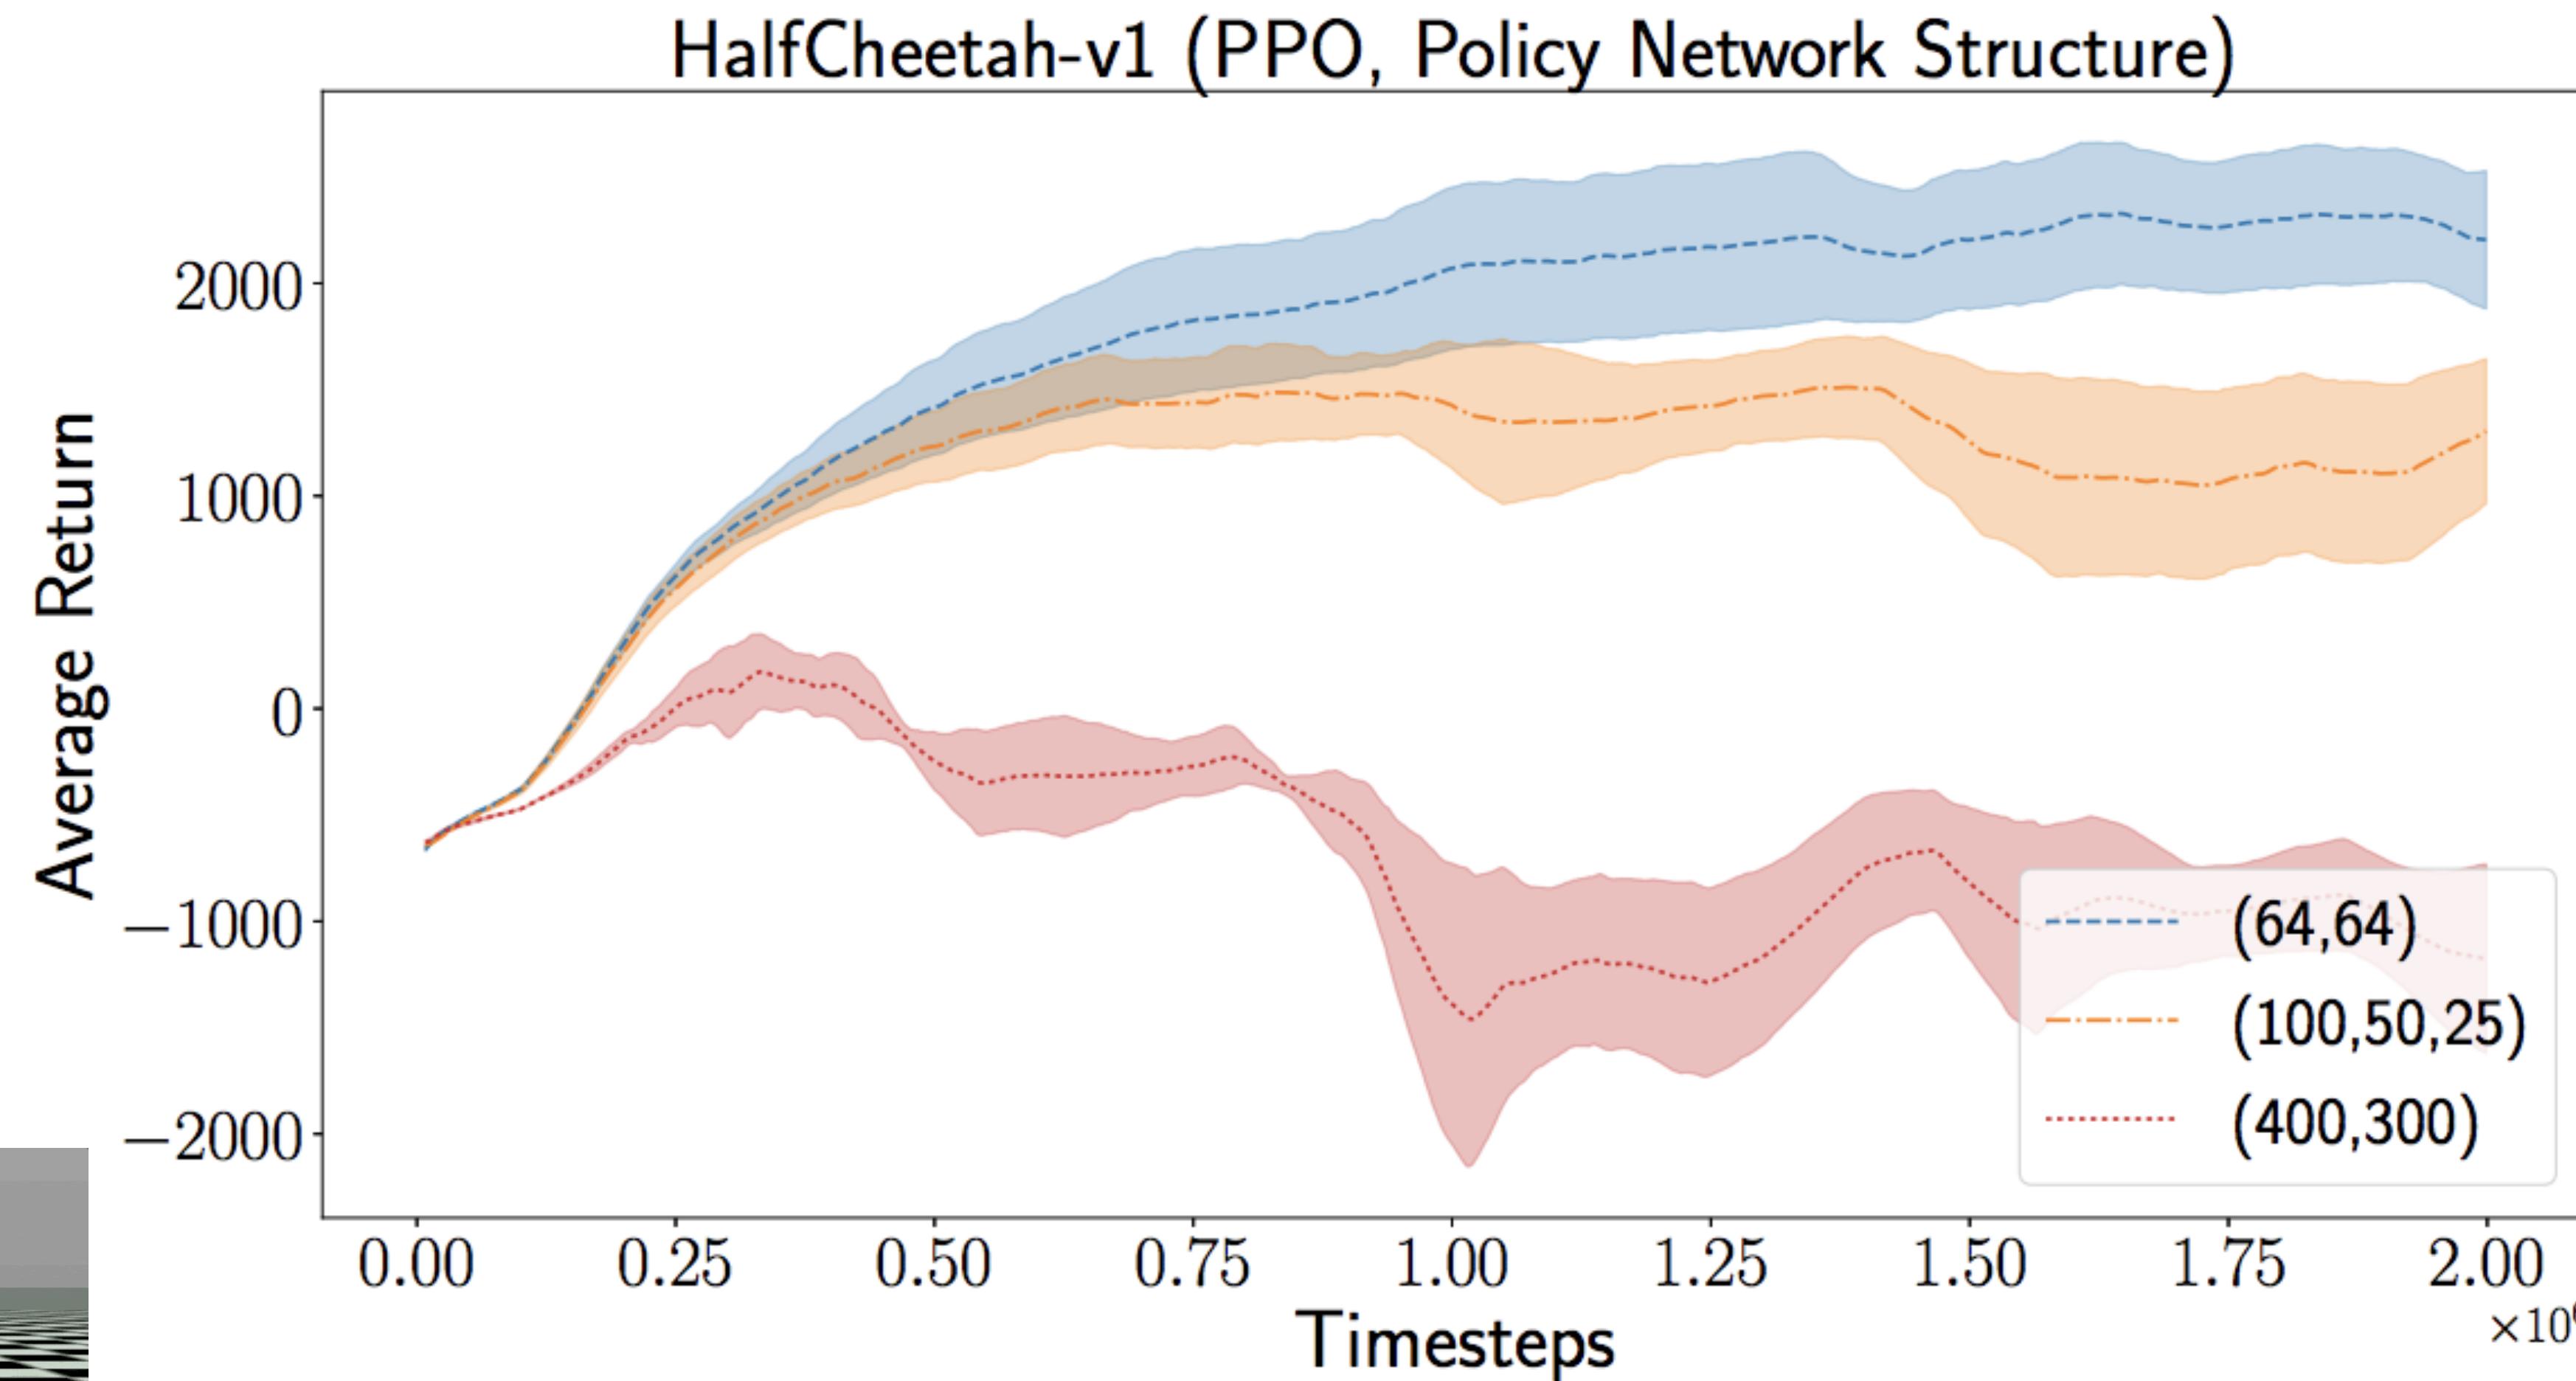


Tune - Hyperparameter Tuning



@deanwampler

Hyperparameters Are Important for Performance



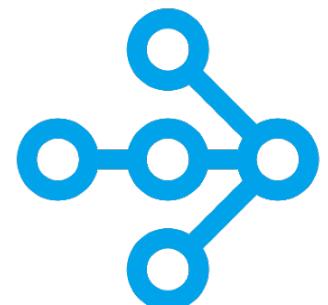
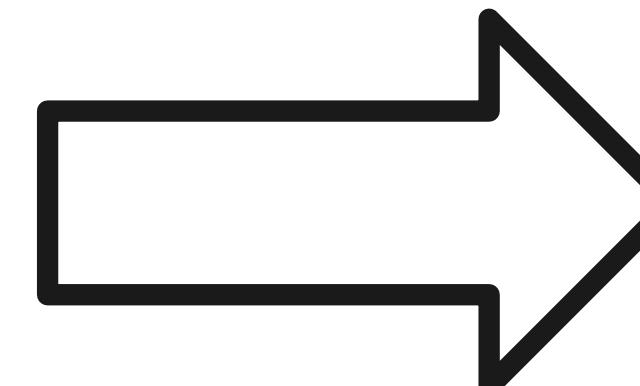
@deanwampler

We Need a Framework for Tuning Hyperparameters

We want the best model

Resources are expensive

Model training is time-consuming

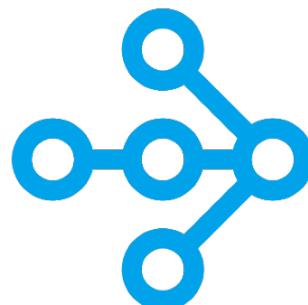
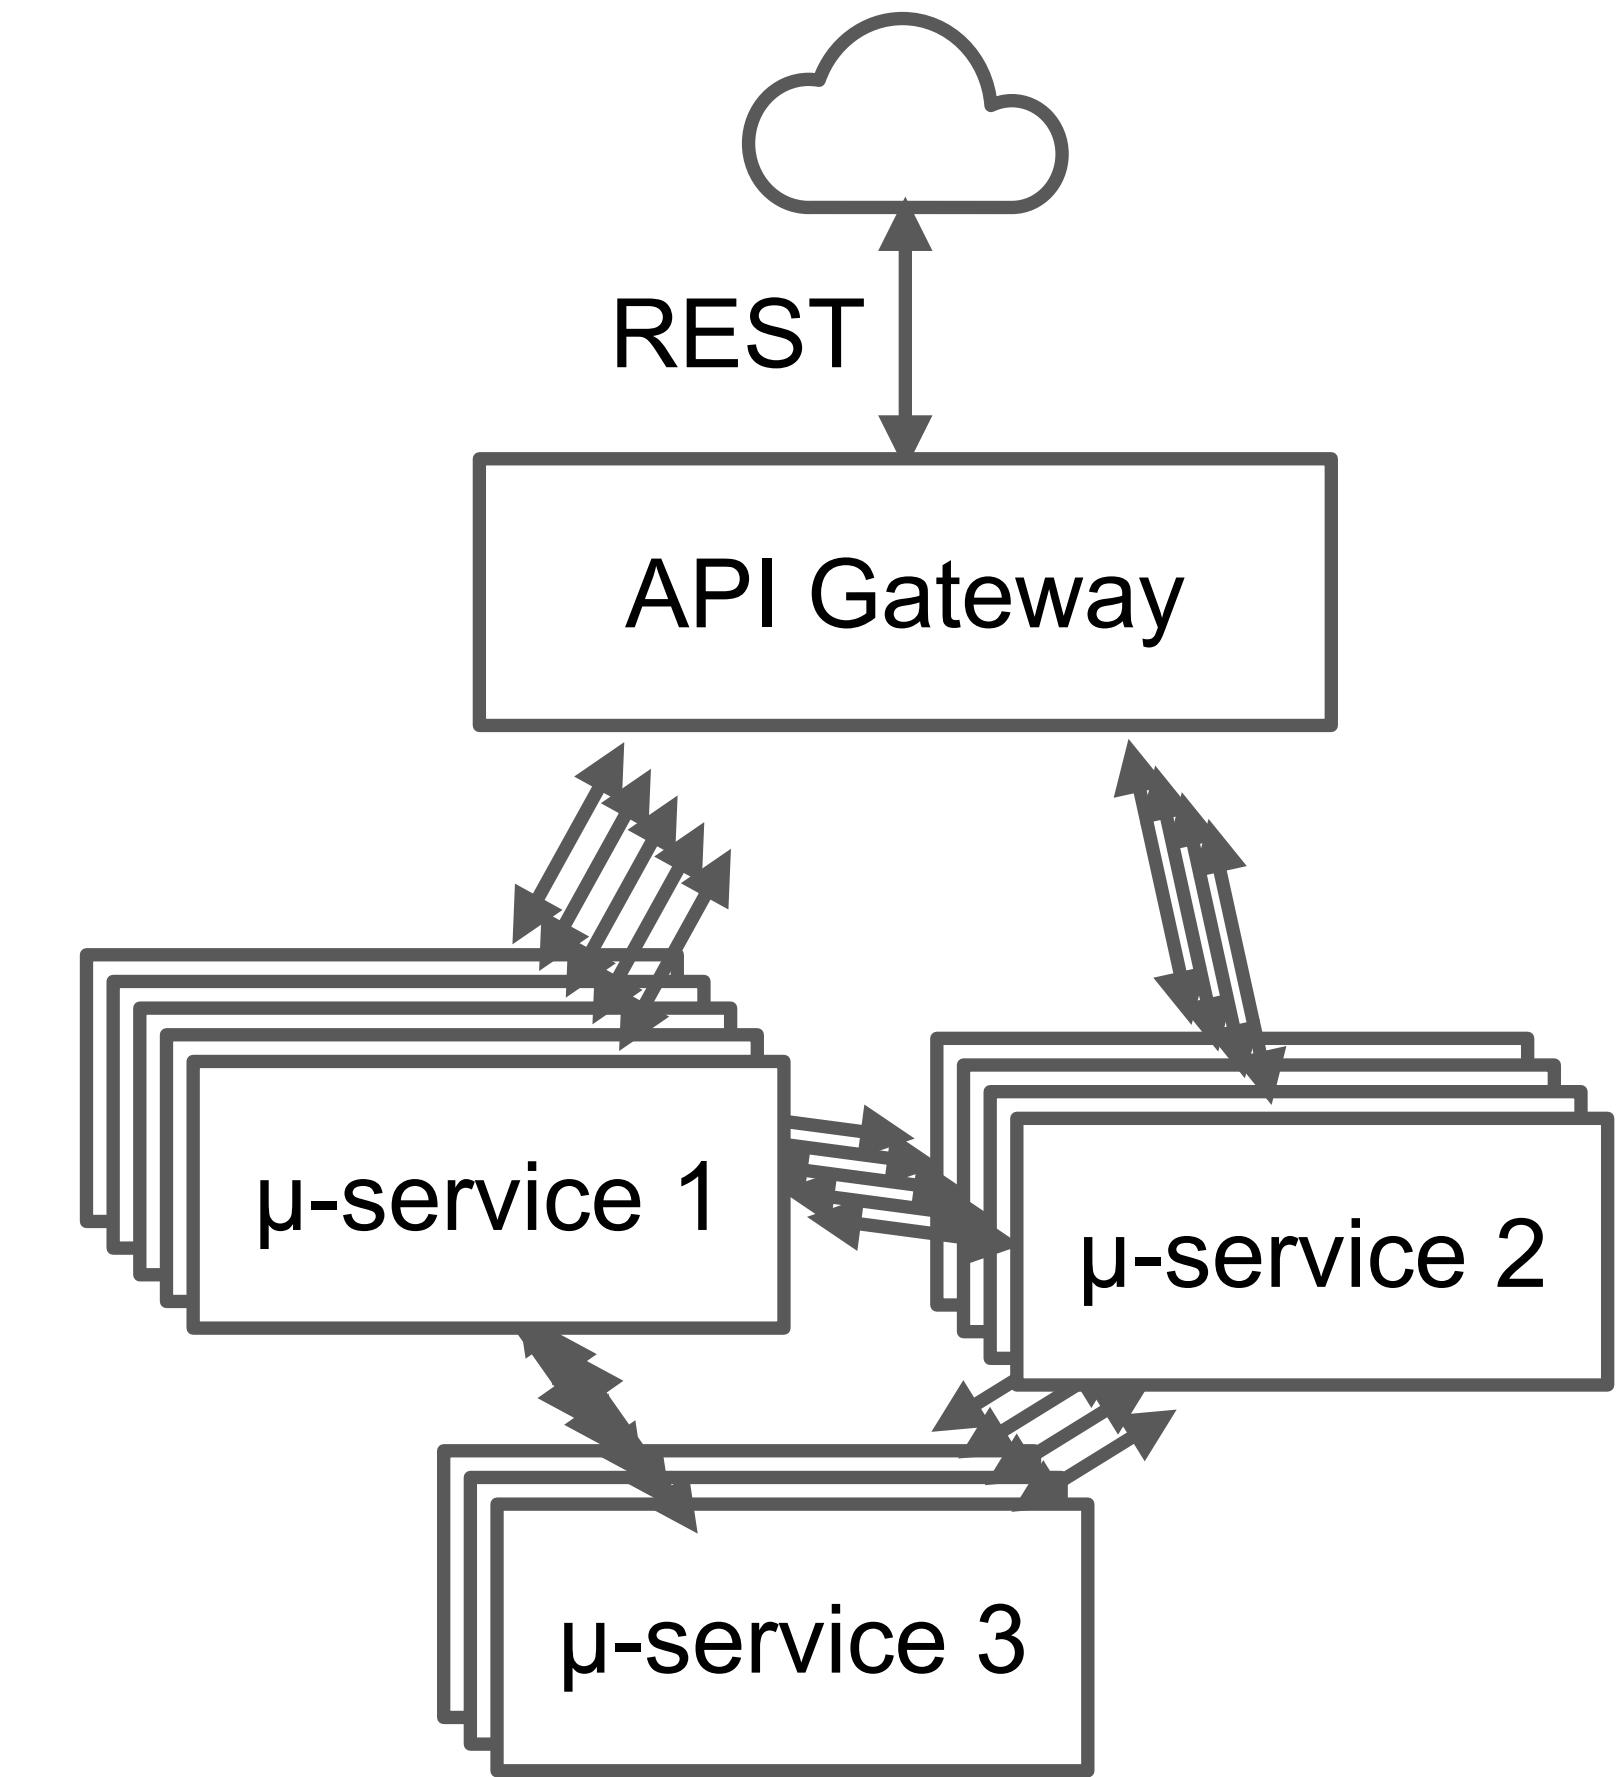




Ray for Microservices?

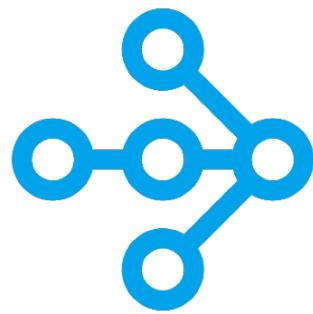
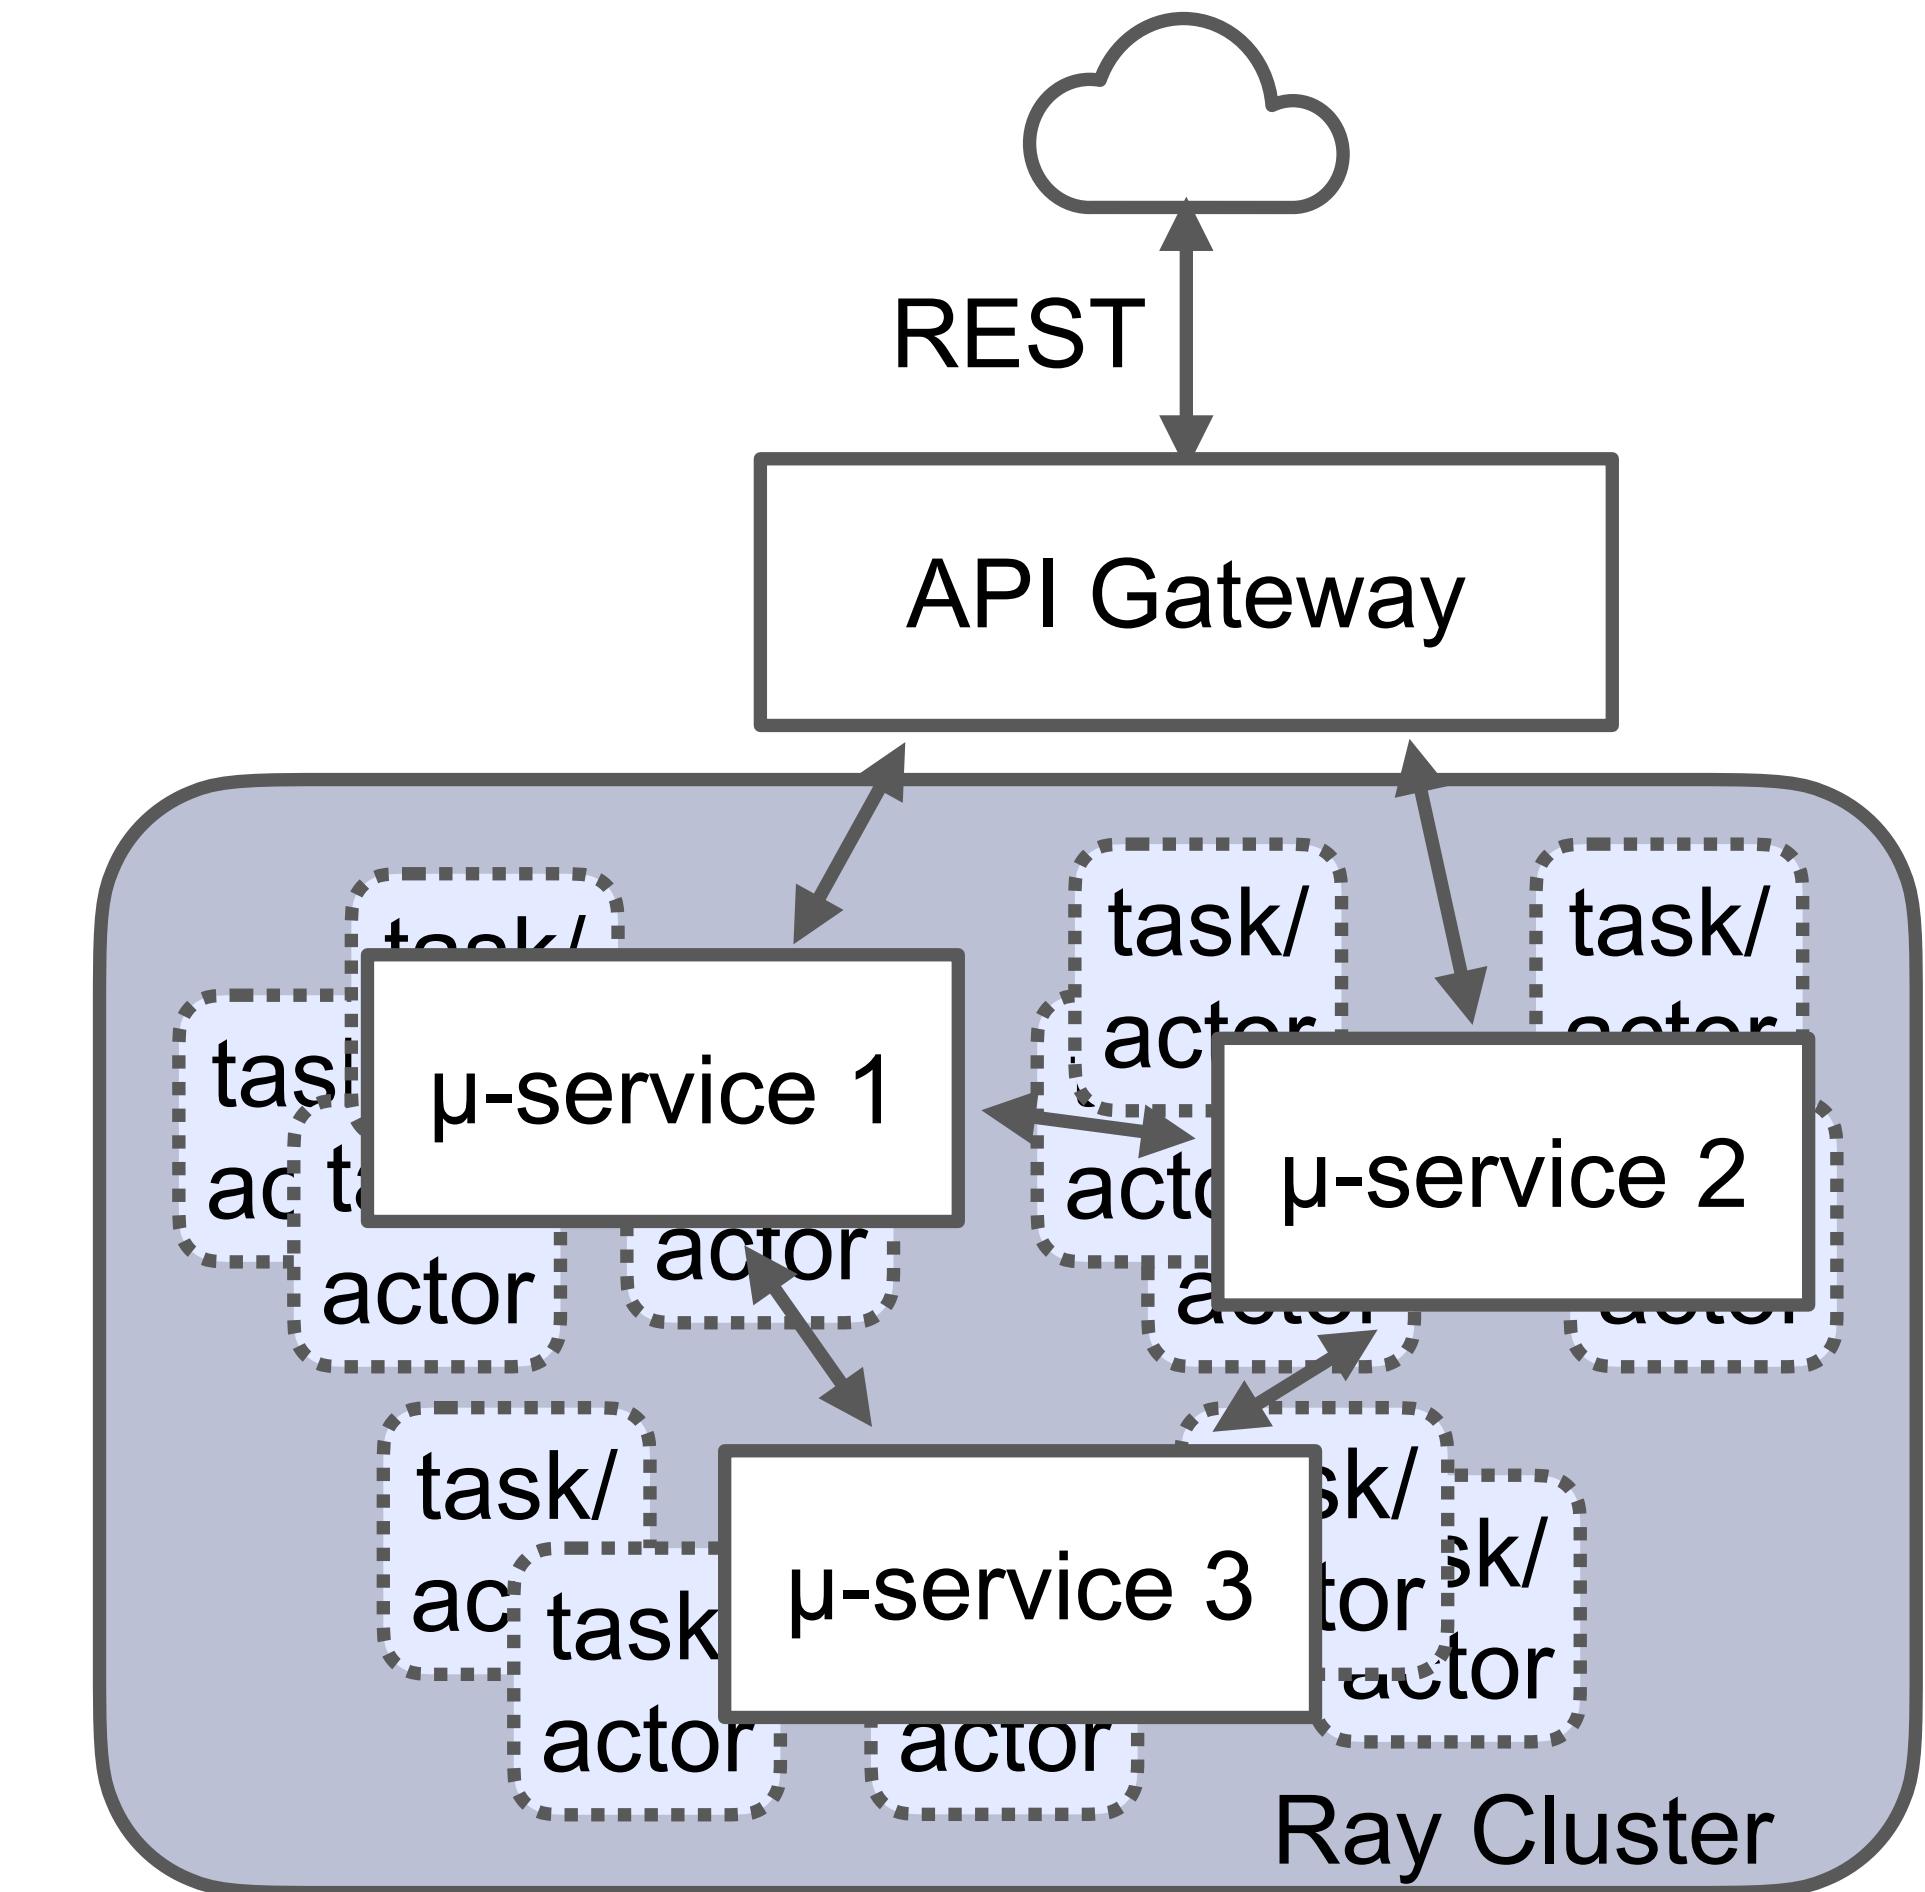
Separate Management

- Each team manages its own instances
- Each microservice has a different number of instances...
- that have to be managed **explicitly**



Simplified Management

- With Ray, you have one “logical” instance to manage and Ray does the cluster-wide scaling for you.



Questions?

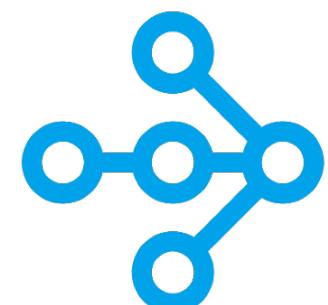
[ray.io](#)

[anyscale.com](#) - We're Hiring!

[anyscale.com/events](#)

[raysummit.org](#)

dean@anyscale.com



@deanwampler



anyscale