





Typesafe

Strata + Hadoop World
San Jose, CA, Feb. 20, 2015
@deanwampler

Why Spark Is
the Next Top
(Compute) Model

Dean Wampler



dean.wampler@typesafe.com
polyglotprogramming.com/talks
@deanwampler

Tuesday, February 17, 15

About me. You can find this presentation and others on Big Data and Scala at polyglotprogramming.com.
Programming Scala, 2nd Edition is forthcoming.
photo: Dusk at 30,000 ft above the Central Plains of the U.S. on a Winter's Day.



Spark is a fast and general engine for large-scale data processing built in Scala

*The Spark logo is the property of the Apache foundation.

[SCROLL DOWN TO LEARN MORE](#)

<http://bit.ly/typesafe-spark>



Spark is a fast and general engine for large-scale
data processing built in Scala

*The Spark logo is the property of the Apache foundation.

SCROLL DOWN TO LEARN MORE


Hadoop circa 2013

Tuesday, February 17, 15

The state of Hadoop as of last year.
Image: Detail of the London Eye




Hadoop v2.X Cluster



Tuesday, February 17, 15

Schematic view of a Hadoop 2 cluster. For a more precise definition of the services and what they do, see e.g., <http://hadoop.apache.org/docs/r2.3.0/hadoop-yarn/hadoop-yarn-site/YARN.html> We aren't interested in great details at this point, but we'll call out a few useful things to know.


Resource and Node Managers



Tuesday, February 17, 15

Hadoop 2 uses YARN to manage resources via the master Resource Manager, which includes a pluggable job scheduler and an Applications Manager. It coordinates with the Node Manager on each node to schedule jobs and provide resources. Other services involved, including application-specific Containers and Application Masters are not shown.

Name Node and Data Nodes



Tuesday, February 17, 15

Hadoop 2 clusters federate the Name node services that manage the file system, HDFS. They provide horizontal scalability of file-system operations and resiliency when service instances fail. The data node services manage individual blocks for files.

MapReduce

The classic compute model
for Hadoop

MapReduce


1 map step + 1 reduce step
(wash, rinse, repeat)

MapReduce

Example:
Inverted Index

Web Crawl

Map Phase



Map Phase

| index | |
|----------------------|--------------------|
| block | |
| ... | ... |
| wikipedia.org/hadoop | Hadoop provides... |
| ... | ... |
| block | |
| ... | ... |
| wikipedia.org/hbase | HBase stores... |
| ... | ... |
| block | |
| ... | ... |
| wikipedia.org/hive | Hive queries... |
| ... | ... |

Map Task

Map Task

Map Task

Sort, Shuffle

Reduce Task

Reduce Task


Reduce Task

Reduce Task

Reduce Phase

Tuesday, February 17, 15

Now we're running MapReduce. In the map step, a task (JVM process) per file *block* (64MB or larger) reads the rows, tokenizes the text and outputs key-value pairs ("tuples")...



Tuesday, February 17, 15

... the keys are each word found and the values are themselves tuples, each URL and the count of the word. In our simplified example, there are typically only single occurrences of each work in each document. The real occurrences are interesting because if a word is mentioned a lot in a document, the chances are higher that you would want to find that document in a search for that word.

Map Phase

| index | |
|----------------------|--------------------|
| block | |
| ... | ... |
| wikipedia.org/hadoop | Hadoop provides... |
| ... | ... |
| block | |
| ... | ... |
| wikipedia.org/hbase | HBase stores... |
| ... | ... |
| block | |
| ... | ... |
| wikipedia.org/hive | Hive queries... |
| ... | ... |

Map Task

Map Task

Map Task

Reduce Task

Reduce Task

Reduce Task

Reduce Task

Sort, Shuffle

Map Phase



Map Task


Map Task

Map Task

Sort, Shuffle

Reduce Task


Reduce Phase



Tuesday, February 17, 15

Finally, each reducer just aggregates all the values it receives for each key, then writes out new files to HDFS with the words and a list of (URL-count) tuples (pairs).

Altogether...



Tuesday, February 17, 15

Finally, each reducer just aggregates all the values it receives for each key, then writes out new files to HDFS with the words and a list of (URL-count) tuples (pairs).



So, what's
not to like?

Tuesday, February 17, 15

This seems okay, right? What's wrong with it?

Awkward

Most algorithms are
much harder to implement
in this restrictive
map-then-reduce model.

Awkward

Lack of flexibility
inhibits optimizations, too.

Performance

Full dump to disk
between jobs.

A photograph of a large Ferris wheel against a backdrop of a cloudy, overcast sky. The Ferris wheel's structure is visible, with its cables and support towers. The perspective is from below, looking up at the wheel.

Enter Spark

spark.apache.org

Cluster Computing

Can be run in:

- YARN (Hadoop 2)
- Mesos (Cluster management)
- EC2
- Standalone mode
- Cassandra
- ...



Tuesday, February 17, 15

If you have a Hadoop cluster, you can run Spark as a seamless compute engine on YARN. (You can also use pre-YARN Hadoop v1 clusters, but there you have manually allocate resources to the embedded Spark cluster vs Hadoop.) Mesos is a general-purpose cluster resource manager that can also be used to manage Hadoop resources. Scripts for running a Spark cluster in EC2 are available. Standalone just means you run Spark's built-in support for clustering (or run locally on a single box - e.g., for development). EC2 deployments are usually standalone... Finally, database vendors like Datastax are integrating Spark.

Compute Model

Fine-grained *operators*
for composing algorithms.



Compute Model

RDDs:


Resilient, Distributed Datasets



Tuesday, February 17, 15

RDDs shard the data over a cluster, like a virtualized, distributed collection (analogous to HDFS). They support intelligent caching, which means no naive flushes of massive datasets to disk. This feature alone allows Spark jobs to run 10-100x faster than comparable MapReduce jobs! The “resilient” part means they will reconstitute shards lost due to process/server crashes.

Compute Model



Spark


Tuesday, February 17, 15

RDDs shard the data over a cluster, like a virtualized, distributed collection (analogous to HDFS). They support intelligent caching, which means no naive flushes of massive datasets to disk. This feature alone allows Spark jobs to run 10-100x faster than comparable MapReduce jobs! The “resilient” part means they will reconstitute shards lost due to process/server crashes.

Compute Model

Written in Scala,
with Java and Python APIs.





Inverted Index in MapReduce (Java).

Tuesday, February 17, 15

Let's see an actual implementation of the inverted index. First, a Hadoop MapReduce (Java) version, adapted from <https://developer.yahoo.com/hadoop/tutorial/module4.html#solution> It's about 90 lines of code, but I reformatted to fit better.

This is also a slightly simpler version than the one I diagrammed. It doesn't record a count of each word in a document; it just writes (word,doc-title) pairs out of the mappers and the final (word,list) output by the reducers just has a list of documentations, hence repeats. A second job would be necessary to count the repeats.

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LineIndexer {

    public static void main(String[] args) {
        JobClient client = new JobClient();
        JobConf conf =
            new JobConf(LineIndexer.class);

        conf.setJobName("LineIndexer");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(Text.class);
    }
}
```

Tuesday, February 17, 15

I've shortened the original code a bit, e.g., using * import statements instead of separate imports for each class.

I'm not going to explain every line ... nor most lines.

Everything is in one outer class. We start with a main routine that sets up the job. Lotta boilerplate...

I used yellow for method calls, because methods do the real work!! But notice that the functions in this code don't really do a whole lot...

```
JobConf conf =  
    new JobConf(LineIndexer.class);  
  
conf.setJobName("LineIndexer");  
conf.setOutputKeyClass(Text.class);  
conf.setOutputValueClass(Text.class);  
FileInputFormat.addInputPath(conf,  
    new Path("input"));  
FileOutputFormat.setOutputPath(conf,  
    new Path("output"));  
conf.setMapperClass(  
    LineIndexMapper.class);  
conf.setReducerClass(  
    LineIndexReducer.class);  
  
client.setConf(conf);  
  
try {
```

```
cont.setReducerClass(  
    LineIndexReducer.class);  
  
client.setConf(conf);  
  
try {  
    JobClient.runJob(conf);  
} catch (Exception e) {  
    e.printStackTrace();  
}  
}  
}  
  
public static class LineIndexMapper  
extends MapReduceBase  
implements Mapper<LongWritable, Text,  
           Text, Text> {  
    private final static Text word =  
        new Text();  
    . . . . .
```

```
extends MapReduceBase
implements Mapper<LongWritable, Text,
           Text, Text> {
private final static Text word =
    new Text();
private final static Text location =
    new Text();

public void map(
    LongWritable key, Text val,
    OutputCollector<Text, Text> output,
    Reporter reporter) throws IOException {

    FileSplit fileSplit =
        (FileSplit)reporter.getInputSplit();
    String fileName =
        fileSplit.getPath().getName();
    location.set(fileName);
```

```
(FileSplit)reporter.getInputSplit(),  
String fileName =  
    fileSplit.getPath().getName();  
location.set(fileName);  
  
String line = val.toString();  
StringTokenizer itr = new  
    StringTokenizer(line.toLowerCase());  
while (itr.hasMoreTokens()) {  
    word.set(itr.nextToken());  
    output.collect(word, location);  
}  
}  
}  
}  
  
public static class LineIndexReducer  
extends MapReduceBase  
implements Reducer<Text, Text,  
                    Text, Text> {
```

```
public void reduce(Text key,
  Iterator<Text> values,
  OutputCollector<Text, Text> output,
  Reporter reporter) throws IOException {
  boolean first = true;
  StringBuilder toReturn =
    new StringBuilder();
  while (values.hasNext()) {
    if (!first)
      toReturn.append(", ");
    first=false;
    toReturn.append(
      values.next().toString());
  }
  output.collect(key,
    new Text(toReturn.toString()));
}
```

Tuesday, February 17, 15

The reducer class, LineIndexReducer, with the reduce method that is called for each key and a list of values for that key. The reducer is stupid; it just reformats the values collection into a long string and writes the final (word,list-string) output.

```
        if (!first)
            toReturn.append(", ");
        first=false;
        toReturn.append(
            values.next().toString());
    }
    output.collect(key,
        new Text(toReturn.toString()));
}
}
```

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LineIndexer {

    public static void main(String[] args) {
        JobClient client = new JobClient();
        JobConf conf =
            new JobConf(LineIndexer.class);

        conf.setJobName("LineIndexer");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(Text.class);
        FileInputFormat.addInputPath(conf,
            new Path("input"));
        FileOutputFormat.setOutputPath(conf,
            new Path("output"));
        conf.setMapperClass(
            LineIndexMapper.class);
        conf.setReducerClass(
            LineIndexReducer.class);

        client.setConf(conf);

        try {
            JobClient.runJob(conf);
        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    public static class LineIndexMapper
        extends MapReduceBase
        implements Mapper<LongWritable, Text,
                    Text, Text> {
        private final static Text word =
            new Text();
        private final static Text location =
            new Text();

        public void map(
            LongWritable key, Text val,
            OutputCollector<Text, Text> output,
            Reporter reporter) throws IOException {

            FileSplit fileSplit =
                (FileSplit)reporter.getInputSplit();
            String fileName =
                fileSplit.getPath().getName();
            location.set(fileName);

            String line = val.toString();
            StringTokenizer itr = new
                StringTokenizer(line.toLowerCase());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                output.collect(word, location);
            }
        }
    }

    public static class LineIndexReducer
        extends MapReduceBase
        implements Reducer<Text, Text,
                           Text, Text> {
        public void reduce(Text key,
                          Iterator<Text> values,
                          OutputCollector<Text, Text> output,
                          Reporter reporter) throws IOException {
            boolean first = true;
            StringBuilder toReturn =
                new StringBuilder();
            while (values.hasNext()) {
                if (!first)
                    toReturn.append(", ");
                first=false;
                toReturn.append(
                    values.next().toString());
            }
            output.collect(key,
                new Text(toReturn.toString()));
        }
    }
}
```

Altogether

Tuesday, February 17, 15

The whole shebang (6pt. font)



Inverted Index in Spark (Scala).

Tuesday, February 17, 15

This code is approximately 45 lines, but it does more than the previous Java example, it implements the original inverted index algorithm I diagrammed where word counts are computed and included in the data.

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
        text.split("""\W+""") map {
```

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
        text.split("""\W+""") map {
```

Tuesday, February 17, 15

It starts with imports, then declares a singleton object (a first-class concept in Scala), with a main routine (as in Java).
The methods are colored yellow again. Note this time how dense with meaning they are this time.

```
def main(args: Array[String]) = {
```

```
    val sc = new SparkContext(  
        "local", "Inverted Index")
```

```
    sc.textFile("data/crawl")  
        .map { line =>  
            val array = line.split("\t", 2)  
            (array(0), array(1))  
        }  
        .flatMap {  
            case (path, text) =>  
                text.split("""\W+""") map {  
                    word => (word, path)  
                }  
        }  
        .map {  
            case (w, p) => ((w, p), 1)
```

Tuesday, February 17, 15

You begin the workflow by declaring a `SparkContext`. We're running in "local" mode, in this case, meaning on a single machine (and using a single core). Normally this argument would be a command-line parameter, so you can develop locally, then submit to a cluster, where "local" would be replaced by the appropriate URI.

```
def main(args: Array[String]) = {  
  
    val sc = new SparkContext(  
        "local", "Inverted Index")  
  
    sc.textFile("data/crawl")  
    .map { line =>  
        val array = line.split("\t", 2)  
        (array(0), array(1))  
    }  
    .flatMap {  
        case (path, text) =>  
            text.split("""\W+""") map {  
                word => (word, path)  
            }  
    }  
    .map {  
        case (w, p) => ((w, p), 1)  
    }  
}
```

Tuesday, February 17, 15

The rest of the program is a sequence of function calls, analogous to “pipes” we connect together to construct the data flow. Data will only start “flowing” when we ask for results. We start by reading one or more text files from the directory “data/crawl”. If running in Hadoop, if there are one or more Hadoop-style “part-NNNNN” files, Spark will process all of them (they will be processed synchronously in “local” mode).

```
sc.textFile("data/crawl")
  .map { line =>
    val array = line.split("\t", 2)
    (array(0), array(1))
  }
  .flatMap {
    case (path, text) =>
    text.split("""\W+""") map {
      word => (word, path)
    }
  }
  .map {
    case (w, p) => ((w, p), 1)
  }
  .reduceByKey {
    case (n1, n2) => n1 + n2
  }
  .map {
```

Tuesday, February 17, 15

sc.textFile returns an RDD with a string for each line of input text. So, the first thing we do is map over these strings to extract the original document id (i.e., file name), followed by the text in the document, all on one line. We assume tab is the separator. "(array(0), array(1))" returns a two-element "tuple". Think of the output RDD has having a schema "fileName: String, text: String".

```
sc.textFile("data/crawl")
  .map { line =>
    val array = line.split("\t", 2)
    (array(0), array(1))
  }
  .flatMap {
    case (path, text) =>
      text.split("""\W+""") map {
        word => (word, path)
      }
  }
  .map {
    case (w, p) => ((w, p), 1)
  }
  .reduceByKey {
    case (n1, n2) => n1 + n2
  }
  .map {
```

Tuesday, February 17, 15

flatMap maps over each of these 2-element tuples. We split the text into words on non-alphanumeric characters, then output collections of word (our ultimate, final “key”) and the path. That is, each line (one thing) is converted to a collection of (word,path) pairs (0 to many things), but we don’t want an output collection of nested collections, so flatMap concatenates nested collections into one long “flat” collection of (word,path) pairs.

```
.map {  
    case (w, p) => ((w, p), 1)  
}  
.reduceByKey {  
    case (n1, n2) => n1 + n2  
}  
  
.map {  
    case ((w, p), n) => (w, (p, n))  
}  
.groupBy {  
    case (w, (p, n)) => w  
}  
.map {  
    case (w, seq) =>  
        val seq2 = seq map {  
            case (_, (p, n)) => (p, n)  
        }  
        sortBy {  
            case (p, n) => (n, p)  
        }  
}
```

((word1, path1), n1)
((word2, path2), n2)
...

Tuesday, February 17, 15

Next, we map over these pairs and add a single “seed” count of 1. Note the structure of the returned tuple; it’s a two-tuple where the first element is itself a two-tuple holding (word, path). The following special method, reduceByKey is like a groupBy, where it groups over those (word, path) “keys” and uses the function to sum the integers. The popup shows what the output data looks like.

```
.map {  
    case (w, p) => ((w, p), 1)  
}  
.reduceByKey {  
    case (n1, n2) => n1 + n2  
}  
.map {  
    case ((w, p), n) => (w, (p, n))  
}  
.groupBy {  
    case (w, (p, n)) => w  
}  
.map {  
    case (w, seq) =>  
        val seq2 = seq map {  
            case (_, (p, n)) => (p, n)  
        }  
        sortBy {  
            case (p, n) => -n  
        }  
}
```

(word1, (path1, n1))
(word2, (path2, n2))
...

```
}

    .groupBy {
        case (w, (p, n)) => w
    }                                (word, seq((word, (path1, n1)), (word, (path2, n2)), ...))

    .map {
        case (w, seq) =>
            val seq2 = seq map {
                case (_, (p, n)) => (p, n)
            }
            .sortBy {
                case (path, n) => (-n, path)
            }
            (w, seq2.mkString("", ""))
    }

    .saveAsTextFile("/path/to/out")

    sc.stop()
}
```

```
}

    .groupBy {
        case (w, (p, n)) => w
    }

    .map {
        case (w, seq) =>
            val seq2 = seq map {
                case (_, (p, n)) => (p, n)
            }
            .sortBy {
                case (path, n) => (-n, path)
            }
            (w, seq2.mkString(", "))
    }

    .saveAsTextFile("/path/to/out")

    sc.stop()
}
```

(word, "(path1, n1), (path2, n2), ...")

Now we do an explicit group by to bring all the same words together. The output will be (word, (word, (path1, n1)), (word, (path2, n2)), ...).

The last map removes the redundant "word" values in the sequences of the previous output and sorts by count descending, path ascending. (Sorting by path is mostly useful for reproducibility, e.g., in tests!) It outputs the sequence as a final string of comma-separated (path,n) pairs.

```
val seq2 = seq map {
    case (_, (p, n)) => (p, n)
}
.sortBy {
    case (path, n) => (-n, path)
}
(w, seq2.mkString(", "))

}

.saveAsTextFile("/path/to/out")

sc.stop()
}
```

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
          text.split("""\w+""") map {
            word => (word, path)
          }
      }
      .map {
        case (w, p) => ((w, p), 1)
      }
      .reduceByKey {
        case (n1, n2) => n1 + n2
      }
      .map {
        case ((w, p), n) => (w, (p, n))
      }
      .groupByKey {
        case (w, (p, n)) => w
      }
      .map {
        case (w, seq) =>
          val seq2 = seq map {
            case (_, (p, n)) => (p, n)
          }
          (w, seq2.mkString(", "))
      }
      .saveAsTextFile("/path/to/out")

    sc.stop()
  }
}
```

Altogether

```
.map { line =>
  val array = line.split("\t", 2)
  (array(0), array(1))
}
.flatMap {
  case (path, text) =>
  text.split("""\W+""") map {
    word => (word, path)
  }
}
.map {
  case (w, p) => ((w, p), 1)
}
.reduceByKey {
  case (n1, n2) => n1 + n2
}
.map {
  case ((w, p), n) => (w, (p, n))
```

*Concise
Operators!*

$$\nabla \cdot \mathbf{D} = \rho$$


$$\nabla \cdot \mathbf{B} = 0$$

$$\nabla \times \mathbf{E} = -\frac{\partial \mathbf{B}}{\partial t}$$

$$\nabla \times \mathbf{H} = \mathbf{J} + \frac{\partial \mathbf{D}}{\partial t}$$

Tuesday, February 17, 15

Another example of a beautiful and profound DSL, in this case from the world of Physics: Maxwell's equations: <http://upload.wikimedia.org/wikipedia/commons/c/c4/Maxwell'sEquations.svg>

The background of the slide features a large satellite dish antenna. The dish itself is a light grey parabolic shape, and the supporting structure is a complex network of dark grey metal beams and cables forming a truss. The dish is mounted on a tall vertical pole. The sky above the dish is a clear blue.

*The Spark version took me
~30 minutes to write.*

Tuesday, February 17, 15

Once you learn the core primitives I used, and a few tricks for manipulating the RDD tuples, you can very quickly build complex algorithms for data processing! The Spark API allowed us to focus almost exclusively on the “domain” of data transformations, while the Java MapReduce version (which does less), forced tedious attention to infrastructure mechanics.



*Use a SQL query
when you can!!*

Spark SQL!

Mix SQL queries with
the RDD API.



Spark SQL!

Create, Read, and Delete Hive Tables



Spark SQL!

Read JSON and
Infer the Schema



Spark SQL!

Read and write
Parquet files




Tuesday, February 17, 15

Parquet is a newer file format developed by Twitter and Cloudera that is becoming very popular. It stores in column order, which is better than row order when you have lots of columns and your queries only need a few of them. Also, columns of the same data types are easier to compress, which Parquet does for you. Finally, Parquet files carry the data schema.

SparkSQL

~10-100x
the performance
of Hive.





Combine SparkSQL with Machine Learning.

Tuesday, February 17, 15

We'll use the Spark "MLlib" in the example, then return to it in a moment.

```
CREATE TABLE Users(  
    userId      INT,  
    name        STRING,  
    email       STRING,  
    age         INT,  
    latitude    DOUBLE,  
    longitude   DOUBLE,  
    subscribed  BOOLEAN);
```

```
CREATE TABLE Events(  
    userId INT,  
    action INT);
```

Equivalent HiveQL Schemas definitions.

Tuesday, February 17, 15

This example adapted from the following blog post announcing Spark SQL:

<http://databricks.com/blog/2014/03/26/Spark-SQL-manipulating-structured-data-using-Spark.html>

Adapted here to use Spark's own SQL, not the integration with Hive. Imagine we have a stream of events from users and the events that have occurred as they used a system.

```
val trainingDataTable = sql("""  
    SELECT e.action, u.age,  
          u.latitude, u.longitude  
    FROM Users u  
  JOIN Events e  
  ON u.userId = e.userId""")
```

```
val trainingData =  
  trainingDataTable map { row =>  
    val features =  
      Array[Double](row(1), row(2), row(3))  
    LabeledPoint(row(0), features)  
  }  
  
val model =  
  new LogisticRegressionWithSGD()  
  .run(trainingData)
```

```
val trainingDataTable = sql("""  
    SELECT e.action, u.age,  
          u.latitude, u.longitude  
    FROM Users u  
    JOIN Events e  
    ON u.userId = e.userId""")
```

```
val trainingData =  
  trainingDataTable map { row =>  
    val features =  
      Array[Double](row(1), row(2), row(3))  
    LabeledPoint(row(0), features)  
  }
```

```
val model =  
  new LogisticRegressionWithSGD()  
  .run(trainingData)
```

```
val model =  
  new LogisticRegressionWithSGD()  
  .run(trainingData)
```

```
val allCandidates = sql("""  
  SELECT userId, age, latitude, longitude  
  FROM Users  
  WHERE subscribed = FALSE""")  
  
case class Score(  
  userId: Int, score: Double)  
val scores = allCandidates map { row =>  
  val features =  
    Array[Double](row(1), row(2), row(3))  
  Score(row(0), model.predict(features))  
}  
  
// In-memory table
```

Tuesday, February 17, 15

Next we train the recommendation engine, using a “logistic regression” fit to the training data, where “stochastic gradient descent” (SGD) is used to train it. (This is a standard tool set for recommendation engines; see for example: <http://www.cs.cmu.edu/~wcohen/10-605/assignments/sgd.pdf>)

```
val model =  
  new LogisticRegressionWithSGD()  
.run(trainingData)
```

```
val allCandidates = sql("""  
SELECT userId, age, latitude, longitude  
FROM Users  
WHERE subscribed = FALSE""")
```

```
case class Score(  
  userId: Int, score: Double)  
val scores = allCandidates map { row =>  
  val features =  
    Array[Double](row(1), row(2), row(3))  
  Score(row(0), model.predict(features))  
}
```

```
// In-memory table
```

Tuesday, February 17, 15

Now run a query against all users who aren't already subscribed to notifications.

```
case class Score(  
  userId: Int, score: Double)  
val scores = allCandidates map { row =>  
  val features =  
    Array[Double](row(1), row(2), row(3))  
  Score(row(0), model.predict(features))  
}
```

```
// In-memory table  
scores.registerTempTable("Scores")
```

```
val topCandidates = sql("""  
SELECT u.name, u.email  
FROM Scores s  
JOIN Users u ON s.userId = u.userId  
ORDER BY score DESC  
LIMIT 100""")
```

```
case class Score(  
  userId: Int, score: Double)  
val scores = allCandidates map { row =>  
  val features =  
    Array[Double](row(1), row(2), row(3))  
  Score(row(0), model.predict(features))  
}
```

```
// In-memory table  
scores.registerTempTable("Scores")
```

```
val topCandidates = sql("""  
SELECT u.name, u.email  
FROM Scores s  
JOIN Users u ON s.userId = u.userId  
ORDER BY score DESC  
LIMIT 100""")
```

```
// In-memory table  
scores.registerTempTable("Scores")
```

```
val topCandidates = sql("""  
SELECT u.name, u.email  
FROM Scores s  
JOIN Users u ON s.userId = u.userId  
ORDER BY score DESC  
LIMIT 100""")
```

```
val trainingDataTable = sql("""  
    SELECT e.action, u.age,  
          u.latitude, u.longitude  
    FROM Users u  
    JOIN Events e  
    ON u.userId = e.userId""")  
  
val trainingData =  
  trainingDataTable map { row =>  
    val features =  
      Array[Double](row(1), row(2), row(3))  
    LabeledPoint(row(0), features)  
  }  
  
val model =  
  new LogisticRegressionWithSGD()  
  .run(trainingData)  
  
val allCandidates = sql("""  
    SELECT userId, age, latitude, longitude  
    FROM Users  
    WHERE subscribed = FALSE""")  
  
case class Score(  
  userId: Int, score: Double)  
val scores = allCandidates map { row =>  
  val features =  
    Array[Double](row(1), row(2), row(3))  
  Score(row(0), model.predict(features))  
}  
  
// In-memory table  
scores.registerTempTable("Scores")  
  
val topCandidates = sql("""  
    SELECT u.name, u.email  
    FROM Scores s  
    JOIN Users u ON s.userId = u.userId  
    ORDER BY score DESC  
    LIMIT 100""")
```

Altogether



Event Stream Processing

Spark Streaming

Use the same abstractions
for near real-time,
event streaming.






Window of 3 RDD Batches #1

Window of 3 RDD Batches #2



Tuesday, February 17, 15

A DStream (discretized stream) wraps the RDDs for each “batch” of events. You can specify the granularity, such as all events in 1 second batches, then your Spark job is passed each batch of data for processing. You can also work with moving windows of batches.



Very similar code...

```
val sc = new SparkContext(...)  
val ssc = new StreamingContext(  
    sc, Seconds(1))  
  
// A DStream that will listen to server:port  
val lines =  
    ssc.socketTextStream(server, port)  
  
// Word Count...  
val words = lines flatMap {  
    line => line.split("""\w+""")  
}  
  
val pairs = words map (word => (word, 1))  
val wordCounts =  
    pairs reduceByKey ((n1, n2) => n1 + n2)
```

```
val sc = new SparkContext(...)  
val ssc = new StreamingContext(  
    sc, Seconds(1))
```

```
// A DStream that will listen to server:port  
val lines =  
    ssc.socketTextStream(server, port)  
  
// Word Count...  
val words = lines flatMap {  
    line => line.split("""\w+""")  
}  
  
val pairs = words map (word => (word, 1))  
val wordCounts =  
    pairs reduceByKey ((n1, n2) => n1 + n2)
```

```
val sc = new SparkContext(...)  
val ssc = new StreamingContext(  
    sc, Seconds(1))
```

```
// A DStream that will listen to server:port  
val lines =  
    ssc.socketTextStream(server, port)
```

```
// Word Count...  
val words = lines flatMap {  
    line => line.split("""\w+""")  
}
```

```
val pairs = words map (word => (word, 1))  
val wordCounts =  
    pairs reduceByKey ((n1, n2) => n1 + n2)
```

```
ssc.socketTextStream(server, port)
```

```
// Word Count...
val words = lines flatMap {
  line => line.split("""\w+""")  
}
```

```
val pairs = words map (word => (word, 1))
val wordCounts =
  pairs reduceByKey ((n1, n2) => n1 + n2)
```

```
wordCount.print() // print a few counts...
```

```
ssc.start()
ssc.awaitTermination()
```

Tuesday, February 17, 15

Now we “count words”. For each mini-batch (1 second’s worth of data), we split the input text into words (on whitespace, which is too crude).

Once we setup the flow, we start it and wait for it to terminate through some means, such as the server socket closing.

```
ssc.socketTextStream(server, port)
```

```
// Word Count...
val words = lines flatMap {
  line => line.split("""\w+""")  
}
```

```
val pairs = words map (word => (word, 1))
val wordCounts =
  pairs reduceByKey ((n1, n2) => n1 + n2)
```

```
wordCount.print() // print a few counts...
```

```
ssc.start()
ssc.awaitTermination()
```

```
ssc.socketTextStream(server, port)
```

```
// Word Count...
val words = lines flatMap {
  line => line.split("""\w+""")
}

val pairs = words map (word => (word, 1))
val wordCounts =
  pairs reduceByKey ((n1, n2) => n1 + n2)

wordCount.print() // print a few counts...
```

```
ssc.start()
ssc.awaitTermination()
```

```
ssc.socketTextStream(server, port)
```

```
// Word Count...
val words = lines flatMap {
  line => line.split("""\w+""")
}

val pairs = words map (word => (word, 1))
val wordCounts =
  pairs reduceByKey ((n1, n2) => n1 + n2)

wordCount.print() // print a few counts...
```

```
ssc.start()
ssc.awaitTermination()
```

Big Data circa 2020



Tuesday, February 17, 15

So where will we be five years from now?

MapReduce vs. Spark



Tuesday, February 17, 15

Spark just replaced MapReduce

YARN? vs. Mesos



Tuesday, February 17, 15

What about YARN. It's somewhat specific to the MapReduce model (batch mode, finite-duration jobs, somewhat static allocation of resources for job life). It's less "universal" and efficient compared to Mesos. As Data environments grow more sophisticated, I believe YARN will reach a point where we need to replace it. Mesos is the most likely contender.

HDFS

vs. ?



Tuesday, February 17, 15

As a distributed file system layered on top of a native filesystem, HDFS is not nearly as efficient as it could be. Its resiliency features are a hack. It fares poorly with small or incrementally-updated files. A distributed file system with better performance, resiliency, and efficiency for a wider variety of scenarios will become essential. Possible replacements are MapR-FS, Ceph, Gluster, and others(?)



Recap



Dean Wampler

@deanwampler

Functional Programming: I came for the concurrency, but I stayed for the data science.

Reply Delete ★ Favorite ... More

RETWEETS

6

FAVORITES

5



Tuesday, February 17, 15

Why is Spark so good (and Java MapReduce so bad)? Because fundamentally, data analytics is Mathematics and programming tools inspired by Mathematics - like Functional Programming - are ideal tools for working with data. This is why Spark code is so concise, yet powerful. This is why it is a great platform for performance optimizations. This is why Spark is a great platform for higher-level tools, like SQL, graphs, etc.

Interest in FP started growing ~10 years ago as a tool to attack concurrency. I believe that data is now driving FP adoption even faster. I know many Java shops that switched to Scala when they adopted tools like Spark and Scalding (<https://github.com/twitter/scalding>).

Spark

A flexible, scalable distributed compute
platform with concise, powerful APIs
and higher-order tools.
spark.apache.org



Why Spark Is the Next Top (Compute) Model

@deanwampler

dean.wampler@typesafe.com

polyglotprogramming.com/talks



Tuesday, February 17, 15

Copyright (c) 2014-2015, Dean Wampler, All Rights Reserved, unless otherwise noted.
Image: The London Eye on one side of the Thames, Parliament on the other.